Selfie: Introduction to the Implementation of Programming Languages, Operating Systems, and Processor Architecture

Christoph Kirsch and Sara Seidl Department of Computer Sciences University of Salzburg 2018

selfie.cs.uni-salzburg.at/slides

Copyright (c) 2015-2018, the Selfie Project authors. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- 1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- 2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

THIS MATERIAL IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS MATERIAL, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

What?

- This presentation is part of a 3-semester introduction to architecture, compilers, and operating systems using the selfie system.
- Selfie is a minimal, fully self-contained 64-bit implementation of a self-compiling C compiler, RISC-V emulator, and RISC-V hypervisor. Selfie can compile, execute, and virtualize itself any number of times!
- Everything needed is implemented in a single, 10-KLOC file of C code.
 There are no includes and no external libraries. For bootstrapping, a C compiler will do.
- The purpose of selfie is to teach the implementation of programming languages (C subset), operating systems (virtual memory, concurrent processes), and processor architecture (RISC-V subset), all using the same code!

Why?

- Selfie shows, from first principles, how a minimal but still realistic hardware and software stack works.
- Anyone understanding elementary arithmetic and Boolean logic may be able to follow!
- The goal is to see and understand how the semantics of programming languages and the concurrent execution of programs is constructed using nothing but bits.
- The self-referential nature of the construction is an important part of selfie. Seeing how to resolve it establishes a well-founded understanding of basic computer science principles.

How?

- We use the selfie code to explain everything. Students will need to read, understand, and modify the code to follow the course.
- Instead of copying code here we provide clickable links to actual selfie code snippets on github.com
- Standard terminology is introduced with clickable links to wikipedia.org

Course Material

- Website: <u>selfie.cs.uni-salzburg.at</u>
- Slides (draft): <u>selfie.cs.uni-salzburg.at/slides</u>
- Book (draft, outdated): <u>leanpub.com/selfie</u>
- Sources (code, slides, book): github.com/cksystemsteaching/selfie

- Install selfie on your machine or in the cloud using the instructions provided in the selfie repository on github.com
- Please note that the slides are incomplete as of 2018 and published incrementally as they become available.

Syllabus

- Programming in C*, the C subset in which selfie is written and compiles.
- Introduction to RISC-U, the RISC-V subset targeted, emulated, and virtualized by selfie.
- 3. Introduction to starc, the selfie compiler (scanner, parser, type checker, register allocator, code generator).
- 4. Introduction to mipster, the selfie emulator (virtual and physical memory, machine contexts).
- 5. Introduction to hypster, the selfie hypervisor (virtual memory, context switching).
- 6. Introduction to monster, the selfie symbolic execution engine (planned).

Programming in C*

- C* is a tiny subset of the <u>programming language C</u>
- C* supports only 2 <u>data types</u>: unsigned integer, uint64_t, and pointer to unsigned integer, uint64_t*. There are no signed integers and no composite data types.
- C* features the <u>unary</u> * <u>operator</u> as the only means to access heap memory hence the name C*. There are no arrays and no structs in C*.
- C* features 5 <u>statements</u> (assignment, if-else, while loop, procedure call, return).
- C* has 3 types of <u>literals</u> (signed decimal number, character, string).
- C* supports 5 <u>arithmetic operators</u> (+, -, *, /, %) and 6 <u>comparison</u> operators (==, !=,<,<=,>,>=). There are no bitwise operators and no Boolean operators.

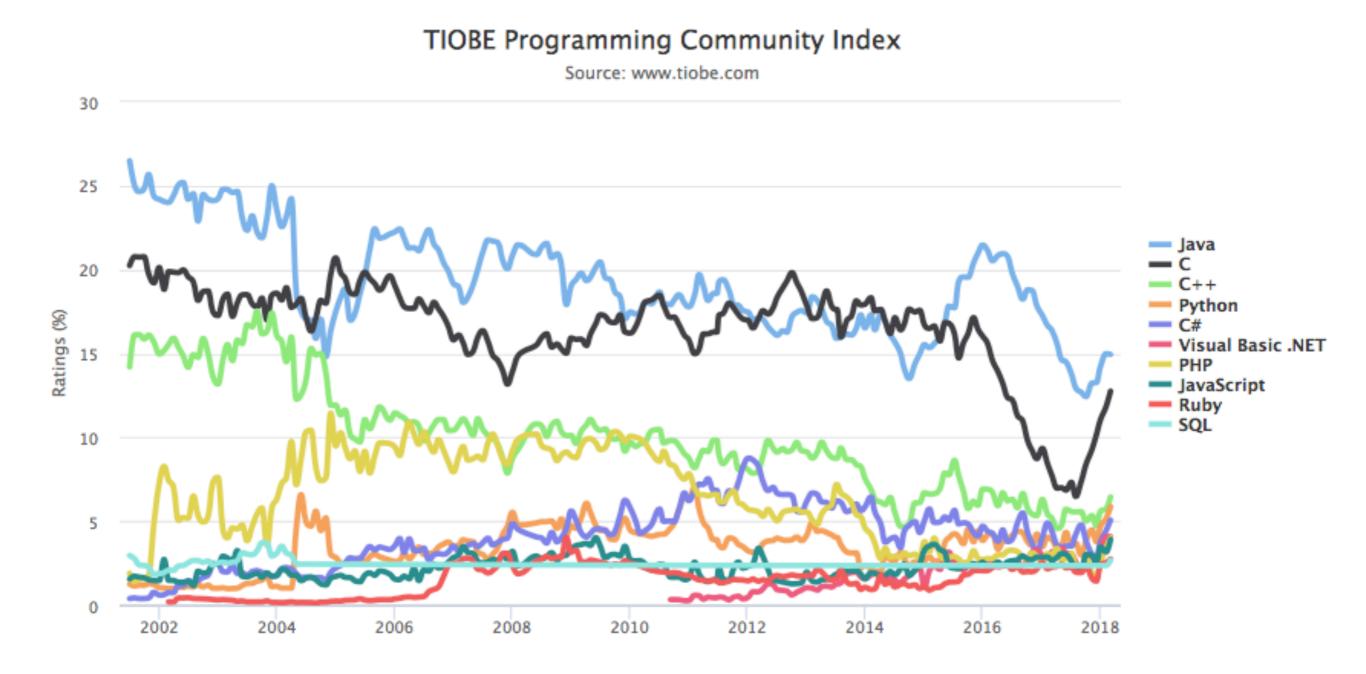
```
uint64 t atoi(uint64 t* s) {
  uint64 t n;
  n = 0;
  // loop until s is terminated
  while (*s != 0) {
    // use base 10, offset by '0'
    n = n * 10 + *s - '0';
    // go to next digit
    s = s + 1;
  return n;
}
```

atoi stands for **ASCII** to integer

Given a string s of decimal digits, the atoi code computes the numerical value n represented by s

Click atoi to see the actual code in selfie. We provide such links throughout the presentation.

Programming Language C

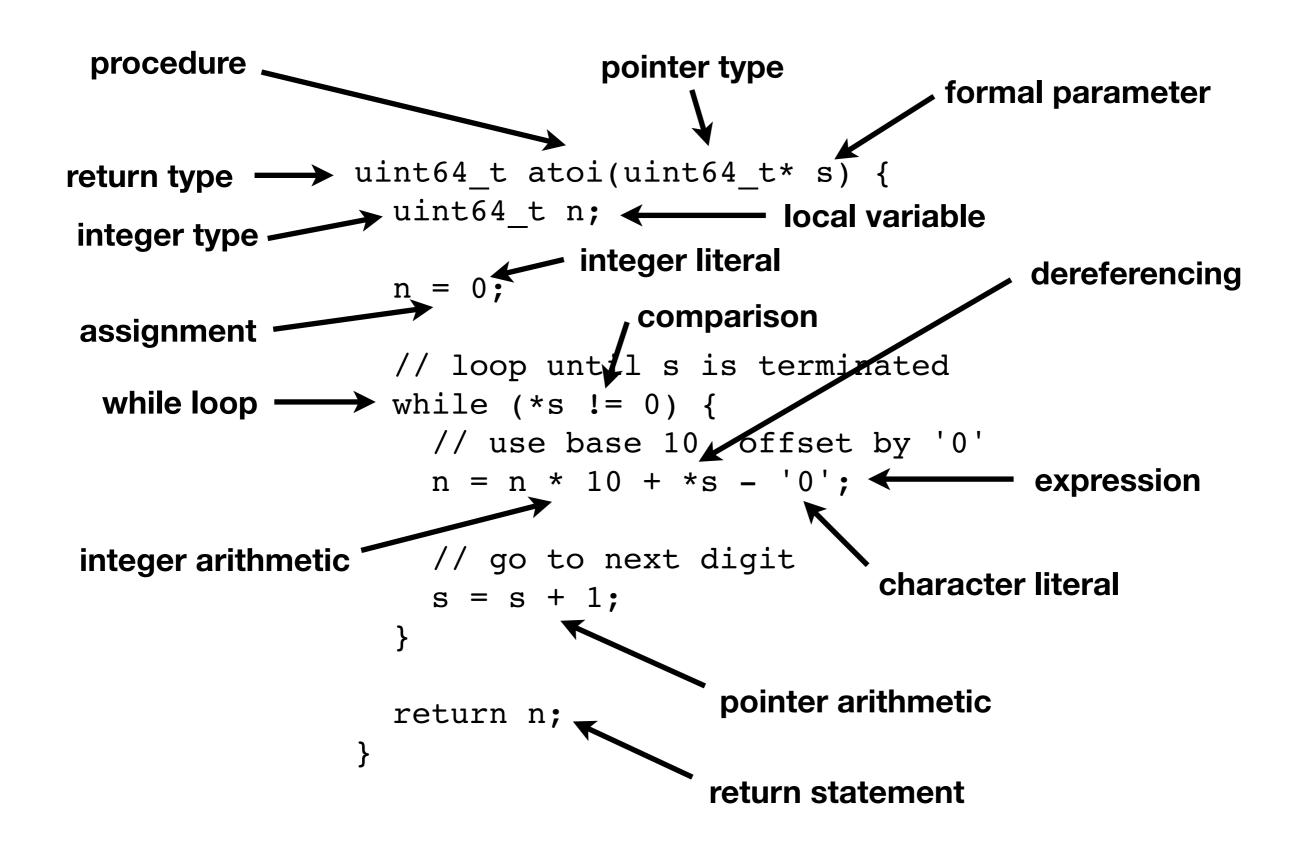


C with all its dialects is still the most popular programming language https://www.tiobe.com/tiobe-index

Why C?

- This course aims at much more than just learning how to code.
- The choice of programming language is important but not more important than many other choices made here.
- C is relevant, C* is simple.
- With <u>imperative programming</u> the notion of program <u>state</u> maps oneto-one to the notion of machine state.
- Seeing the pitfalls of imperative programming enables students truly appreciate other paradigms such as <u>functional programming</u>.
- In fact, awareness of state makes students understand the value of functional programming. This is nevertheless left to others to teach.

C* by Example



C* Integers and Pointers

- C* integers are unsigned 64-bit integers
- ► In C* there are five arithmetic operators: +, -, *, /, %
- And six comparison operators: ==, !=, <, <=, >, >=
- C* pointers are 64-bit pointers to C* integers
- And pointer arithmetic: +, -

C* versus C Integer Literals

- C* integer literals are unsigned 64-bit
- C integer literals are signed 32-bit
- For example, 1/-1==0 in C* but 1/-1==−1 in C
- ► And, 1%-1==1 in C* but 1%-1==0 in C
- Also, 1<-1 and 1<=-1 hold in C* but 1>-1 and 1>=-1 do not whereas the opposite is true in C
- ► The semantics of / and % as well as <, <=, >, and >= is different for signed and unsigned integers!

C* Characters and Strings

- C* characters are <u>ASCII</u>-encoded.
- C* character literals are characters in code like 'c'.
- C* strings are stored as null-terminated sequences of characters. Alternatively the end of a string could be identified by storing the number of characters at its beginning.
- C* string literals are strings in code like "this".
- ► The difference between 'a' and "a".

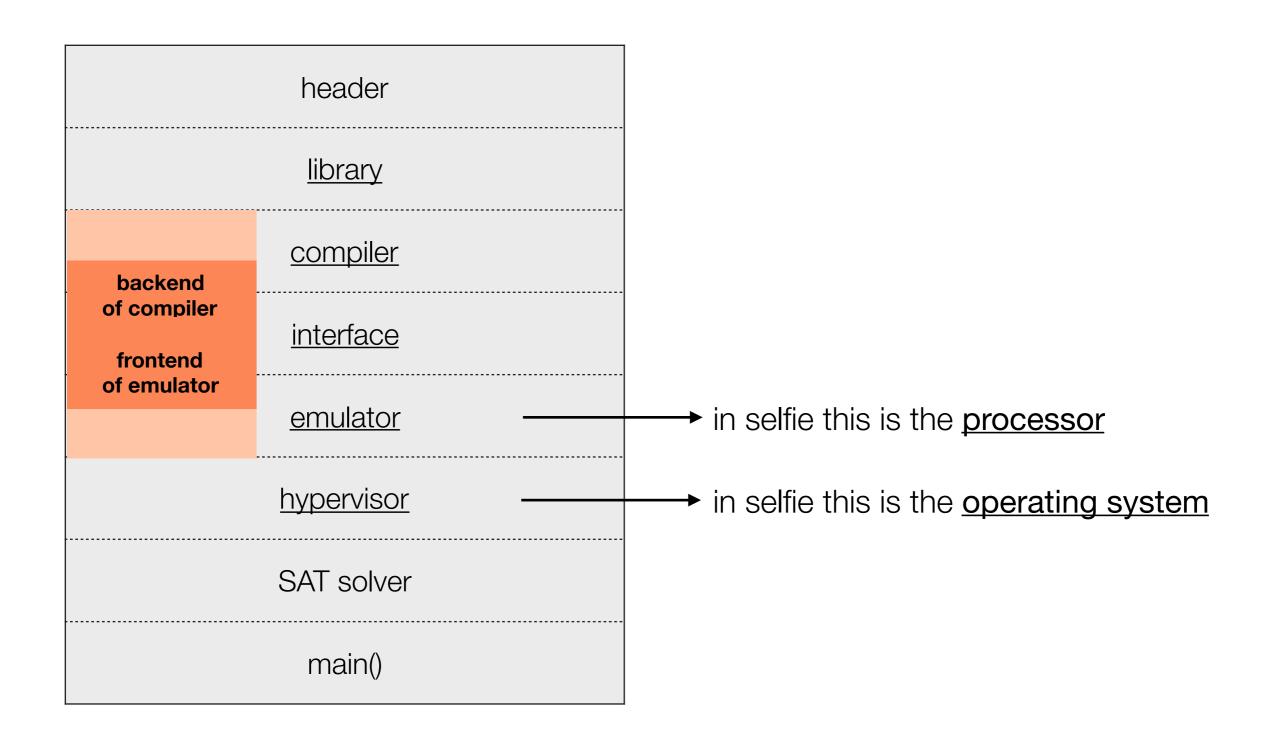
ascii representation (**numerical value**) in memory

pointer to first word of where the string is stored

C* versus C Strings

- C* strings are arrays of unsigned 64-bit integers.
- C strings are arrays of characters, that is, of type char.
- ► For example, *"Hello World!" is equal to 0x6F57206F6C6C6548 in C* but 0x48 in C.
- Note that 0x48 is ASCII for H while 0x65, 0x6C, 0x6F, 0x20, and 0x57 are ASCII for e, 1, o, space, and W, respectively.

Selfie Code Structure

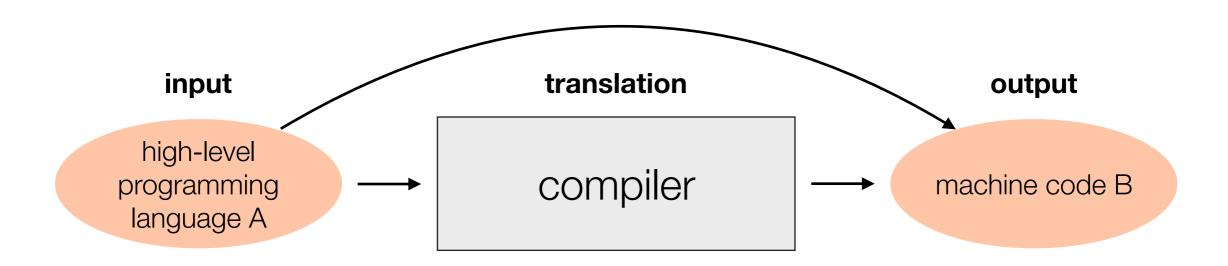


The Selfie Library

```
uint64 t leftShift(uint64 t n, uint64 t b);
uint64 t rightShift(uint64 t n, uint64 t b);
uint64 t getBits(uint64 t n, uint64 t i, uint64 t b);
uint64 t getLowWord(uint64 t n);
uint64 t getHighWord(uint64 t n);
uint64 t abs(uint64 t n);
uint64 t signedLessThan(uint64 t a, uint64 t b);
uint64 t signedDivision(uint64 t a, uint64 t b);
uint64 t isSignedInteger(uint64 t n, uint64 t b);
uint64 t signExtend(uint64 t n, uint64 t b);
uint64 t signShrink(uint64 t n, uint64 t b);
uint64 t loadCharacter(uint64 t* s, uint64 t i);
uint64 t* storeCharacter(uint64 t* s, uint64_t i, uint64_t c);
uint64 t stringLength(uint64 t* s);
void
         stringReverse(uint64 t* s);
uint64 t stringCompare(uint64 t* s, uint64 t* t);
uint64 t atoi(uint64 t* s);
uint64 t* itoa(uint64 t n, uint64 t* s, uint64 t b, uint64 t a, uint64 t p);
```

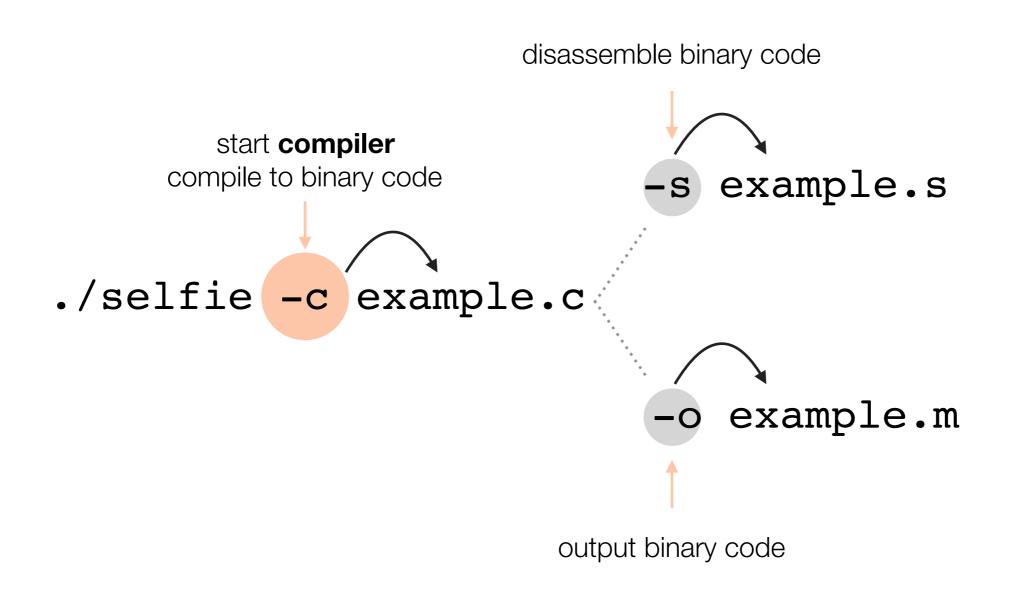
Introduction to Compilers

A Compiler

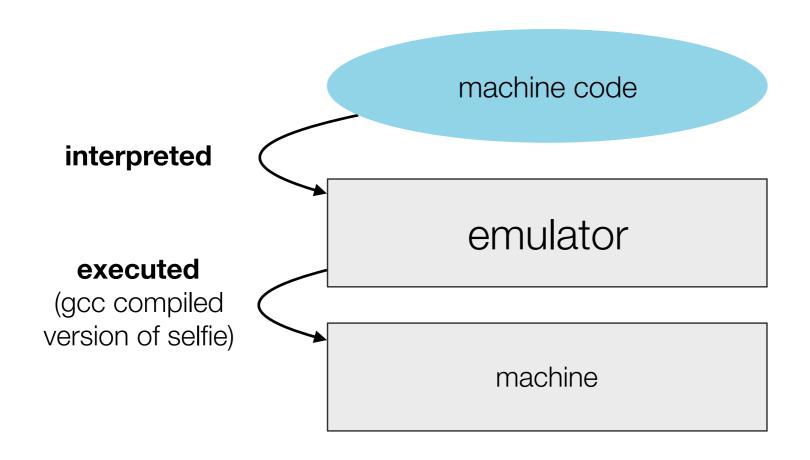


- The selfie compiler written in C* translates C* code (self-referential).
- High-level languages have a structure that defines the control flow.
- Machine code has no structure, it is just a sequence of instructions.
- The compiler reads an input program, which is a sequence of characters (ASCII, UTF-8-encoded), and writes machine code.

Using the Selfie Compiler



An Emulator

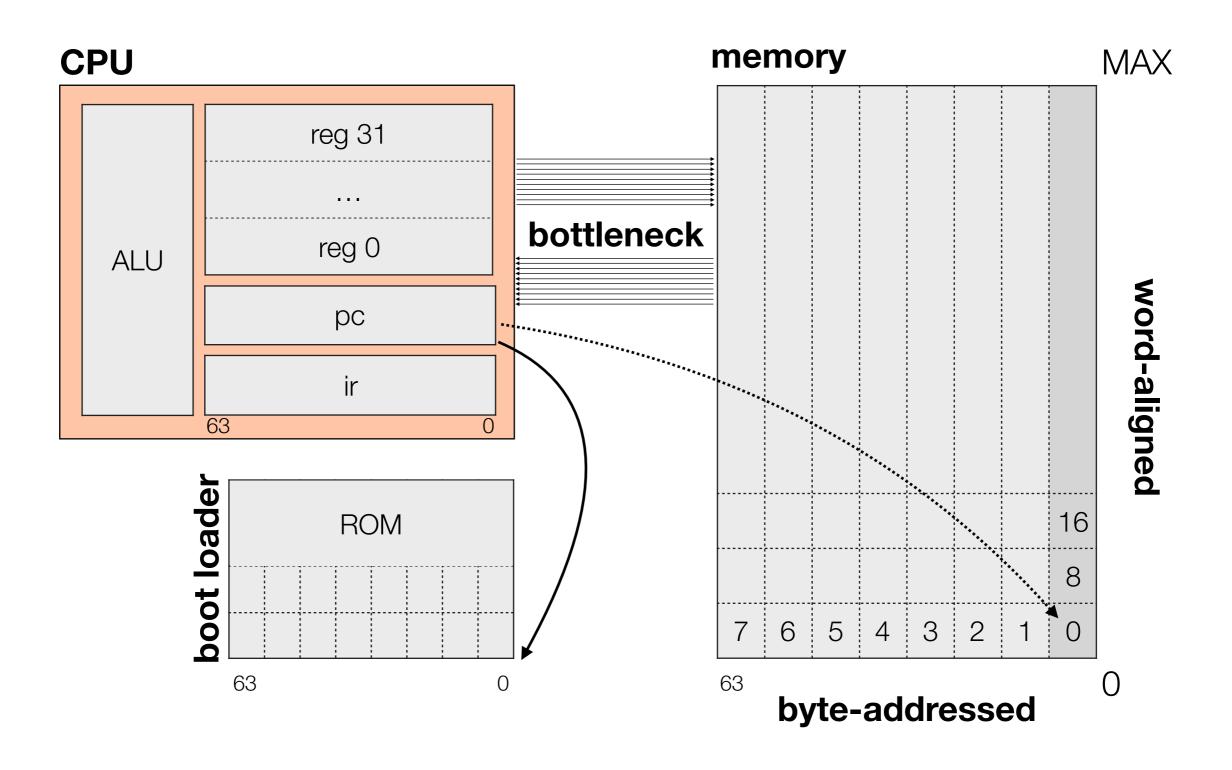


The selfie <u>emulator</u> interprets machine code generated by the selfie compiler.

Using the Selfie Emulator

load binary code start mipster and emulate binary code with 1MB of virtual memory ./selfie example.m ./selfie -c example.c start starc compiler and compile to binary code

Von Neumann Machine



More about the Von Neumann architecture

Von Neumann Machine

- Key idea code and data in same memory
- How do we know what code is and what data is?
 - The program counter points to an instruction in memory making it code. This instruction may instruct the processor to load bits from memory and modify them making these bits data.
- Bootstrapping Loading the first program
 - At the very beginning hardware support is needed to set the PC to address 0 of the memory where the boot-loader code is stored.
 - The boot-loader is code stored in non-volatile memory (ROM) that instructs the processor to load code.
 - Last instruction sets the PC to address 0 of main memory.

States Machine State and Program State

Machine State

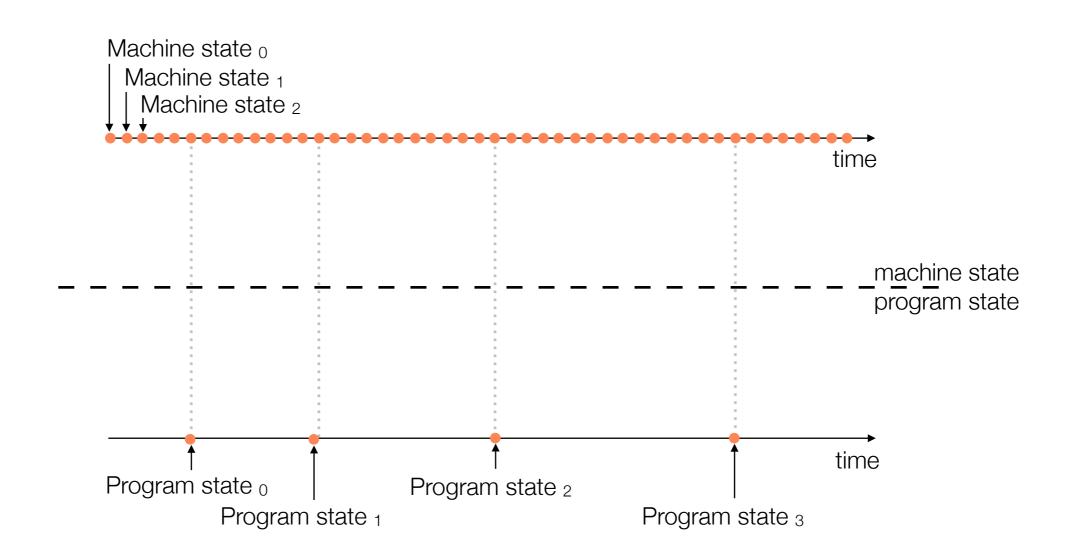
- defined by the all the bits in the machine
- #machine states = #bits ²

Program State

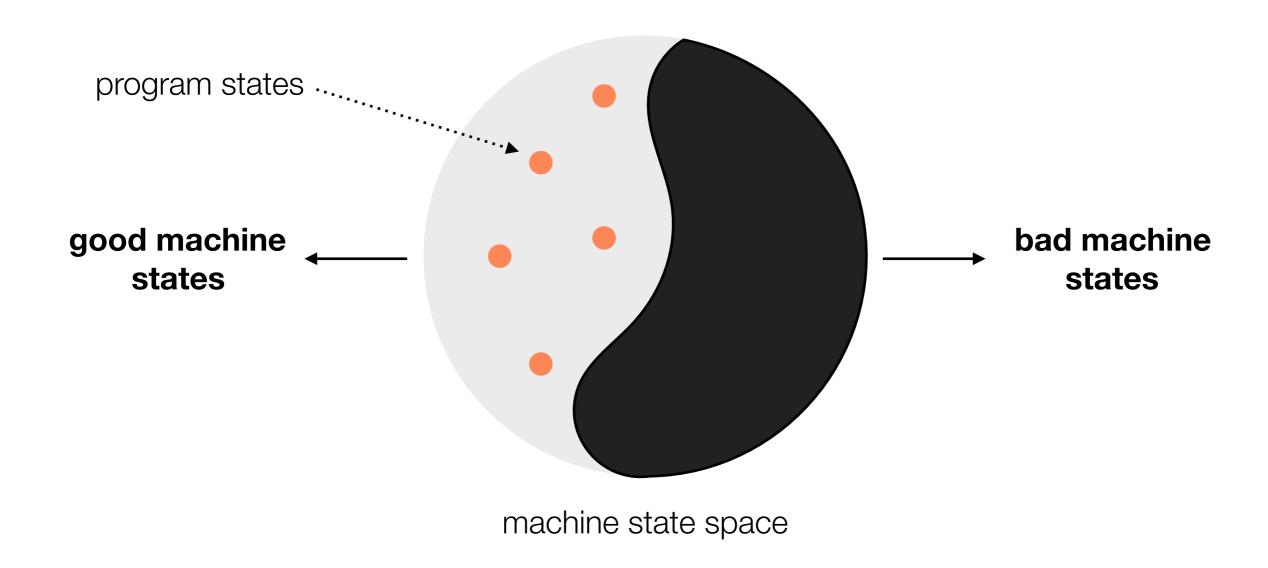
- global variables
- call stack current statement
- dynamically allocated memory (state space explodes)
- #program states much smaller then #machine states

State space High Level Programming Languages

- each program state corresponds to some machine state
- far less program states (high-level) than there are machine states



Correctness



A correct program never takes the machine into the "bad state" space.
 (for any execution)

Introduction to RISC-U

- RISC-U is the <u>RISC-V</u> subset targeted, emulated, and virtualized by selfie.
- There are 14 instructions, each 32-bit wide, the processor knows and the compiler generates code for.
- When talking about formal languages it is important to distinguish between the <u>syntax</u> and the <u>semantics</u> of that language.

Syntax of RISC-U

At machine code level we distinguish further between the human readable <u>assembly</u> code and the <u>binary</u> code the CPU executes.

addi \$t02 \$t01 4

0000000010000110000001110010011

Syntax of RISC-U

- Different instructions have different binary encodings. Instructions are encoded using special formats.
- The format of an instruction specifies how the 32 bits are interpreted. It is designed in a way that allows fast decoding of the instructions.
- Selfie always stores binary code which can then be <u>disassembled</u> to obtain the assembly code.

Semantics of RISC-U

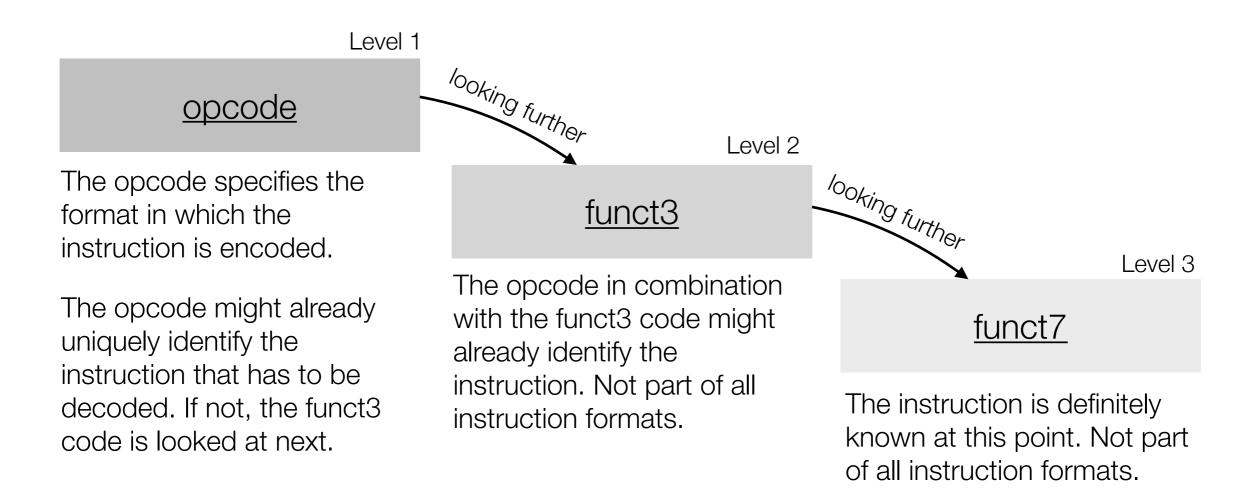


- The semantics of an instruction is determined by how the processor implements it.
- ► In selfie this implementation is done in the do_.. 's, like do_addi() or do_sltu().

Decoding Instructions



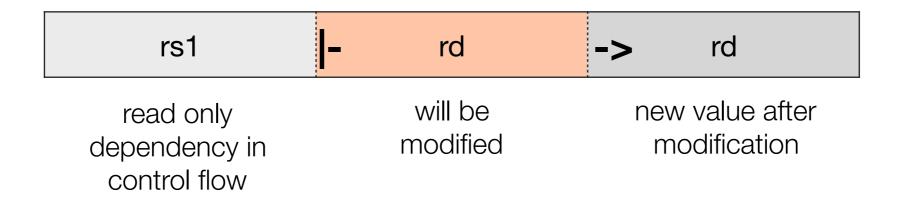
- Each instruction can be uniquely identified by certain parts of the 32 bits called opcode, funct3 and funct7 code.
- When the instruction is known, the meaning of the remaining bits of the instruction becomes clear.



Instructions and Machine State



Selfie uses special procedures (<u>before()</u> & <u>after()</u>) that show on which part of the machine state they depend, which part they modify and the modification itself.



- This information is enough to determine the machine state at any point of execution (completely deterministic).
- The only way to inject information from outside that is not known beforehand is through the read call.

Instructions and Formats



R	ISC	:- U

lui	beq	
addi	jal	
add	jalr	
sub	ld	
mul	sd	
divu	sltu	
remu	ecall	

Formats

R-format
<u>I-format</u>
<u>S-format</u>
<u>B-format</u>
<u>J-format</u>
<u>U-format</u>

special case of addi

nop

lui	\$rd	imm	
addi	\$rd	\$rs1	imm

- Load upper immediate loads imm value shifted by 12 bits into \$rd and add immediate adds imm to the content of \$rs1 and stores the result in \$rd.
- These two instructions are used to initialize registers (\$rs1 = \$zero) and to load addresses into a register in order to read values from memory.
- ▶ lui is used to load the upper and addi to load the lower bits see load integer(uint64 t value).
- ► A special case of addi is nop, with \$zero = \$zero + 0.

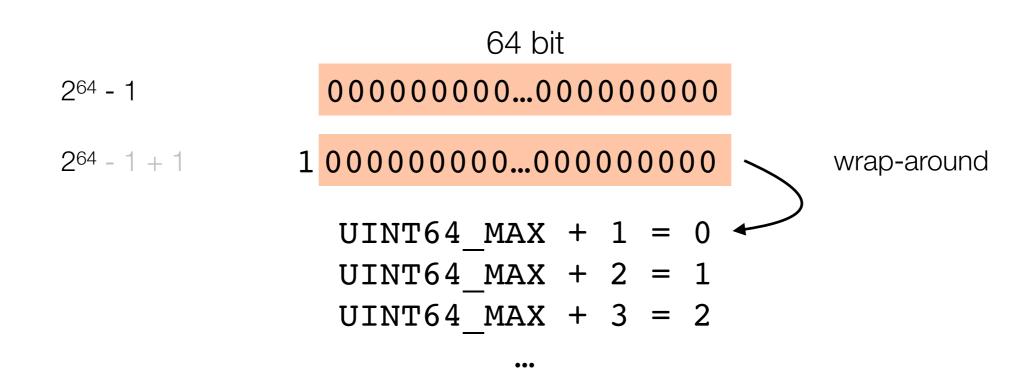
Arithmetic Instructions

add	\$rd	\$rs1	\$rs2
sub	\$rd	\$rs1	\$rs2
mul	\$rd	\$rs1	\$rs2
divu	\$rd	\$rs1	\$rs2
remu	\$rd	\$rs1	\$rs2

The processor executes these instructions using <u>unsigned integer</u> <u>arithmetic</u> with <u>wrap-around semantics</u>.

Wrap-around Semantics

- Cause of unbelievably expensive bugs, e.g. the <u>Ariane 5 Flight 501</u>.
- ▶ 2⁶⁴ 1 is the largest value that can be represented by 64 bits. In selfie this value is denoted UINT64 MAX.
- Adding 1 to UINT64_MAX leads to a wrap-around where only the 64 LSB are considered.



Arithmetic Instructions

sltu \$rd \$rs1 \$rs2

- Set \$rd to 1 if \$rs1 < \$rs2.</p>
- ► This is the only instruction needed to implement <, >, <=, >=, == and !=.
- How this is done:
 - '==' is implemented using b a < 1.

In unsigned arithmetic only 0 satisfies this condition.

Memory Instructions

ld	\$rd	offset(\$rs1)
sd	\$rs2	offset(\$rs1)

- ▶ Load into \$rd the value that is stored at the address that is obtained by adding the immediate offset to the content of \$rs1.
- ▶ **Store** the content of \$rs2 at the address that is obtained by adding the immediate offset to the content of \$rs1.
- ► The <u>addressing mode</u> used for these instructions is called registerrelative addressing.

Control-Flow Instructions

- Control flow, at machine code level, is the order in which instructions are executed.
- All previous instructions feature implicit trivial control flow, that is, they simply set the program counter to the next instruction. Their main purpose is data manipulation.
- The following instructions have a more sophisticated effect on control flow.

Control-Flow Instructions

7	beq	\$rs1	\$rs2	imm
	jal	\$rd	imm	
	jalr	\$rd	\$rs1	imm

- The first two instructions use a different addressing mode called pcrelative addressing at the resolution of 12 bit.
- ▶ Branch on equal sets the pc to pc + imm if the content of \$rs1 matches \$rs2.
- Jump and link is used for procedure calls and stores the return address (address of next instruction) in \$rd.
- Jump and link register is similar to jal, except that it uses registerrelative addressing.

Link Register

```
uint64_t increment(uint64_t inc) {
  return inc + 1;
}

uint64_t main() {
  uint64_t a;

a = 0
  a = increment(a);

return a;
}
```

The next **instruction** that is linked is the assignment into a!

Operating System Support



ecall

- This is not a standard machine instruction but rather a programmable instruction that allows you to request services from the operating system, also known as system call.
- The arguments that specify what action we need the OS to take are provided in \$a7.
- Selfie supports these ecalls: read, write, open, malloc, and exit.
- ► The implementation of each of these ecalls is in <u>handle system call()</u>.

Introducing Language Operators

- In selfie there are no <u>bitwise operators</u> and no native machine instructions for bitwise operations.
- The next section shows how new language operators can be systematically implemented.

Bitwise Operators

- Bitwise operators in C perform operations on bit level. There are:
 - logic bitwise operators: & (bitwise AND), | (bitwise OR) and ~ (bitwise NOT).
 - shift bitwise operators: << (bitwise left shift) and >> (bitwise right shift).
- Logical shifts (zeros are shifted in) or <u>arithmetic shifts</u> (sign bit is shifted in).

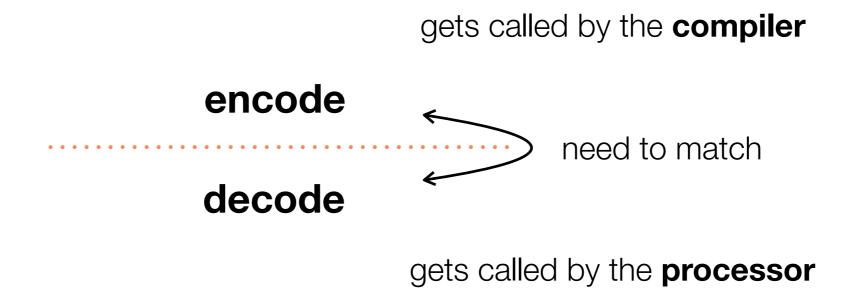
Bitwise Operators in Selfie



- The semantics of the bitwise operators << and >> on the uint64_t type is that of a logical shift.
- In selfie there are library functions that perform shift operations (see <u>here</u>):
 - left_shift(): n << b = n * 2b
 - right_shift(): $n \gg b = n/2^b$

Language Operators

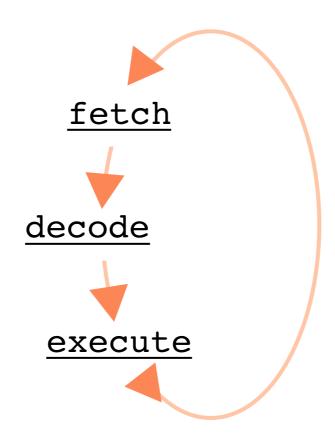
- The compiler has to recognize the operator and generate code for it.
- The processor has to understand the instructions encoded by the compiler and executes them.
- We first expand the processor by implementing a new machine instruction that can then be used by the compiler to generate code.



Implementing Machine Instructions

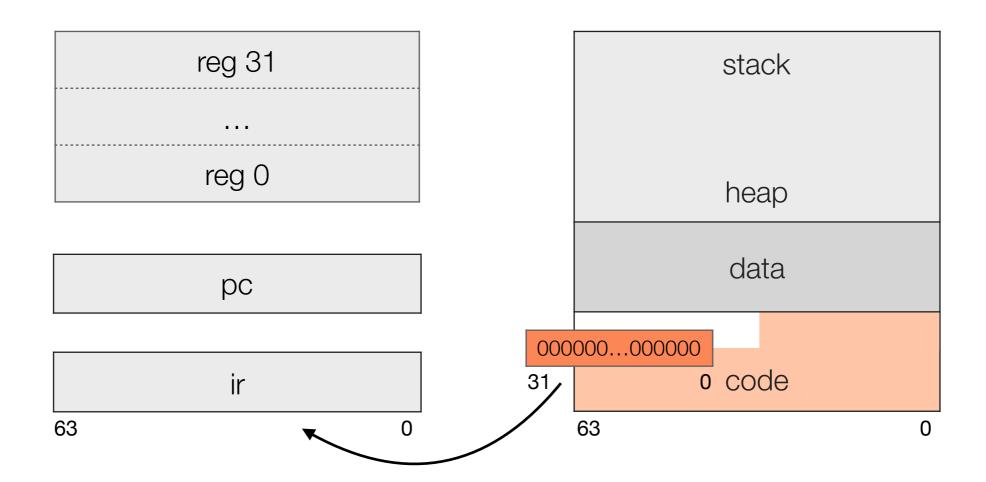
- Extending the Processor Implementing machine instructions for SLL and SRL.
- Before implementing an instruction the semantics it is supposed to have must be clear.
- The implementation of an instruction determines its semantics. Truly understanding each instruction is key to understanding the target language generated by the compiler and executed by the processor.

The heart of the selfie emulator is the procedure run until exception()



Fetch

- A 32-bit instruction gets fetched from memory and is then stored in the instruction register (ir)
- 2 instructions stored in memory using hi/low word



Decode Remember



- The opcode specifies the instruction format in which the instruction is encoded.
- After decoding the binary code of the instruction, the processor knows what instruction and parameters it is dealing with.
- Example <u>ADDI</u>:
 - Immediate instructions contain bits that are interpreted as numerical value that is retrieved using shifting operations.
 - 5 bit are enough to encode 32 registers.

immediate	rs1	funct3	rd	opcode
12 bit	5 bit	3 bit	5 bit	7 bit

Execute

- The execution of every RISC-U instruction has a well-defined effect. It changes the state of the machine only at a specific location involving little data.
- At most two registers or one register and one memory location are modified by an instruction.
 - Every instruction modifies the PC.
 - Most instructions which modify data (another register or memory location) have trivial control flow (PC to next instruction).
 - Control-flow instructions have a more sophisticated control flow, that is, they may change the PC using relative or absolute addressing.

Execute



semantics: 64-bit unsigned addition with wrap-around semantics

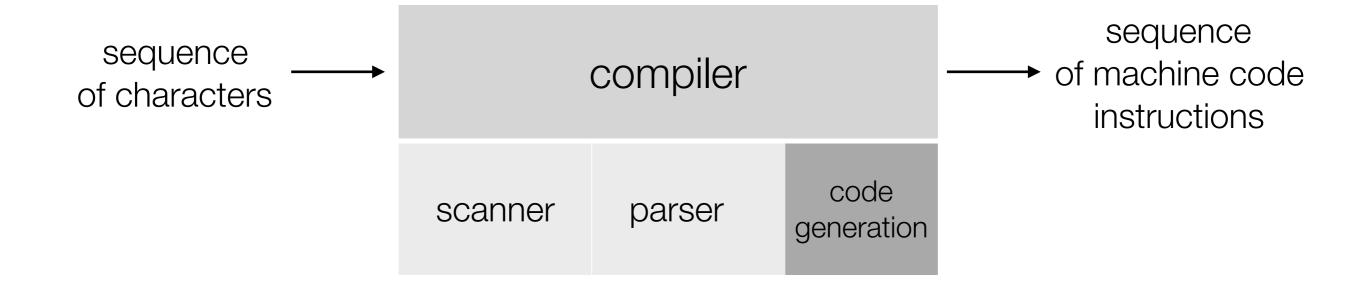
Example <u>ADDI</u>:

- Bits in \$rd are overwritten with \$rs1 + imm.
- PC = PC + INSTRUCTIONSIZE.
- Used for initialization loading constants into registers.

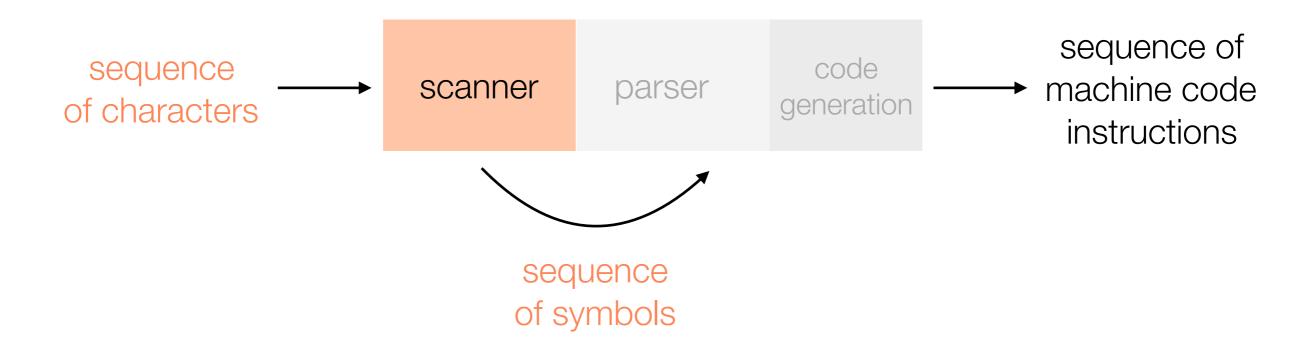
Realize...

- This is how virtually every general-purpose processor operates.
- Instructions only have a tiny effect on the overall machine state.
- What makes modern computation so powerful is:
 - the incredible speed.
 - the enormous size of memory.

The Compiler



The Scanner



The Scanner

The scanner performs a **membership test** on a sequence of characters. It checks whether the symbols it sees are **valid symbols** of the programming language.

- Reading only one character at a time the scanner transforms a sequence of characters into a sequence of symbols that has no structure.
- There are single-character symbols ('<', x) and multi-character symbols ('<=','variable'). All symbols are uniquely identified by a token (integer).
- The whole sequence of symbols is not stored. Only the last read symbol is remembered.

What are valid symbols?



What is needed:

A specification of valid symbols -> <u>regular expression</u> and a formalism to write such a specification.

In selfie...

the chosen formalism is **EBNF** (see in the grammar.md file).

Regular Expression

- Regular expression is a formalism that defines a <u>regular language</u> (= a set of symbols defined by regular expression).
- The production rules consists of <u>terminal symbols</u>, <u>non-terminal</u> <u>symbols</u> and operators (repetition { }, concatenation □, xor | , ...).
- A regular expression can be reduced to a single rule by replacing every non-terminal symbol with its right-hand side until no non-terminal symbols are left.

non-terminal symbol

terminal symbol (literals of language)

EBNF Niklaus Wirth

```
ebnf = { production } .
production = identifier "=" expression "." .
expression = term { " | " term } .
term = factor { factor } .
factor = identifier | string | "(" expression ")" |
         "[" expression "]" | "{" expression "}" .
string = """ printableCharacter { printableCharacter }""" .
printableCharacters = "a" | ... | "!" .
identifier = "ebnf" | "production" | ... | "identifier .
```

The Scanner

The problem statement:

Given a specification of valid symbols where symbols may consist of multiple characters...

...Is a given input (source code) in this specification, or in other words do the characters in the source file form valid symbols according to the specification?

Abstraction

- A scanner for the English language accepts this sentence.
- A sequence of characters gets assembled symbols (words and punctuation).

This_sense_makes_!_no_sentence

Specification and Implementation

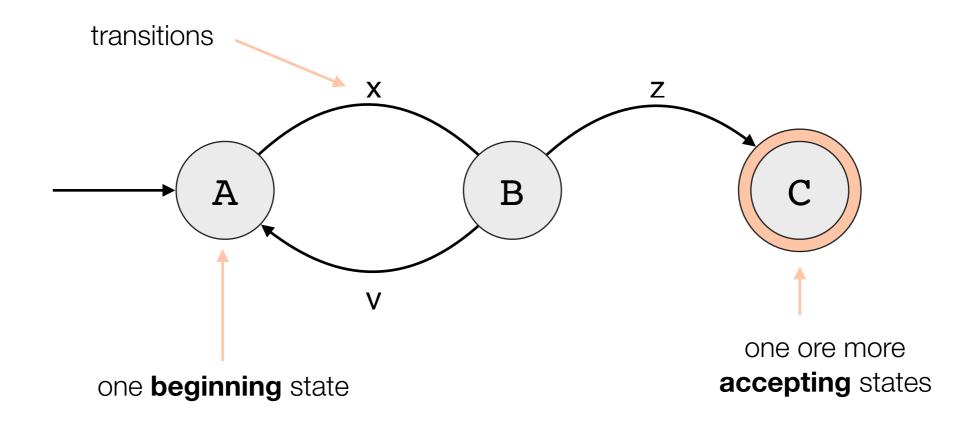
Specification

The specification of valid symbols is a regular expression written in EBNF. It is easier to reason about the correctness of regular expression and it is expressive enough. The number of valid symbols is infinite and their size is unbounded.

Implementation

The scanner implements this specification using a <u>finite state machine</u> (FSM), a computational model, that recognizes/accepts the set of regular expressions. The implementation is simple and can be done very efficiently.

FSM



- Transitions describe how to get from one state to another.
- ► E.x. When in state B the condition z (may be "reading z") is true, a move to state C happens.

```
get_symbol()

calls

get_character()

find_next_character()
```

- The heart of the scanner is the procedure get_symbol(), which finds the next valid symbol in a sequence of characters and stores it in a global variable (symbol).
- It internally uses two procedures get_character() and find next character().
- get_character() simply reads the next character of the input stream and stores it into a global variable (character).
- find_next_character() is the implementation of whitespace.

Sequence of Characters

- The procedure get_character() is used to read the next byte into a buffer.
- The return value is either the number of read bytes or 0, which indicates end of file (EOF).
- Any input file is nothing more than a sequence of characters to the scanner. It will even accept meaningless sequences as long as it consists of valid symbols in the specified language.

Whitespace



- Whitespace and comments do not change the semantics of code and could even be omitted (minification). Their purpose is to make the code more readable.
- The scanner implements whitespace and comments by simply ignoring them using the procedure <u>find_next_character()</u>.

The Scanner getSymbol()



The procedure isCharacter Identifier Letter0r get symbol() is the ACCEPT or Keyword else DigitOr implementation of the Underscore() i >= maxscanner as a FSM. IdentifierLength $I\K$ **ERROR** ischaracterletter isCharacter atoi() Digit() else isCharacterDigit() Number N ACCEPT character != EOF CHAR SINGLEQUOTE Character Start ACCEPT C c_{HAR_PLUS} Ρ **ERROR** Every case for which there is no transition, leads to SYNTAX ERROR case.

Properties

The scanner never crashes and always terminates: In code the ESM is simply a huge if/else, if/else, states

In code the FSM is simply a huge if/else if/else statement that covers every possible case. Syntax errors fall through to the else case.

► The scanner is good at forgetting:

As soon as the scanner reaches a state it has been in before everything in between is forgotten.

Correctness:

The state space of a machine is usually huge or even infinite and therefore reasoning about it is hard. An FSM however has only very few states (finitely many) and reasoning about them is much easier.

Properties

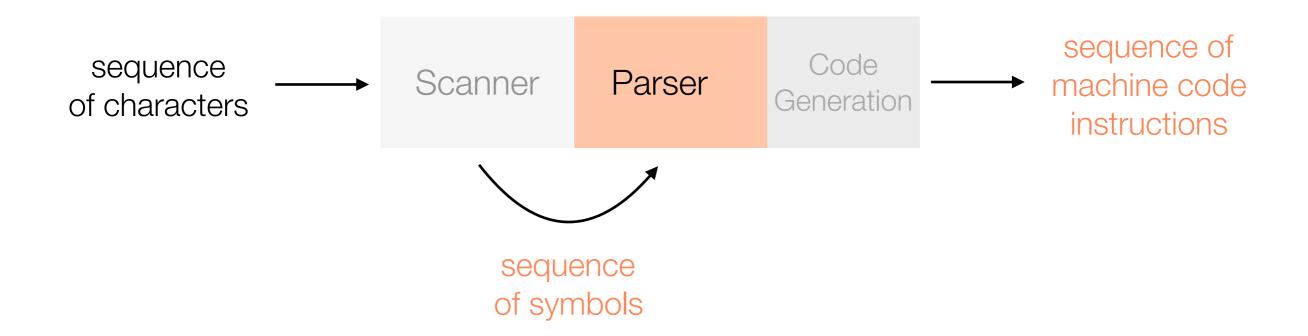
The scanner can not recognize structure:

It just operates on a sequence of characters and turns it into a sequence of symbols.

To recognize structure, the scanner would need to be able to **count** and remember (e.g. the number of braces).

Therefore we need something more powerful called a parser.

The Parser



The Parser

The parser performs a **membership test** on a sequence of symbols. It checks whether the symbols form a **syntactically valid** program.

- The set of syntactically valid programs is infinite.
- The parser builds an internal representation of the structure of the input.
- The technique used in selfie is <u>single-pass</u>, <u>recursive-descent</u> parsing.

What are syntactically valid programs?

What is needed:

A specification of syntactically valid programs -> Grammar and a Formalism to write such a specification.

In selfie...

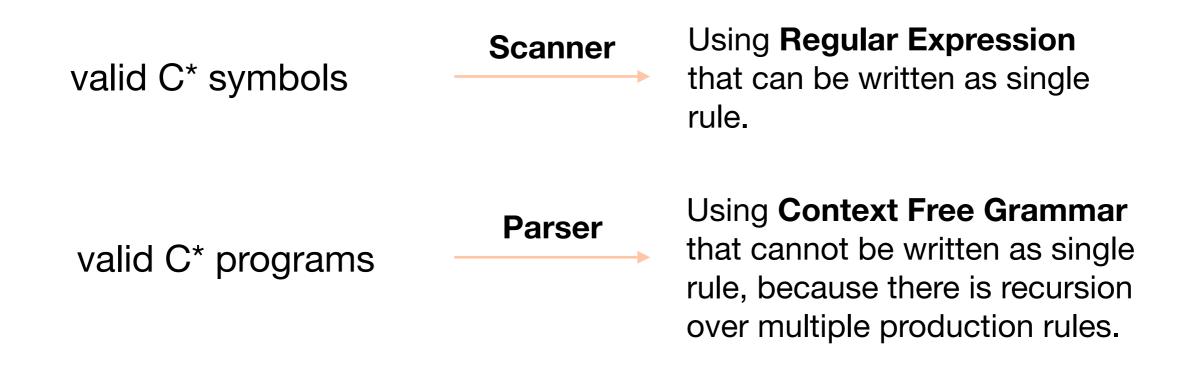
the specification is a <u>context-free grammar</u> (CNF) described using EBNF.

Abstraction

A parser for the English language accepts this syntactically correct sentence (subject, verb, object, punctuation).

This_sense_makes_no_sentence!

Context-Free Grammar



- A CFG is a set of <u>production rules</u>.
- There is only one <u>non-terminal symbol</u> on the left hand side. This symbol can be **substituted in any context** in which it appears on the right hand side.

Parsing Techniques single-pass & multi-pass

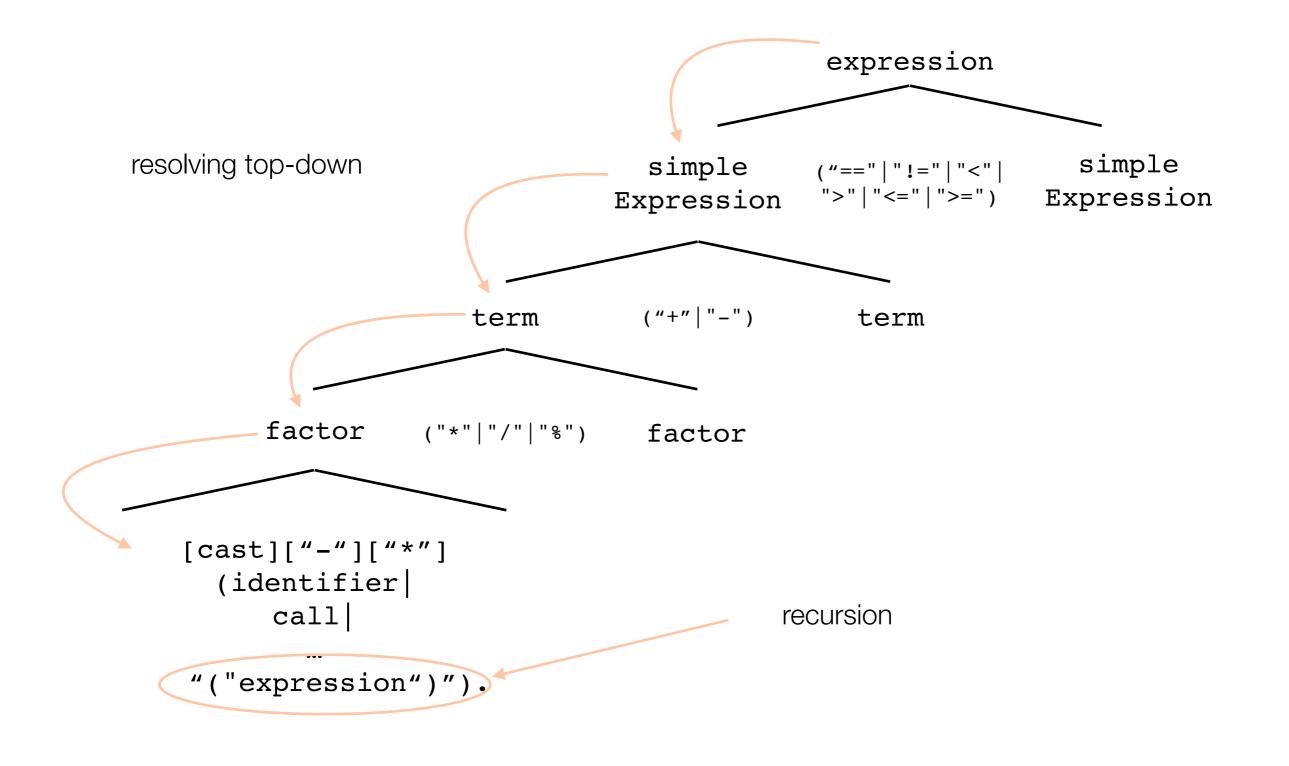
- The internal representation build by the parser is called a syntax tree.
- A <u>single-pass</u> parser parses the input once and never looks back. At any point it only represents a single branch of the syntax tree and never the whole context. This only allows local performance optimization.
- A <u>multi-pass</u> parser builds an internal representation of the whole syntax tree and walks over it multiple times, trying to optimize the performance on a global level.

Parsing Techniques recursive-descent

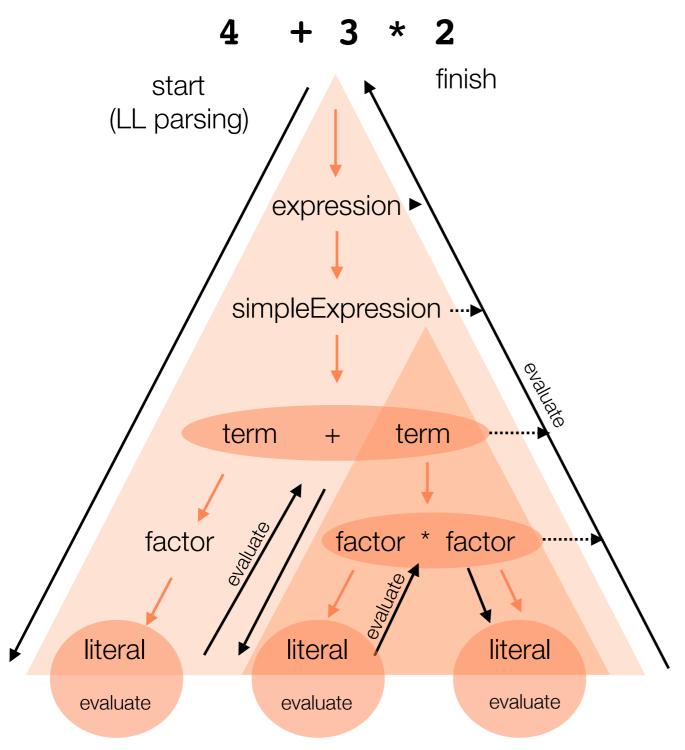
- A <u>recursive descent parser</u> uses recursion and a top-down approach.
- The top-down approach starts by looking at one start symbol. Every symbols that appears on the right-hand side of a production rule is resolved and matched to the input.
- The recursive nature of the parsing rules allows that a symbol that has been resolved in a previous step can appear again on the right-hand side of a production rule.

Example: expression()

```
expression = simpleExpression [("=="|"!="|"<"|">"|"<="|">=") simpleExpression]
```



Parsing by Example



- ► To evaluate an expression we keep resolving all non-terminal symbols until only a terminal symbols is left (e.g literal, '+'). This terminal symbol is then matched against the input (4,+).
- So each branch/subtree is completely evaluated before the result of a parent node can be calculated.

Parsing by Example

- So each branch/subtree is completely evaluated before the result of a parent node can be calculated.
- In the next slides we will see the implications of this statement (on a single-parse parser).

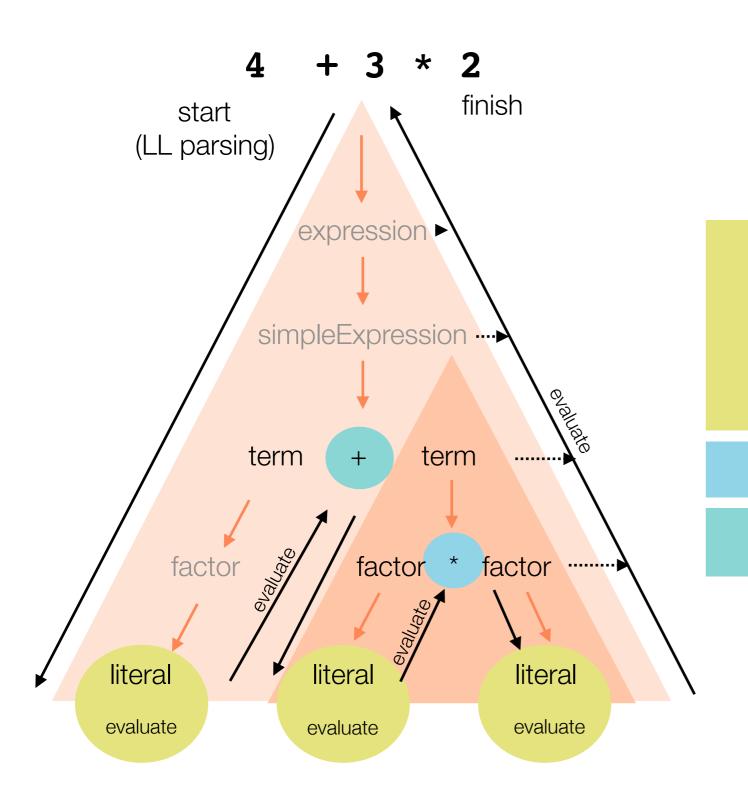
Parsing Technique Implications

- ► The selfie parser is a single-pass parser: The input is read only once, without ever looking back. This means that code must be already generated during the parsing process.
 - → The code generator is part of the parser.
- ► The selfie parser is a top-down parser:
 - Compiling an expression using a top-down approach may require to store intermediate computations temporary in registers. But we need to keep in mind that there is only a very limited amount of those (7 in selfie).
 - → We have to deal with the <u>register allocation problem</u>

Code Generation

- The following example will show...
 - ...(again) how an expression is evaluated by resolving all nonterminal symbols and matching the terminal symbols against the input.
 - that the compiler generates code as soon as possible.

Code Generation



```
addi $t0, $zero, 4

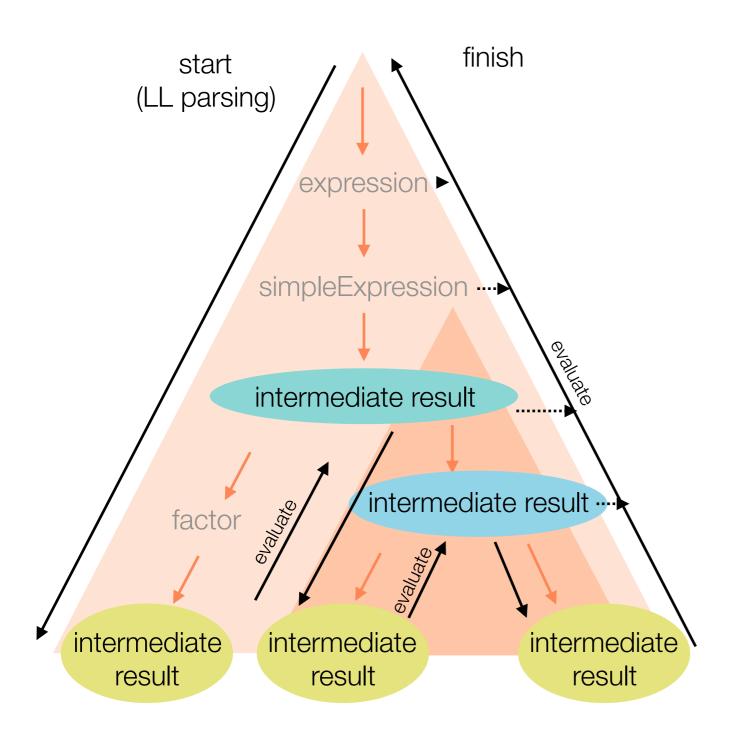
addi $t1, $zero, 3

addi $t2, $zero, 2

mul $t1, $t1, $t2

add $t0, $t0, $t1
```

Register Allocation Problem



```
addi $t0, $zero, 4
addi $t1, $zero, 3
addi $t2, $zero, 2
mul $t1, $t1, $t2
add $t0, $t0, $t1
```

Register Allocation Problem

Register allocation is a compile time problem.

"How do we know which registers are **used** and which are **free** at runtime?"

- When compiling we need to think of the 2 time axis'
 - At compile time we compile and generate code...
 - ...but we have to keep in mind that at runtime there is only a very limited amount of registers available for actually executing the generated code and computing expressions.

Register Allocation Problem



- stack allocator (in selfie using talloc(), tfree()
 - limitation: freeing in reverse order.
 in selfie, this is sufficiently enough and works because of the recursive structure of the parser and code generator.
 - complexity: free, allocate in constant time
- bit vector marking used registers
 - note on complexity: free in constant time, allocate linear in the #registers.

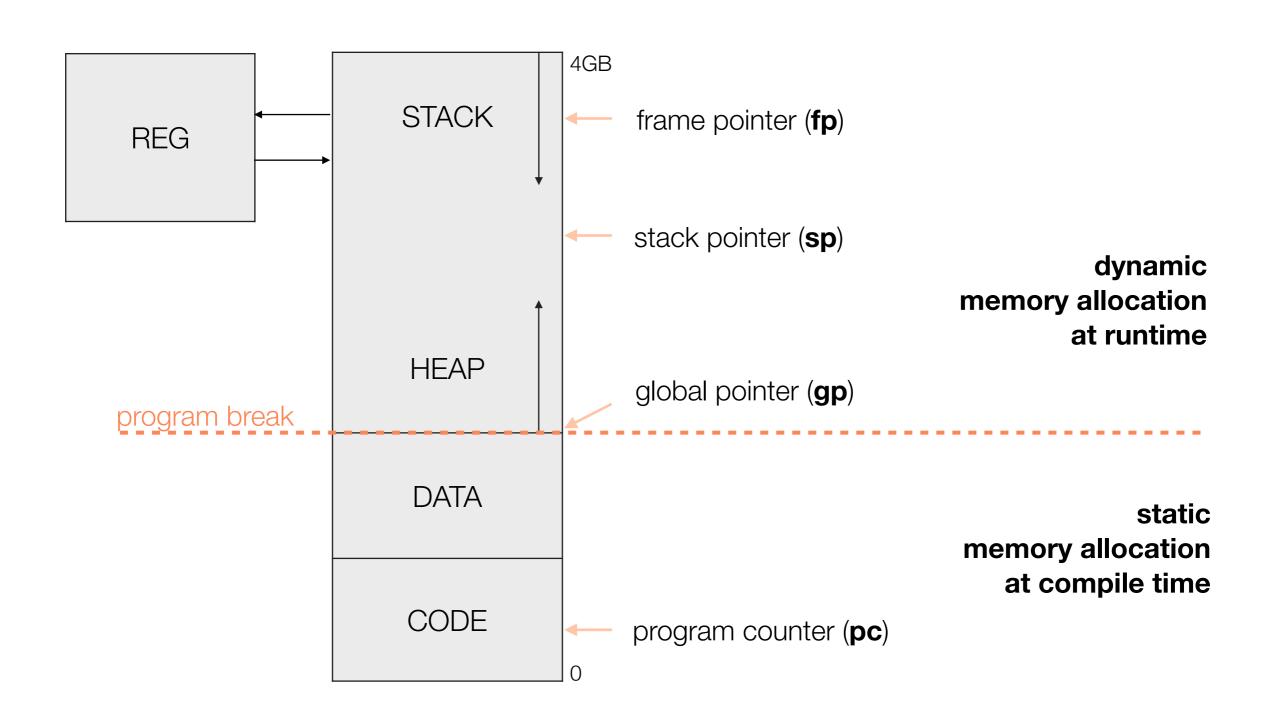
advanced approach

 this is a question of <u>liveness of variables</u> (in registers) and can be reduced to the graph coloring problem.
 Find the most efficient mapping from registers to variables.

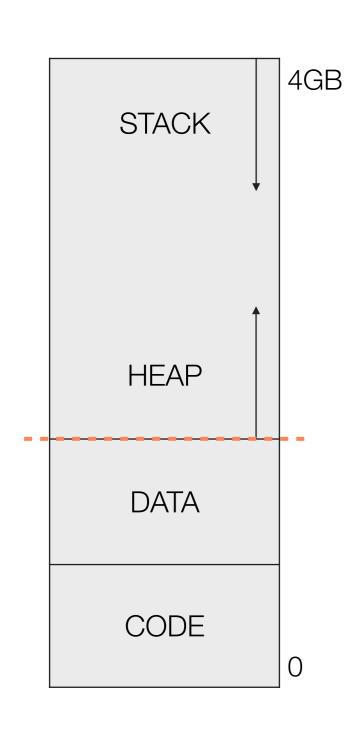
Prerequisites

Memory Model, Variables

Selfie Memory Model

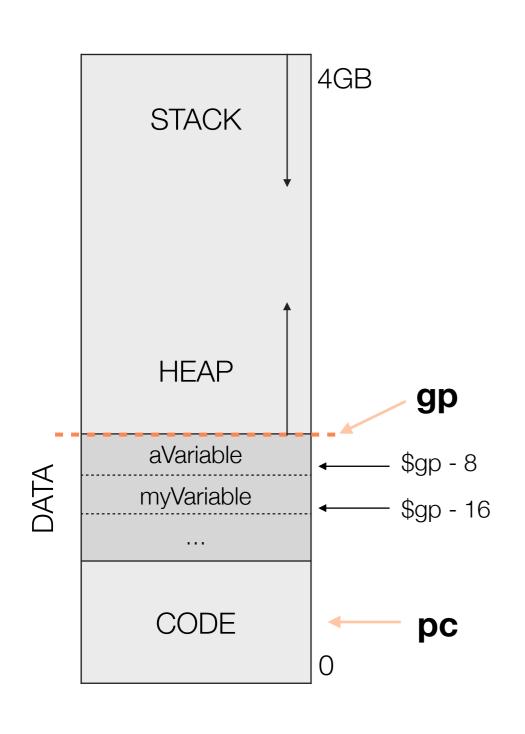


The Memory Layout



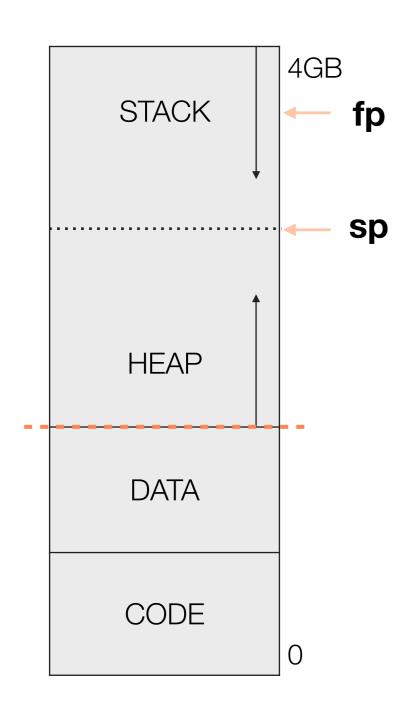
- The <u>stack</u> is used for procedure calls. It hosts local variables and procedure parameters.
- Memory allocated by malloc() is on the <u>heap</u>.
- ► The <u>data segment</u> is where global variables, strings and big integers are stored.
- Machine code instructions are stored in the code segment of our memory model.

The Pointers



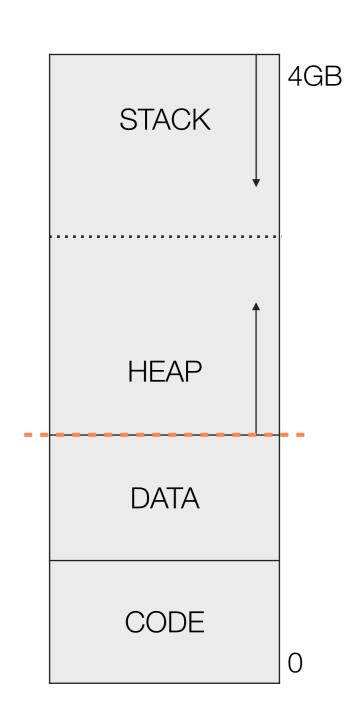
- ► Global variables, strings and big integers are accessed using the **global pointer** (**gp**) and the negative offset into the data segment that is stored in the symbol table. The address of the first memory cell is \$gp 8.
- The program counter (pc) is set to the next instruction that will be loaded into the instruction register.

The Pointers



- The **stack pointer** always points to the top of the stack and is used for memory allocation and deallocation on the stack. In our model the stack grows downwards from high memory addresses to lower ones.
- The frame pointer provides a convenient way to access local variables and procedure parameters. This will be explained in more detail when <u>procedures</u> are discussed.

Static and Dynamic Memory



- The program break splits the memory logically into two parts.
- Memory below the program break is static
 - the layout is fully determined at compile time and is fixed during execution.
- Memory above the program break is dynamic
 - the layout changes during execution.

Memory Allocation

Different concepts depending on when and where memory is allocated.

The When

- static memory allocation compile-time concept.
- dynamic memory allocation runtime concept.

The Where

- using malloc() on the heap.
- procedure calls in-order on the stack
 simplest way to dynamically allocate memory.
- at compile time by the compiler in the data segment.

Variables

Variables are **mappings** from names (identifiers) to memory locations (memory addresses) that hold a information (value).

- a variable is defined within a <u>scope</u> that defines its visibility and when/ where space is allocated.
- in <u>statically typed languages</u> a variable has a type that gives information about the value associated with that variable (size of the type,...) and how the programmer intends to use the variable.
- in selfie the only types are uint64_t and uint64_t*
 (both 8 byte = REGISTERSIZE).

```
uint64 t aVariable;
                           Declaration
uint64 t myVariable;
uint64 t main() {
  myVariable = 2;
                           Access
  aVariable = myVariable
```

Local Variables

```
uint64_t func() {
    uint64_t myLocal;
}
```

Declaration

- declared inside a procedure local scope.
- visible (accessible) within the procedure it is declared in.

Allocation

space is allocated upon each procedure call at runtime on the stack.

Global Variables

Declaration

- declared outside any procedure global scope.
- visible (accessible) throughout program execution.

Allocation

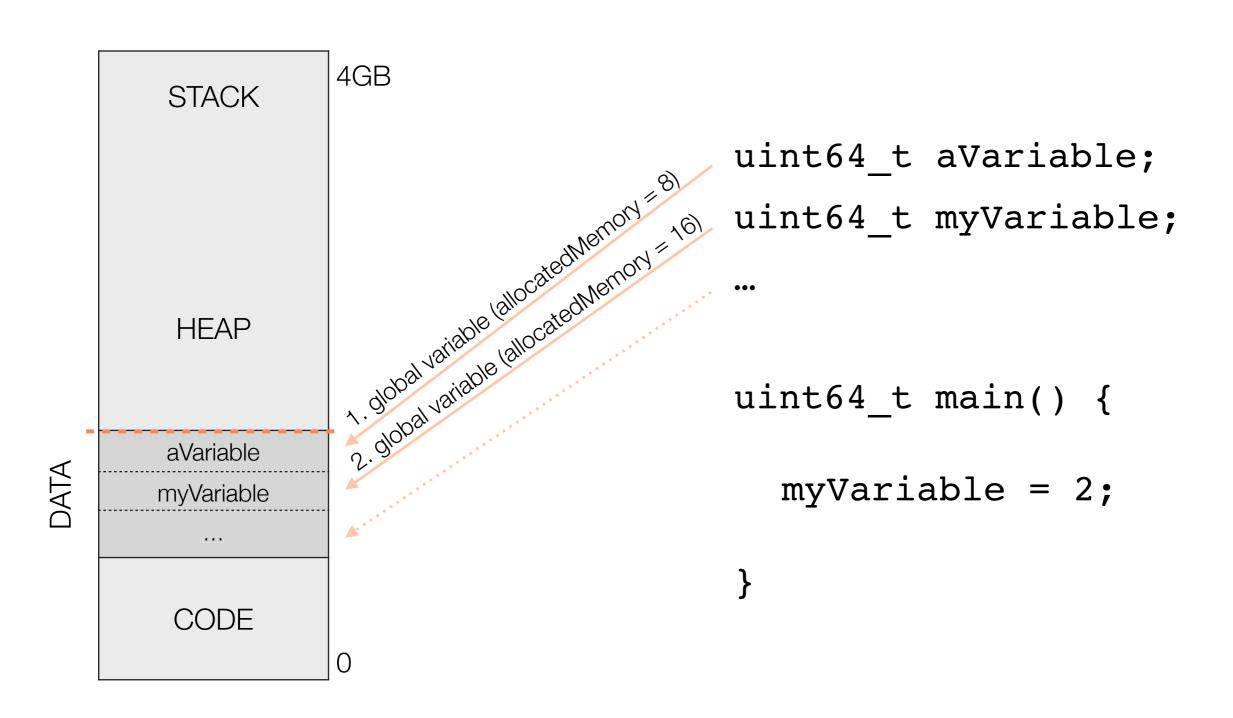
- done once at compile time in the data segment.
- space for variables is allocated in the order in which they appear.
- allocated_memory holds the size (in byte) of the currently allocated memory in the data segment.

- Parsing a declaration in selfie
 - allocate space for 1 value of given type
 global: increase allocated_memory by 1*REGISTERSIZE (variables in selfie are of same size)
 = reverse bump pointer.

local: decreasing the value of stack pointer (this is part of the procedure prologue which is discussed later).

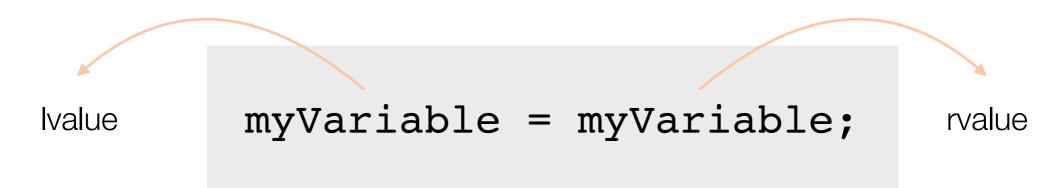
introduce mapping from name to memory location.

```
uint64 t aVariable;
                               Declaration
                               allocate memory and
                               introduce mapping
uint64 t myVariable;
                               (address know)
uint64 t main() {
  myVariable = 2;
                               Access
  aVariable = myVariable
```



- Parsing an access in selfie
 - to access a variable (i.e. to get the value associated with it) the memory address of that variable is needed.
 - therefore we need a way to remember the previously introduced mapping from name to address.
 - a data structure called <u>symbol table</u> is used to remember these mappings.
- A variable that appears on the left-hand side of an assignment represents a storage location whereas a variable that appears on the right-hand side refers to the actual value associated with it.

Ivalue and rvalue



Ivalue

- left-hand side of assignment.
- a storage location (address).
- when parsing the address is loaded into a register.

rvalue

- right-hand side of assignment.
- value in traditional sense.
- when parsing the address is loaded and dereferenced to retrieve the actual value from memory.

```
Declaration
uint64 t aVariable;
                                  allocate memory and
                                  introduce mapping
uint64 t myVariable;
                                  (address know)
uint64 t main() {
  myVariable = 2;
                                  Access
                                  address not known
  aVariable = myVariable
                                  therefore we need to
                                  remember the mapping in
                                  the symbol table
```

Symbol Table

- symbol table
- Provides 2 main functionalities
 - mapping
 associates a name with an entities so we can refer to it by its name.
 - scope
 enables us to have names with different scope.
- data structures used in selfie hash table and singly linked list.
- stored on the heap (using malloc()).

Symbol Table

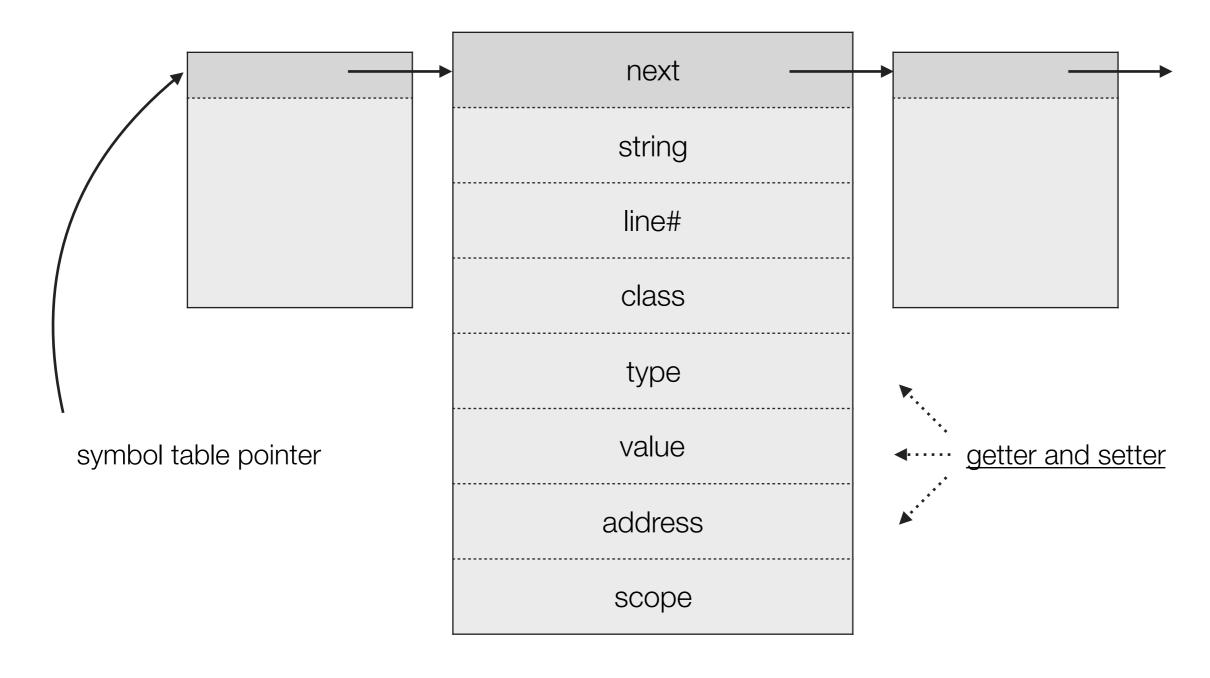
Why on the heap?

- the symbol table does not conform to the semantics of the stack
 - lifetime is not stack like no in-order (de)allocation
- a symbol table of dynamic size cannot be stored in the data segment (static size)
 - a feature of our symbol table is unbound size (idea)
 (a fixed-size symbol table could be stored in the data segment)
 - the only other choice is to store the symbol table on the heap

Symbol Table Entry (as a list element)



symbol table in selfie

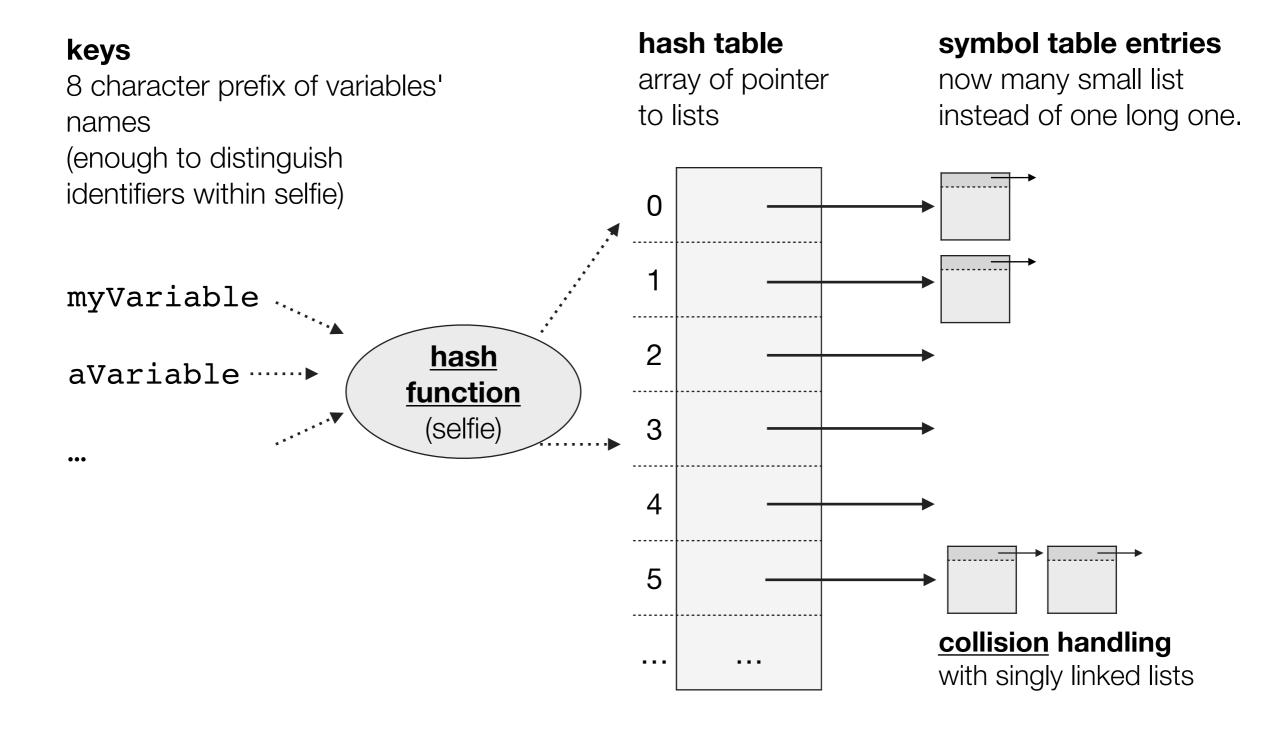


Hash Table

- ► a <u>hash table</u> is a data structure similar to an **array...**
 - it can be seen as array of buckets, each holding data
 - each bucket is uniquely identified by an index
- but it uses a different access method
 - each data value is associated with a key value
 - a technique called <u>hashing</u> is applied to convert a key value into an index using a <u>hash function</u>
- searching values in a hash table (computing the index) can be done in constant time (iff collision free) whereas the search complexity for arrays is linear in the number of entries

Hash Table





Symbol Table API

- insert and search
- getter/setter for adding an element

singly linked list:

- insert complexity constant time
- search complexity linear in the number of entries
- BUT below 1000 list elements linear is OK

hash table

- insert complexity constant time
- search complexity constant time iff collision free (worst-case still linear in the number of entries)

Symbol Table Scope



- names can have different scope (visibility) and therefore they do not need to be unique and can be <u>shadowed</u>
- In selfie we simply use 3 symbol tables get scoped symbol table entry
 - local symbol table (singly linked list)
 first search local table (local variables shadow global variables
 - library symbol table (singly linked list)
 library procedures override defined/declared procedures in global
 symbol table. This table is used for bootstrapping syscalls using
 wrapper code.
 - global symbol table (hash map)
 search last if entry not already found
- Due to its enormous size the global symbol table is implemented as a hash table (speed up for make target of 30%)

Back to Variables

```
uint64_t aVariable;
uint64_t myVariable;
uint64_t main() {
   myVariable = 2;
}
```

,	
next	next
string	myVariable
line#	2
class	VARIABLE
type	uint64_t
value	0
address	-16
scope	GLOBAL_TABLE

- remember negated offset = the value of allocated_memory
- enough to compute the address from global pointer and offset

Variables Conclusion

- The examples so far show variables that hold a single value (uint64_t my_variable, uint64_t* my_pointer)
- A variable can also refer to multiple values
 - array
 - struct

Data Structures and Data Types

Arrays and Structs

Array and Struct

- Selfie does not support arrays or structs.
- However, this section should give you a good understanding of arrays and structs and the modifications necessary to implement support for global and local arrays and structs in selfie.
- We will therefore look at basic concepts, but also discuss details concerning the implementation of global arrays and structs in selfie.
- This should help you understand how data is stored in memory.

Data Structure

A <u>data structure</u> is a way to **organize a collection of data** in memory that simplifies access and modification of that data.

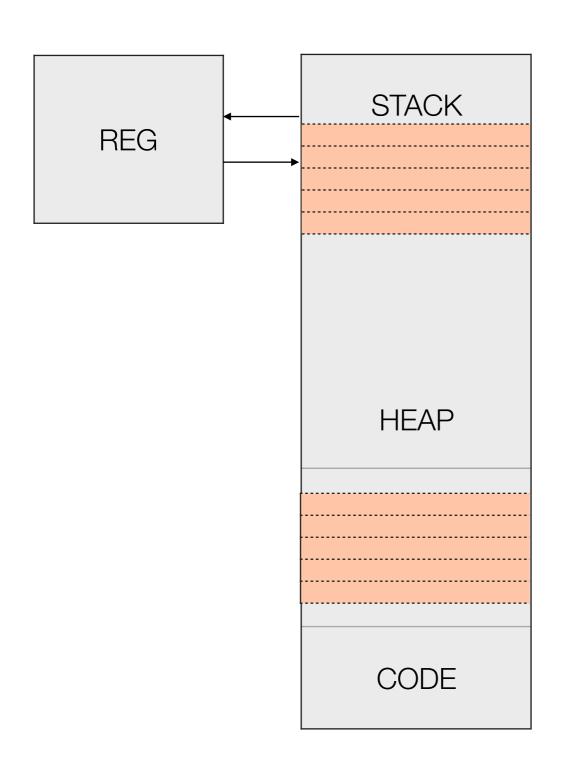
- defines relationship between data values
- defines functions/operations on the data structure
- some data structures: <u>static arrays</u>, <u>dynamic arrays</u>, <u>linked lists</u>, <u>stack</u>, ...
- we will only consider static arrays (referred to as array)

Array

An **array** is a **variable** that is associated with a fixed number of values of the same data type. An element of the array is selected using an index.

- container for multiple values of the same data type
- the length is fixed when the array is created
- the index to access an element an array is computed at runtime
- arrays can be implemented by using contiguous space or a pointer structure

Array



- Arrays declared within procedures are **local** to that procedure.
 - → Space is allocated on the stack at runtime (no pointers to arrays)
- Arrays declared outside any procedure are global.
 - → Space is allocated once in the data segment at compile time
- From now on we will only consider global arrays implemented as contiguous space of addresses.

"Why do we want to support arrays in a programming language?"

"To answer this question we will take a look at the first two of three* metrics a we as computer scientists are generally interested in."

time

space

energy

^{*} there are others, but for us these are the most important

Time

► The big advantage is constant time access when the array is implemented as contiguous space of addresses.

► The access is constant time because

... we assume that the operations to compute the address can be done in constant time.

Space

- ► The downside is the potential for <u>fragmentation</u> when the array is implemented as contiguous space of addresses.
- ► The problem with fragmented memory...
 - We request space of size n.
 - Take n = 6 for the example below.
 - There is enough free space (> n), but there is not enough
 contiguous free space to accommodate a block of size n, so the
 request fails.



U

Space

- Conditions that create fragmentation are...
 - allow allocating space of different size together with deallocating out of order
 - → no fragmentation on stack
- Using lists would solve the problem of fragmentation, since the nodes have the same size, but
 - we would have drawback in size (pointer structure)
 - and we would loose constant time access.

Array

```
uint64_t myVariable; Declaration
uint64_t myArray[5];
uint64_t main() {
  myArray[1] = 2;
}
Access
```

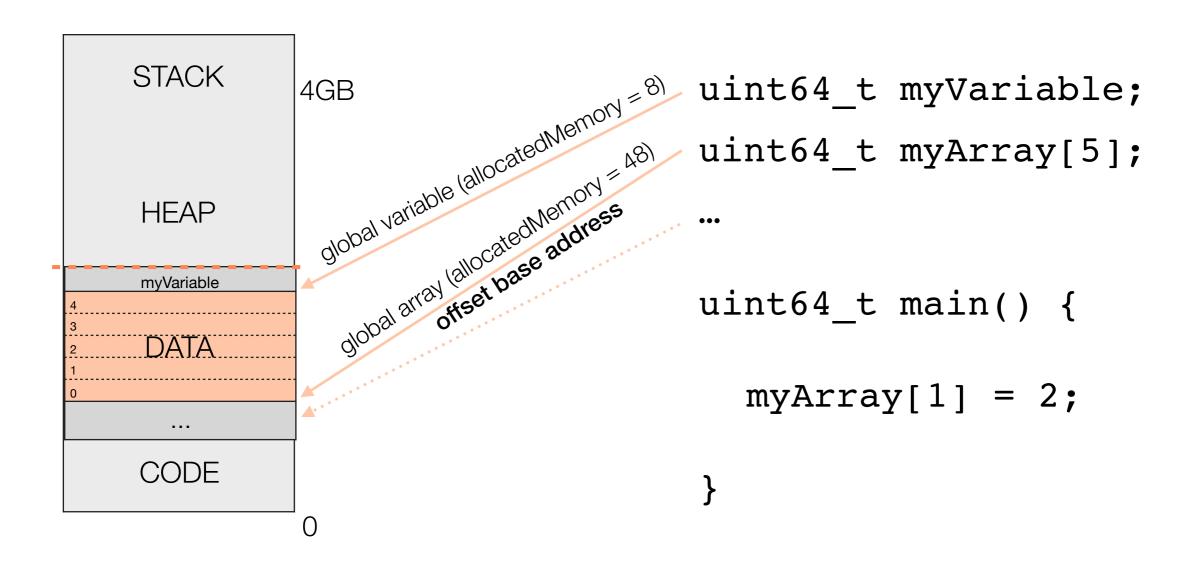
Array Declaration

```
base address
uint64_t myArray[5];
```

- declare a variable and allocate space
 - instead of allocating space for 1 value (as before) we allocate
 contiguous space for 5 values of given type
- calculating size of array myArray[2]...[4]
 - size must be given using constant values
 - formula: 2 * ... * 4 * size of element type

Array Declaration

Arrays are organized upside-down (conform to C standard)



Selfie Declaration

```
uint64_t myArray[5];
```

What is new?

- scanner does not know '[', ']'
 - → new characters
 - → new symbols
- parser does not recognize the selector syntax
 - → new parsing procedure
- symbol table can not handle variables of size other then REGISTERSIZE
 - → modify symbol table, store additional information
- no code generation for declarations

Array

```
uint64_t myVariable; compiler allocates
uint64_t myArray[10];

uint64_t main() {
  myArray[1] = 2;
}
Access
```

Declaration

Array Access

base address

myArray[1] = 2;

yourArray[x] = 3;

positive offset into the array

access a value within the array

- using integer constants or integer variables in between brackets
- address needs to be calculated from the base address and the offset into the array at runtime
- like dereferencing using the star operator
- calculating the address myArray[2]...[4]
 - the offset into the array has to be <u>calculated</u> at runtime
 - the code to do so is generated at compile time

Selfie Access

```
myArray[1] = 2;
yourArray[x] = 3;
```

What is new?

- scanner does not know '[', ']'
 - → new characters
 - → new symbols
- parser does not recognize the selector syntax
 - → new parsing procedure
- code generation to implement behavior according to specification
 - → calculate address
 - → load address (Ivalue) or value (rvalue)

Array

```
Declaration
uint64_t myVariable; compiler allocates
memory

uint64_t myArray[10];

uint64_t main() {

   myArray[1] = 2;

   Access
   calculate address at runtime
```

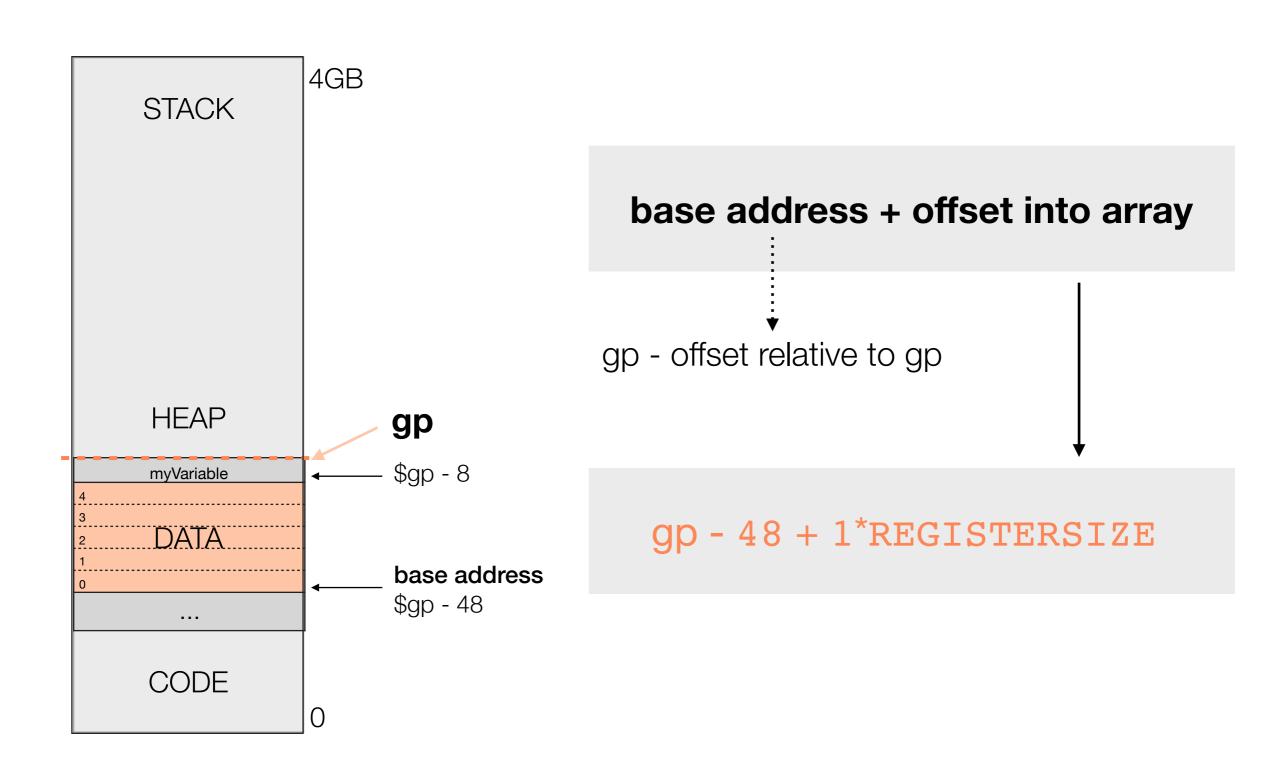
Array Access and Pointer Arithmetic

```
myArray[1] = 2;
*(myArray + 1) = 2;
```

equivalent semantic

- load address of myArray then load 1
- array semantics/pointer arithmetic add 1 * REGISTERSIZE
 → address
- store 2 at that address
- depending the implementation even the emitted code might be equivalent

Calculate Address



Calculate Address

- the offset into the array has to be calculated and added to the base address
- therefore we need to remember the structure of the array
 - number and size of dimensions
 - in the symbol table
- storing multidimensional arrays in one-dimensional memory we can either use <u>row-major or column-major order</u>

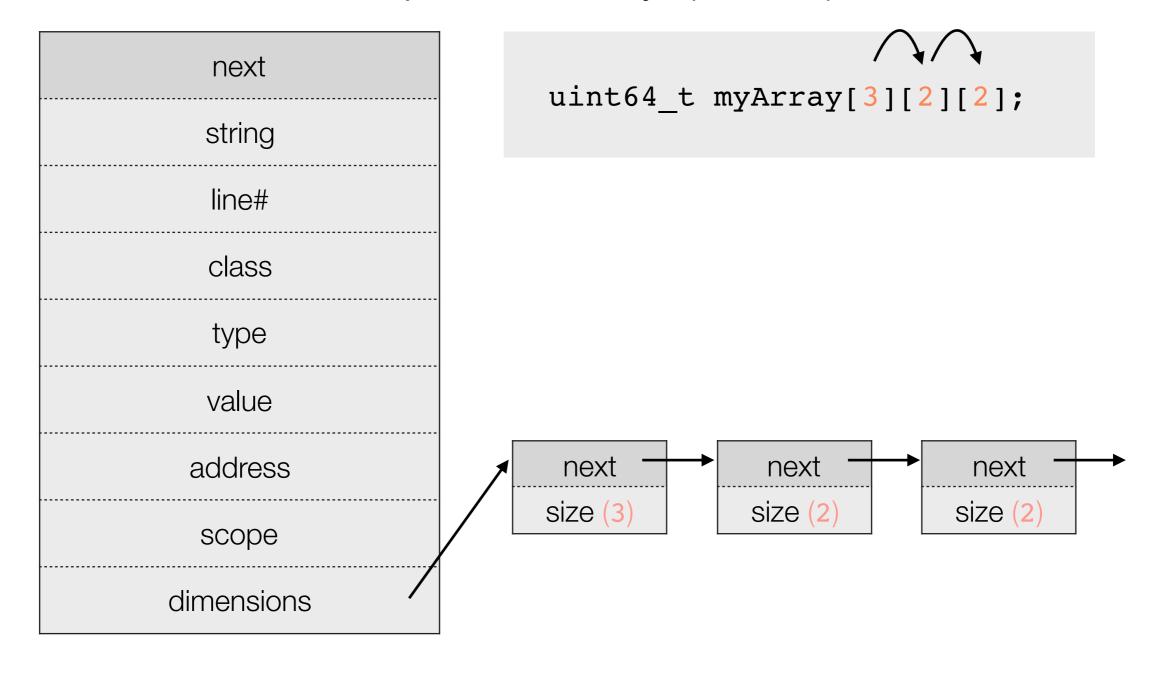
IMPORTANT:

address calculation is done at **runtime**, because variables might be used to access the array.

(access involving only constant values could be optimized)

Symbol Table Entry

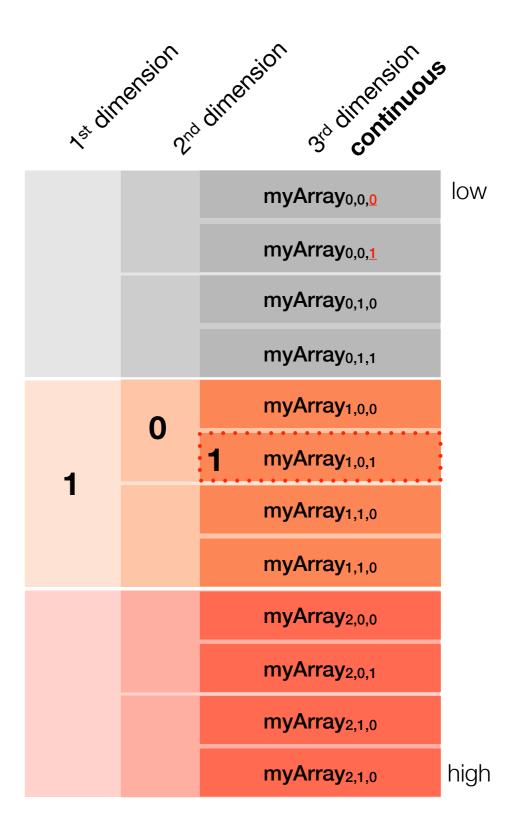
- store the information while parsing the selector
- this can be done in multiple different ways (order, ...)



Row-Major

row-major

Address Calculation



Column-Major

column-major

Address Calculation

low

high

 $myArray_{\underline{0},0,0}$

myArray_{1,0,0}

myArray_{2,0,0}

 $myArray_{0,1,0} \\$

myArray_{1,1,0}

myArray_{2,1,0}

myArray_{0,0,1}

myArray_{1,0,1}

myArray_{2,0,1}

myArray_{0,1,1}

myArray_{1,1,1}

myArray_{2,1,1}

1st dimension **continuous**

2nd dimension

3rd dimension

Implementation Tips

Divide and conquer

1. Syntax

modify compiler to support **parsing** declaration and access without generating code

- new parsing procedure(s) for selector syntax '[..]'
- test on examples
- Information (could be part of 1.)
 modify symbol table and parser and store the necessary information
 - store dimensions and size

3. Semantics

modify compiler to generate code for access computation

- use information from symbol table to emit instructions in the right places to enable address calculation at runtime
- test on examples

Implementation Tips

keep in mind the different semantics of the selector for

declaration

- no code generation
- only symbol table entry
- calculate size of array

access

- code generation
- calculate address

Compile Time and Runtime

- ► declaration is a compile time concept
 - create symbol table entry and allocate memory
- ► access is a runtime concept
 - the address calculation can only happen at runtime (variables)
 - at runtime there is code that computes the address
 - this code is emitted at compile time using the selector currently parsed and symbol table information

Array Conclusion



- Looking at the implementation of the <u>power_of_two_table</u>, <u>SYMBOLS</u>,... now, you should realize and understand
 - practically a one-dimensional arrays
 - implemented by a pointer to a block of addresses (malloced) and accessed using pointer arithmetic
- Using arrays instead does not change the semantics but the syntax.

Array Conclusion

Advantages

 With arrays we have a convenient way to allocate space for multiple values and access them using indexes (no external pointer arithmetic)

Disadvantages

- An array cannot hold values of different type without type casting
- Access to an array is **not "safe"**. Without checking bound, illegal access is possible. (access outside arrays)

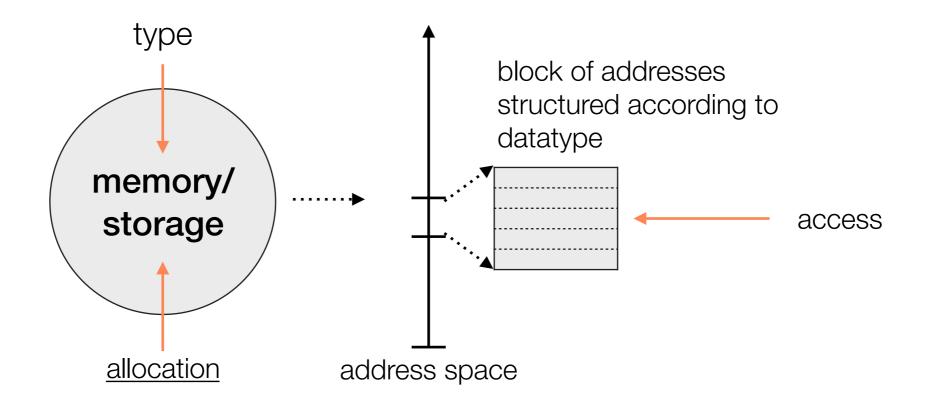
Data Type

A <u>data type</u> defines the type and structure of a data item and tells the compiler (at **compile time**) how the programmer intends to use.

- defines operations on a variable
- defines the domain and size of a variable
- primitive data types: uint64_t, uint64_t*
- composite data types: struct

compile time information

Data Type



- The data type of a variable defines the structure of data in memory and tells the compiler how the programmer intends to use that variable.
- Different data types use different access strategies that have different properties.

Struct also Record, Object in Java

A **struct type** is a user defined composite **data type** that allows to combine data elements (variables) of different types under one name/identifier.

A **struct** is an instance of such a struct type. It is a ordered collection of variables called fields or members that store data.

- a field is identified by a name/identifier
- a field can be of any data type composite or primitive
- a struct instance is implemented as contiguous block of memory

We could ask the same Question again:

"Why do we want to support structs in a programming language?"

"To answer would be the same as it was for arrays ><u>time and space</u><

But here another important concept comes into play >compiler awareness*<"

^{*} no need to understand this now

Using Structs

1. declare struct type (optional)

2. <u>define struct type</u>

3. <u>declare instance of struct type</u> we only discuss **pointers to structs**

4. <u>allocate space for struct</u>

5. access struct

```
struct record_t;
```

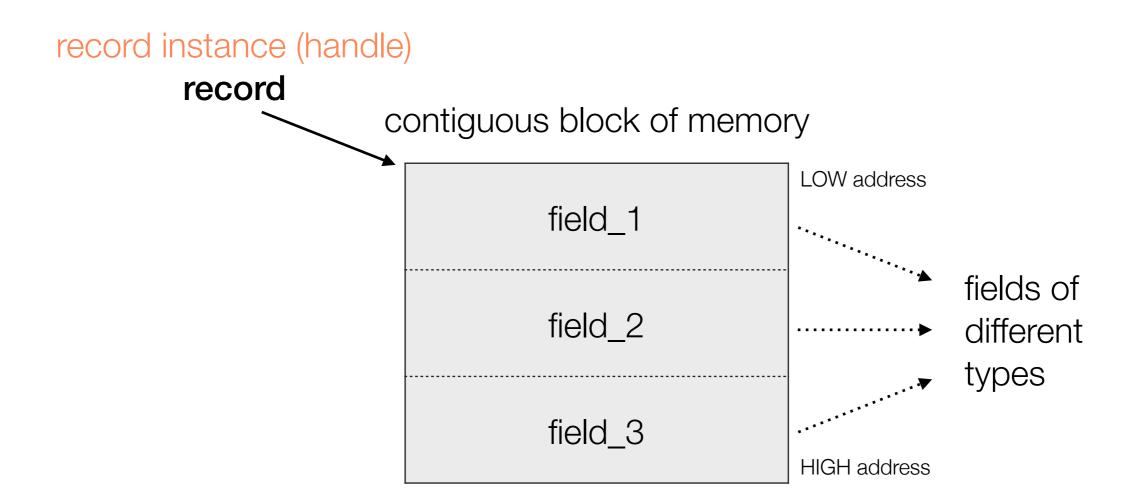
```
struct record_t {
    struct myStruct* field1;
    uint64_t field2;
};
```

```
struct record_t* record;
```

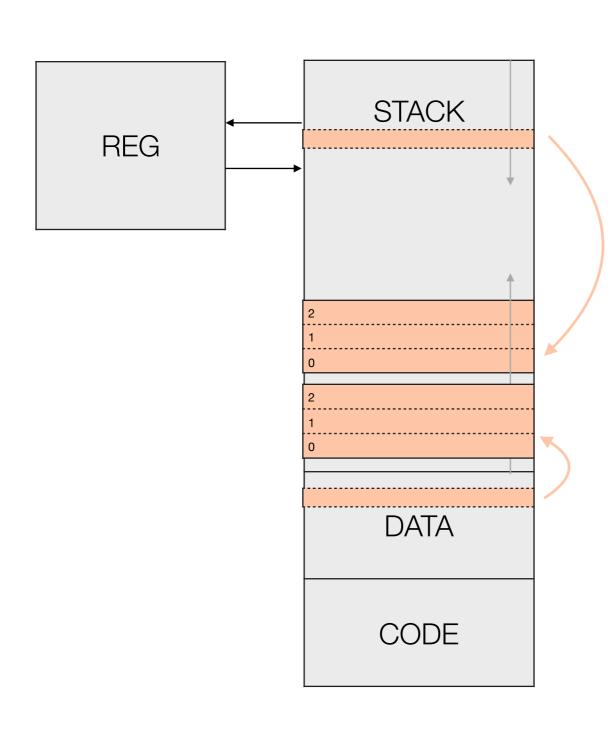
```
record = malloc(16);
```

Struct Instance (abstract)

we only discuss pointers to structs



Struct



- Pointers to structs can be declared locally and globally.
- The pointers will be put on the stack (local) or in the data segment (global).
- Before the struct can be accessed memory has to be allocated (malloc) by the programmer on the heap.
 - → Every struct is put on the heap and is organized up-side down

Struct Definition Type

```
struct record_t {

struct myStruct* field1;

uint64_t field2;

};

new type name

type and name of fields
```

- Defining a new struct type
 - name/identifier for new type
 - name/identifier and type of each field
 - no code generation
- To declare a variable of this type later, we need to store and manage this type information

Struct Definition Type Information

```
struct record_t {

struct myStruct* field1;

uint64_t field2;
};

new type name

type and name of fields
```

- ▶ we need a data structure that is globally visible to store type information → use symbol table to store:
 - name/identifier of new type
 - associate new type with the types structure
- type structure is necessary to access fields of the struct
 - field name/identifier
 - field type
 - · field offset within the struct

Offset

- The offset of a field is the accumulated size of the fields that come before it
 - offset of first field is always 0
 - here offset of second field is size of field1

```
struct record_t {
offset 0
    struct myStruct* field1;
offset 8
    uint64_t field2;
};
```

Struct Definition Symbol Table Entry



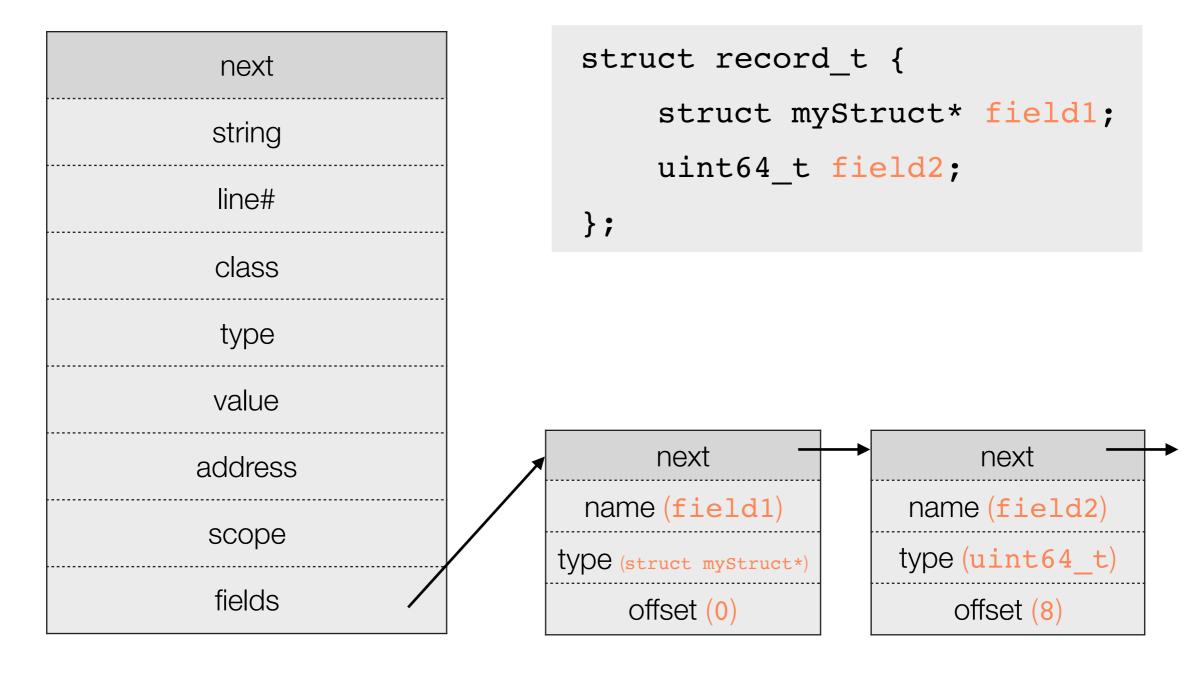
- Until now we only stored "simple" variables and array variables in the symbol table.
- Storing type information requires us to be able to distinguish between variables and types
 - this can achieved by labeling the entry
 (it actually is already done using the class field of the entry)

Classes

- besides variable information the symbol table is used to store bigints, strings and procedure information (<u>classes</u>: VARIABLE, BIGINT, STRING, PROCEDURE)
- we can introduce a new class TYPE

Struct Definition Symbol Table Entry

- store the information while parsing the definition
- one possibility to do so



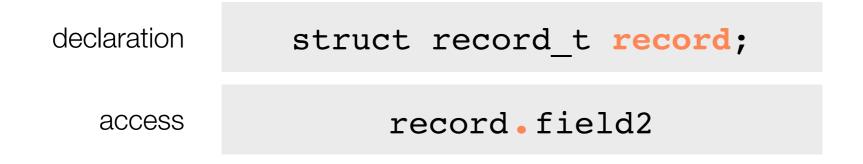
Selfie Definition

```
struct record_t {
    struct myStruct* field1;
    uint64_t field2;
};
```

What is new?

- scanner does not know "struct" keyword
 - → new keyword
- parser does not recognize this syntax
 - → new parsing procedure for struct definition
- symbol table can not store type information
 - → modify symbol table, store type information
- ► no code generation for definition

Struct Declaration in C

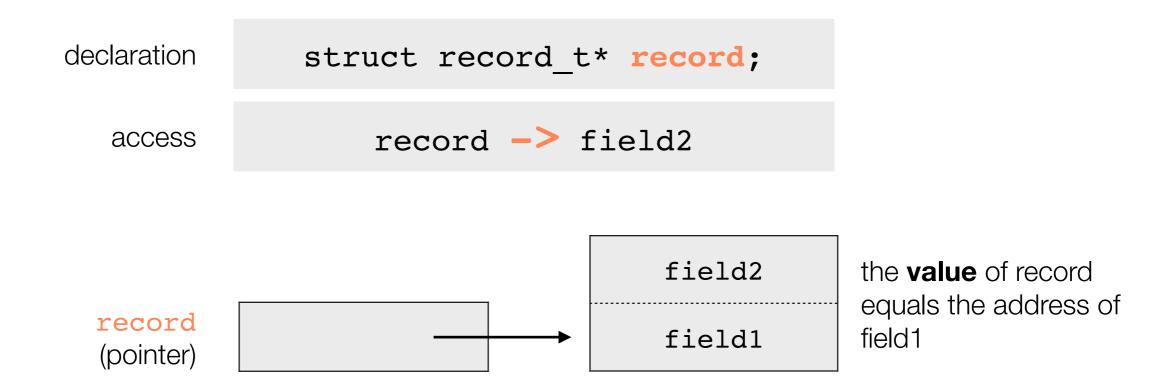


field2
record field1

address of record equals address of field1

- Declaring a variable of type struct record_t
 - the variable is this type
 - the size of this variable is the accumulated size of its fields
 - accessed using the "." notation

Struct Declaration in C



- ► Declaring a variable of type struct record t*
 - variable is a pointer to an instance of type struct record_t
 - the size of this variable (a pointer) is one word (REGISTERSIZE)
 - accessed using '->' notation
- We only consider pointers to structs

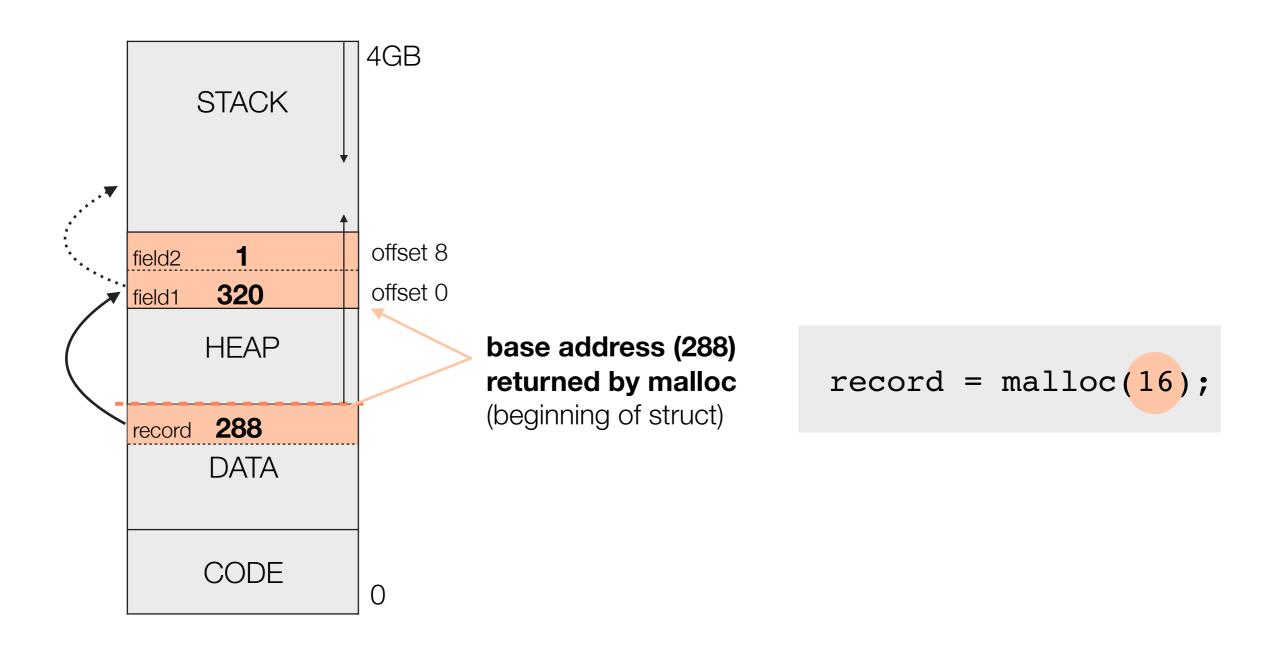


```
struct record_t* record;
```

What is new?

- scanner does not know 'struct' keyword
 - → new keyword
- parser does not recognize the struct types
 - → parse new type
- symbol table can not define variables of type "struct type" (struct types are represented by symbol table entries)
 - → modify the symbol tables type field so it can refer to "struct type"
- no code generation for declarations

Allocate Space for Struct



- allocate space for a struct on the heap
- size of the struct is the accumulated size of its fields

Struct Access



access a value within the struct

- use field name/identifier
- address is calculated from the base address and the offset of the field inside the struct at runtime

calculating the address

- the offset of a field from the beginning of the struct is known at compile time
- the base address of the struct returned by malloc is not known before runtime
- code is generated to add the offset to the base address at compile time

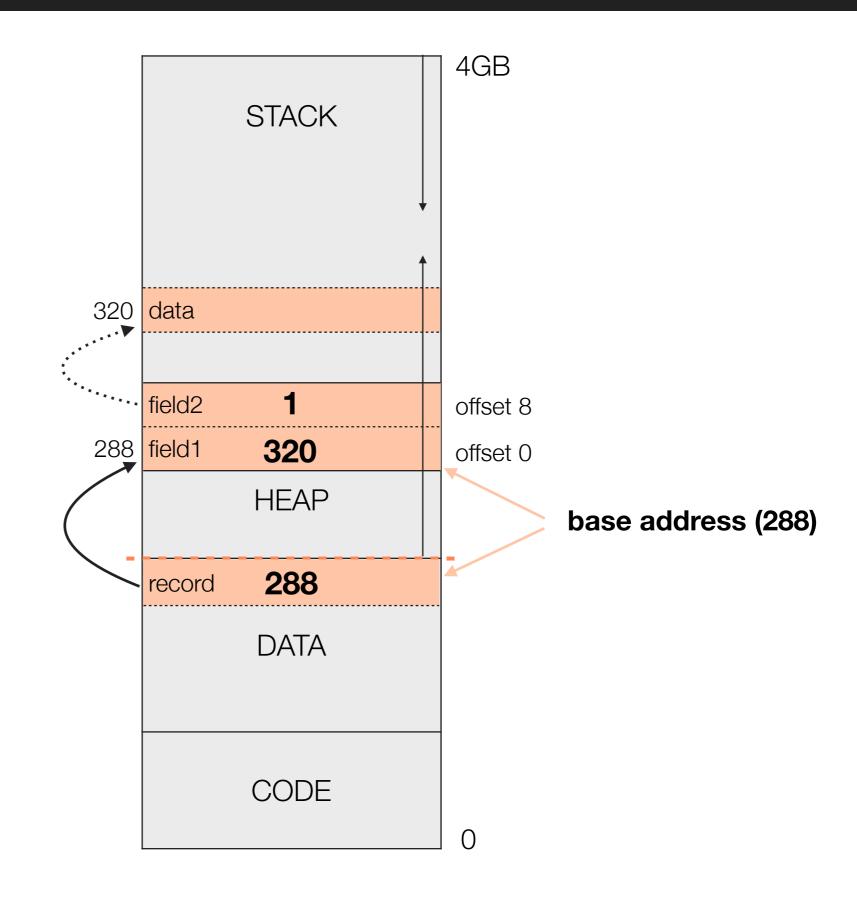
Selfie Access

record -> field2;

What is new?

- scanner does not know '->'
 - → new symbol
- parser does not recognize the selector syntax
 - → new parsing procedure
- code generation to implement behavior according to specification
 - → calculate address
 - → load address (Ivalue) or value (rvalue)

Struct Access Example



Structs Code Generation

1. 2. record -> field2

1. load base address of struct (record)

$$1d $t0 -8($gp)$$

2. add offset of field (field2) to the beginning of the struct

- → now \$t0 holds address (not value) of field(field2)
- 3. optional: dereference to get value

ld \$t0 0(\$t0)

Structs within Structs Code Generation

1. **load base address** of struct (record)

$$1d $t0 -8($gp)$$

- 2. **add offset** of field (field2) to the beginning of the struct addi \$t0 \$t0 8
 - → now \$t0 holds address (not value) of field(field2)
- 3. **dereference** to get the base address of myStruct (320)

4. add offset of field (data) to the beginning of the struct

- → now \$t0 holds address (not value) of field (data)
- 5. optional: dereference to get value of data

Assumption:

Before we see an arrow the register contains the address (not the value) of the field.

see arrow

→ dereference

and repeat

see arrow → dereference

and repeat

. . .

Implementation Tips

Divide and conquer

1. Syntax

modify compiler to support **parsing** definition, declaration and access without generating code

- 'struct' keyword
- new parsing procedure for struct definition and declaration
- new parsing procedures for selector syntax '-> '
- test on examples
- 2. **Information** (could be part of 1.) modify symbol table and parser and store the necessary **information**
 - new class for TYPE
 - store type information
 - → field names, types and offset relative to beginning of struct
 - store variables with a struct type

Implementation Tips

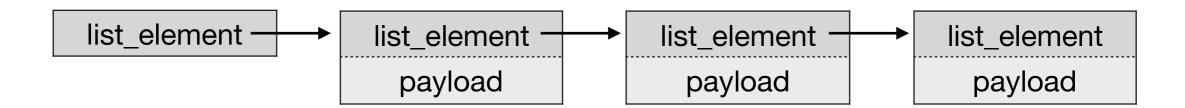
3. **Semantics**

modify compiler to generate code for access computation

- use information from symbol table to emit instructions in the right places to enable address calculation at runtime
- test on examples

Potential for Recursion

- with structs we can define <u>recursive data types</u>
- recursive data types are used to build <u>recursive data structures</u> that are dynamic in size - lists, trees,...



Potential for Recursion

```
struct list_element_t {
    list_element_t* list_element;
    uint64_t payload;
};
```

- forward declaration of data type unavoidable
- It is important that a pointer is used in the above case
 - computing the size of this type would not be possible otherwise
 → dynamic size

Array vs Struct

- arrays and structs share some characteristics
 - contiguous space
 - constant time access
 - automated address calculation
 - size fixed at compile time (as we discussed)
 - the concept of local and global structs

Array vs Struct

 We will now look again at the key differences between arrays and structs

Data

- struct: data elements of different type allowed
- array: data elements have to be of same type

Access

- <u>struct</u>: accessed using field name/identifier
- <u>array</u>: accessed using **index values** (integers) using pointer arithmetic

Array vs Struct

Address computation

- struct: offset to beginning of struct always known at compile time
- <u>array</u>: necessary information might not be available before runtime

myArray[i]

information only available at runtime

record -> field1

information always available at compile time → offset is always known

Array Access

- The key disadvantage of array access is the use of pointer arithmetic
 - illegal access outside the is easily possible
 - accidentally use wrong index
 - compiler does not understand the intentions of the programmer and therefore cannot prevent the above problems

Struct Access

- The key advantage of struct access is the use of field names/ identifiers
 - only fields of that array can be accessed
 (no illegal access possible when using correct syntax)
 - using the wrong name is not possible
 - compiler does understand the intentions of the programmer and therefore has the capability of giving feedback on correctness

Structs Conclusion

- By Implementing/supporting structs we teach the compiler.
- We make the compiler become aware of the programmers intention and give it the capability to give feedback on correctness.
- With structs we now have an unbound type space.

Memory Allocation

Procedures, Garbage Collection

Procedures



Procedure Declaration

doMagic(7, 1)

```
uint64 t doMagic(uint64 t power, uint64 t wand);
                               formal parameters
Procedure Definition
uint64_t doMagic(uint64_t power, uint64_t wand) {
    uint64 t trick;
                                           no code generation - just symbol table
                                                             code generation
    trick = power * wand;
    return trick;
Procedure Call
```

actual parameters

Formal/Actual Parameters & Local Variables

- Local variables represent types and names. They are a special case of formal parameters.
- Formal parameter represent types and names. They are more general and more expressive than local variables, because they enable us to access and modify information outside the procedure.
- Actual parameters represent the actual type of the parameters when the procedure is called.

Procedure Declaration

```
uint64_t doMagic(uint64_t power, uint64_t wand);
```

- symbol table entry
 - introduce a name/identifier for procedure (global)
 - introduce name/identifier for formal parameters (local)
- no code generation

Procedure Definition

```
uint64_t doMagic(uint64_t power, uint64_t wand) {
    uint64_t trick;
    ... body ...
}
```

symbol table entry

- introduce a name/identifier for procedure (global)
- introduce formal parameters (local)
- introduce local variables (local)

code generation

- generate code to prepare the machine for the execution of the procedure (prologue)
- for body of the procedure (its implementation)

Procedure Call

```
doMagic(7, 1);
```

The calling procedure is also known as caller, whereas the called procedure we consider to be the callee.

code generation

- generate code to prepare the machine for jump (parameters and saving certain registers)
- for actual jump to the code implementing the procedure

Parameters

- The caller hands the parameters over to the procedure
- There are two strategies how this can be done:
 - using special registers
 - → fewer parameters possible
 - placing them somewhere the callee can access them
 - → more general
 - → more parameters possible

This is a <u>convention</u> that must be agreed on by the caller and the callee.

In selfie the later one is implemented - pushing parameters on the stack

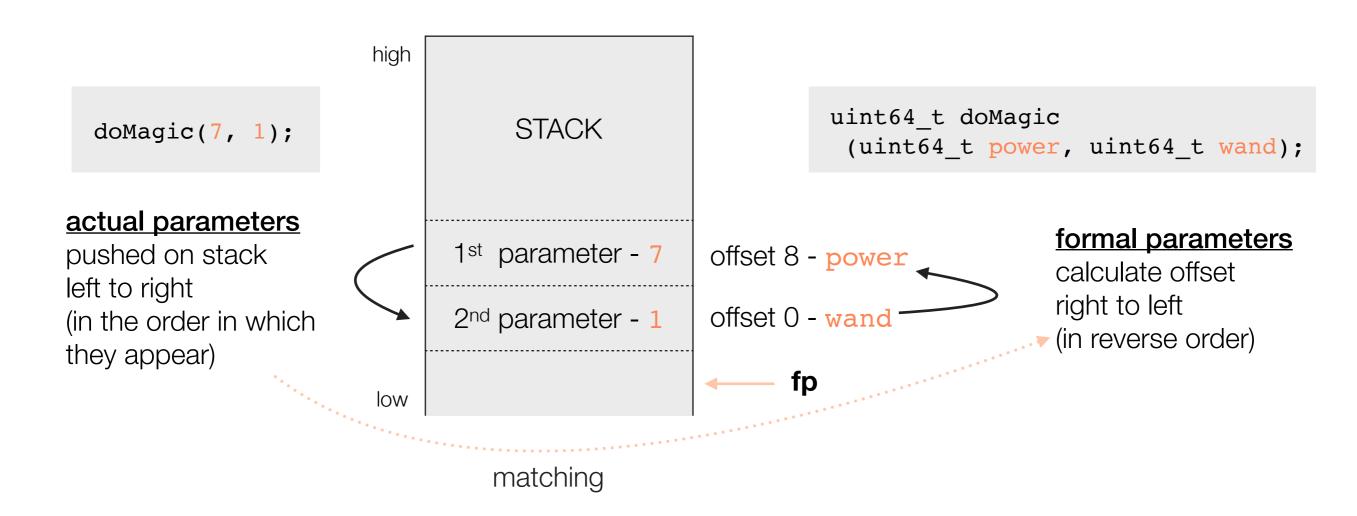
Parsing Parameters

- Matching between actual parameters and formal parameters based on the order in which they appear
- Parsing a declaration
 - formal parameters parsed from left to right
 - offset calculated from right to left
- Parsing a call
 - actual parameters parsed from left to right
 - pushed on stack from left to right



Matching Parameters





- Why not parse declaration the other way around?
 - tedious to match for every call (pushing parameters)
 - Trade-off: either calls are easy OR declaration is easy

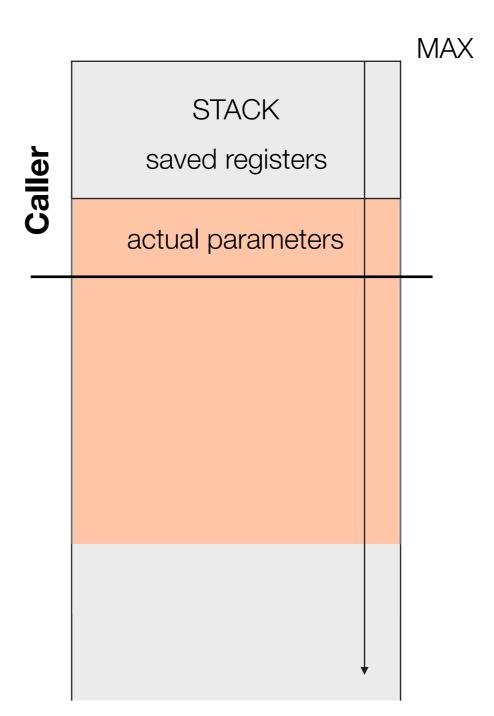
Caller vs Callee Save and Restore

changed.



- Saving and restoring registers in case of a procedure call can be done by either the caller or the callee. There are <u>conventions</u> about how to do this. More important are the **assumptions** that can be made in both cases.
 - Caller-saved registers are saved by the caller right before a
 procedure call and restored immediately after. It is only necessary to
 save the registers that contain values that are used by the caller after
 the call.
 - Advantage callee can assume that it can use all registers.
 - <u>Callee-saved registers</u> are saved by the callee right before they are
 modified and get restored before returning to the caller. Those
 registers can be used to pass information to the callee.
 Advantage <u>caller</u> can assume that none of the register values have

Procedure Frame

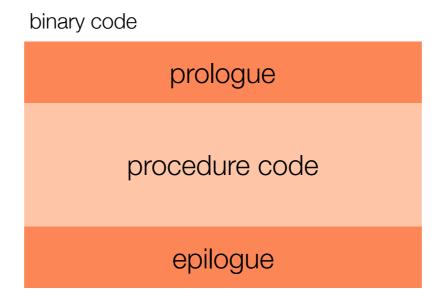


part of the stack accessible to the procedure

Prologue & Epilogue



- Every procedure is encapsulated by
 - a prologue that prepares the stack and registers for the execution of the procedure
 - an epilogue that prepares the stack and registers to return after the execution of the procedure.
- They are compiled as part of the procedure definition
 (compile procedure()) and hence they are part of the callee.



Prologue

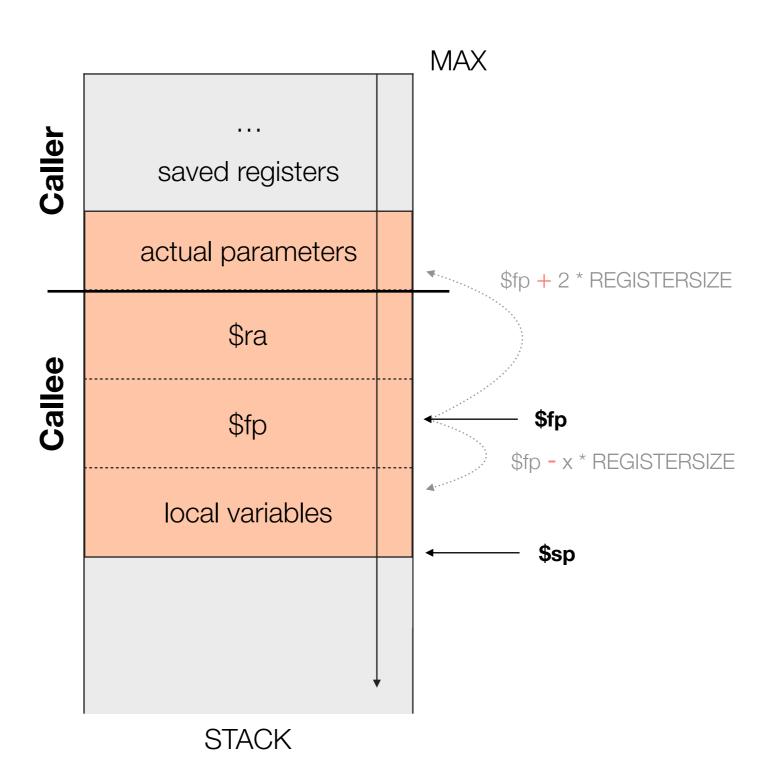
- The prologue of a procedure prepares the stack and registers for the execution of the procedure.
 - setting the frame for the callee
 - memory allocation for local variables
 - saving callers frame pointer and return address.
- The callers frame pointer needs to be saved because it will be overwritten when the callees frame is set.
- The \$ra register contains the address of the next instruction to be executed after the procedure. This register will be overwritten iff another procedure is called within the current procedure. Therefore it needs to be saved too.

Prologue

```
void help procedure prologue(uint64 t bytesOfLocalVariables) {
  // allocate memory for return address
  emitADDI(REG SP, REG SP, -REGISTERSIZE);
  // save return address
 emitSD(REG SP, 0, REG RA);
 // allocate memory for caller's frame pointer
  emitADDI(REG SP, REG SP, -REGISTERSIZE);
  // save caller's frame pointer
  emitSD(REG SP, 0, REG FP);
  // set callee's frame pointer
  emitADDI(REG FP, REG_SP, 0);
 // allocate memory for callee's local variables
  if (number of local variable bytes > 0) {
```

Procedure Frame after Prologue

- part of the stack allocated/ accessible to the procedure
- access with
 - frame pointer
 (parameters, local variables)
 - stack pointer (top of stack)



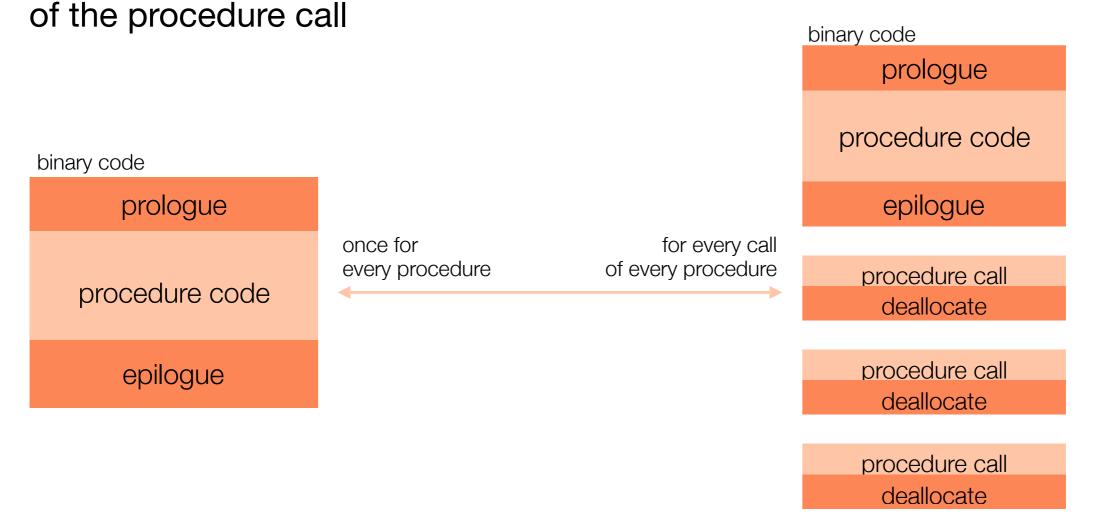
Epilogue

- The epilogue of a procedure prepares the stack and registers for the return of the procedure.
 - setting the frame (back) for the caller
 - deallocate memory of local variables
 - restore callers frame pointer and return address.
 - deallocate memory for parameters
- To deallocate memory for local variables the stack pointer is reset to the beginning of the frame.
- Notice: The memory for parameters is deallocated by the callee even though it was allocated by the caller. Why?

Epilogue

 Deallocating memory for parameters in the epilogue (callee): code for doing so is emitted just once as part of the procedure definition

Deallocating memory for parameters by the caller:
 code for doing so is emitted each time the procedure is called as part

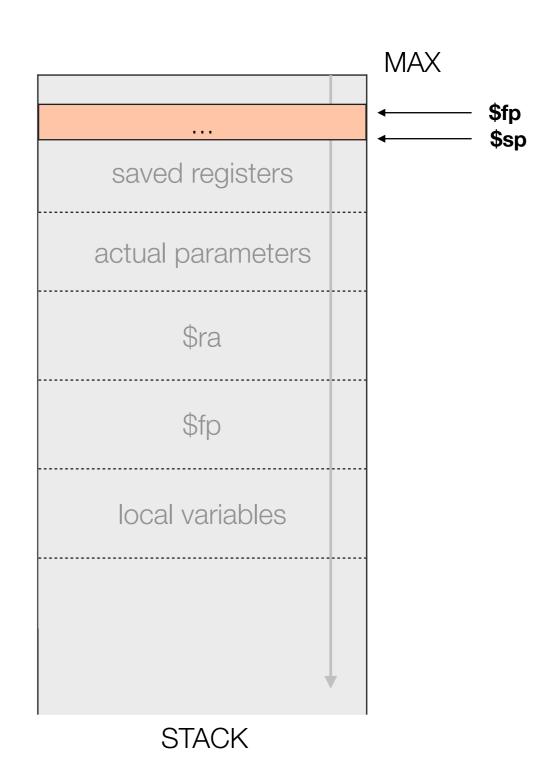


Epilogue

```
void help procedure epilogue(uint64 t parameters) {
  // deallocate memory for callee's frame pointer and local
    variables
  emitADDI(REG SP, REG FP, 0);
  // restore caller's frame pointer
  emitLD(REG FP, REG SP, 0);
  // deallocate memory for caller's frame pointer
  emitADDI(REG SP, REG SP, REGISTERSIZE);
  // restore return address
  emitLD(REG RA, REG SP, 0);
  // deallocate memory for return address and parameters
  emitADDI(REG SP, REG SP, REGISTERSIZE + parameters *
 REGISTERSIZE);
 // return
  emitJALR(REG ZR, REG RA, 0);
```

Procedure Frame after Epilogue

- part of the stack allocated/ accessible to the procedure
- callers frame restored and memory deallocated by resetting \$sp



Iteration

```
uint64_t add(uint64_t x, uint64_t y) {
    while (y > 0) {
        x = x + 1;
        y = y - 1;
    }
    return x;
}
```

- This procedure defines <u>iterative</u> addition of x + y
- The same procedure can be defined recursively.

Recursion

```
uint64_t add(uint64_t x, uint64_t y) {
    uint64_t sum;
    if (y > 0)
        sum = add(x, y - 1) + 1;
        recursive call
    else

base case
    sum = x;
    return sum;
}
```

- The definition of a <u>recursive procedure</u> contains a call to itself.
- The recursive call usually simplifies the problem and leads towards a base case which is used to break out of the recursion.
- ► Taking a look at the **call stack** helps to really understand what happens during a recursive procedure call.

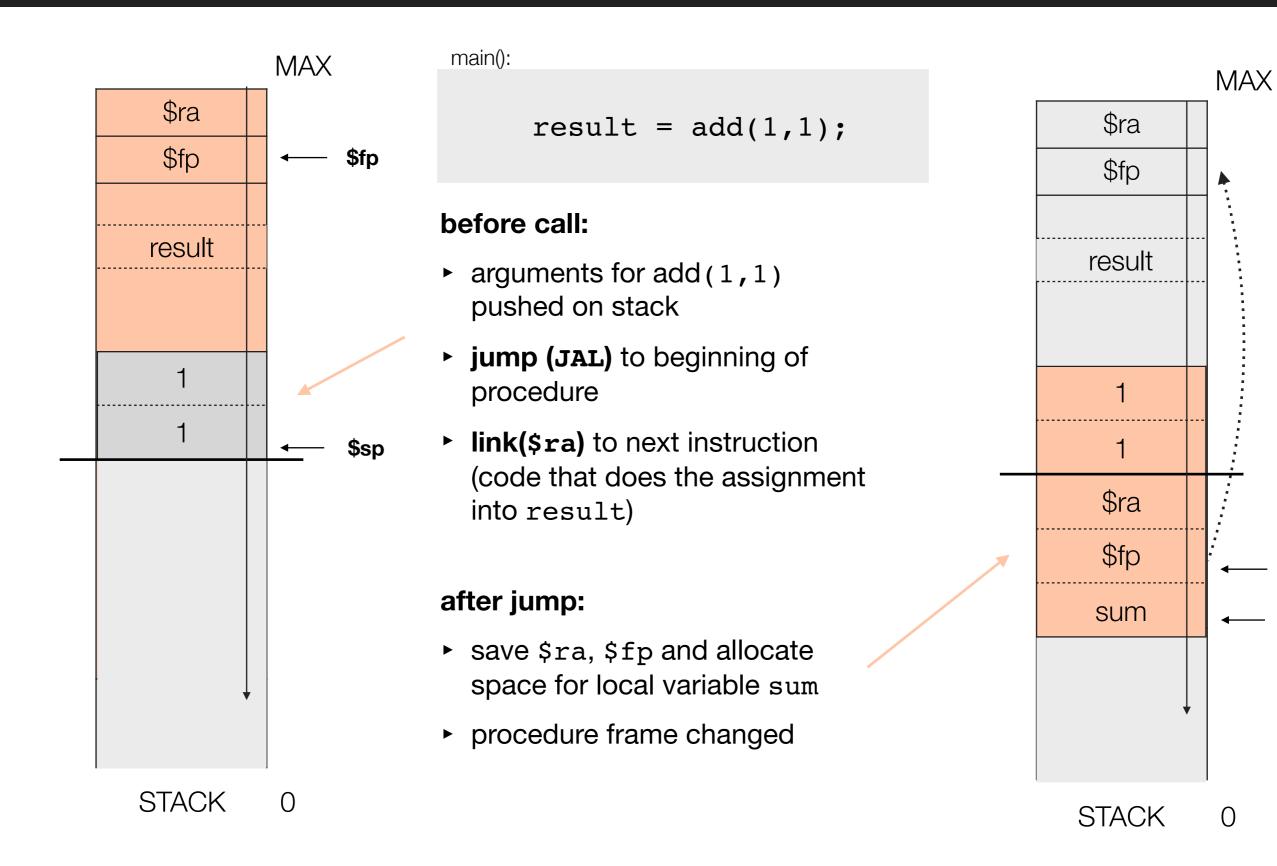
Recursion Call Recursive Procedure

```
result = add(1,1);
```

```
to solve add(1,1); simplify with recursive call solve add(1,0) + 1; base case base case add(1,0) is simple
```

We will look at the call stack during this call.

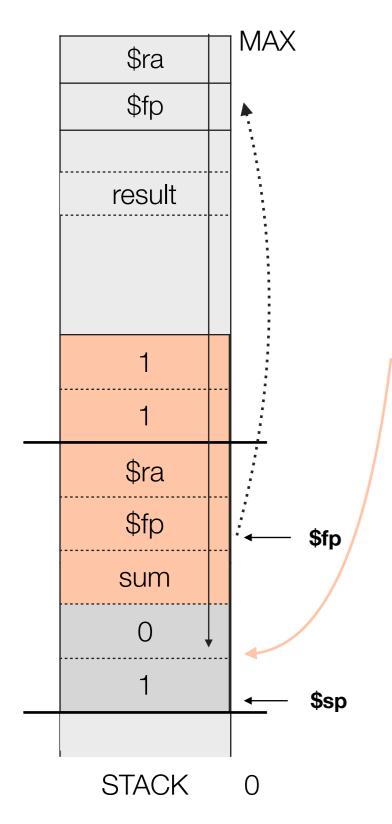
Recursive Procedure First Call



\$fp

\$sp

Recursive Procedure Recursive Call



add(1,1):

$$sum = add(0,1) + 1;$$

before call:

- arguments for add(0,1) pushed on stack
- jump (JAL) to beginning of procedure
- link(\$ra) to next instruction (code that adds 1)

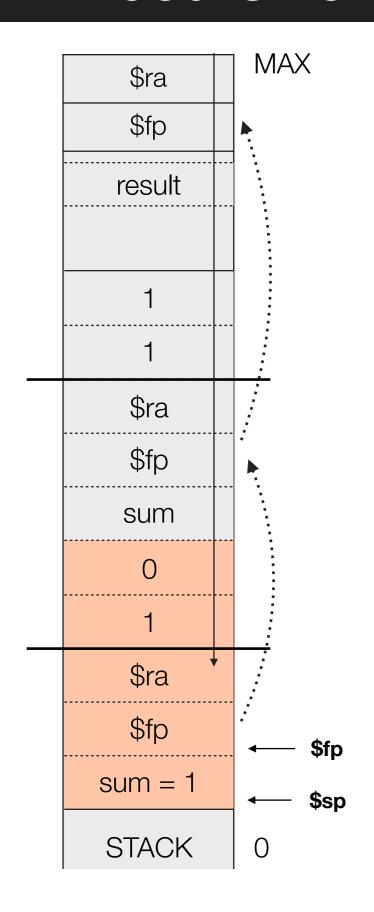
after jump:

- save \$ra, \$fp and allocate space for local variable sum
- procedure frame changed

MAX \$ra \$fp result \$ra \$fp sum \$ra \$fp \$fp sum \$sp STACK 0

different frame = different local variable

Recursive Procedure Base Case and Return



before return:

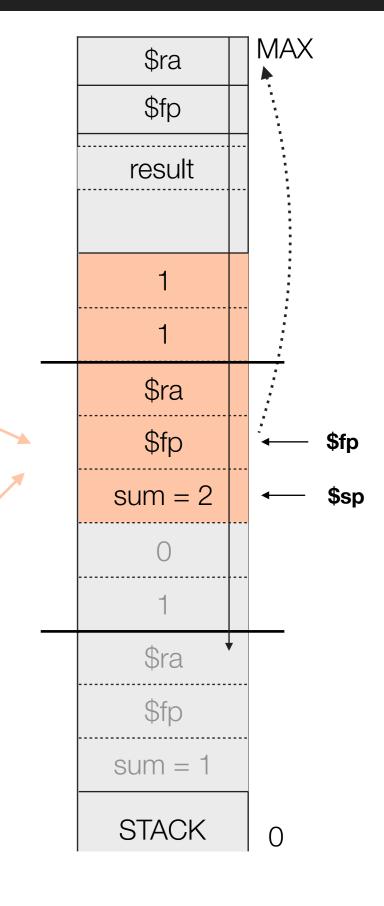
- store return value sum in \$a0 (register for return values)
- deallocate memory
- restore \$fp
- restore \$ra

after return jump:

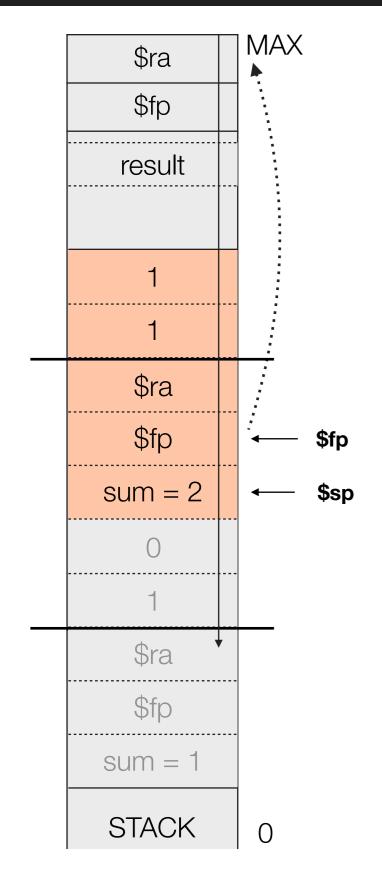
- procedure frame changed
- code that adds 1 to return value (0 + 1)

add(1,1):

$$sum = add(0,1) + 1;$$



Recursive Procedure Return



```
add(1,1):
    sum = add(0,1) + 1;
    ...
    return sum;
```

before return:

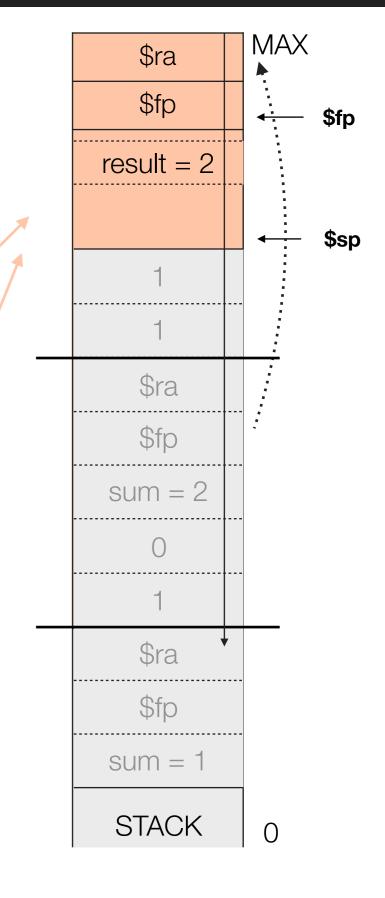
- store return value sum in \$a0 (register for return values)
- deallocate memory
- restore \$fp
- restore \$ra

after return jump:

- procedure frame changed
- code that assigns return value to result

main():

result =
$$add(1,1)$$
;



"Why talk about procedures in the 'Memory Allocation' Section?"

"Procedures are the most efficient way to allocate memory"

Memory Allocation Procedures

Memory for procedures is allocated on the stack (stack allocator).

Efficiency

- using procedures there is no need for explicitly mallocing or deallocating memory
 - → prologue and epilogue
- memory management
 - → deciding if memory is <u>live or dead</u>

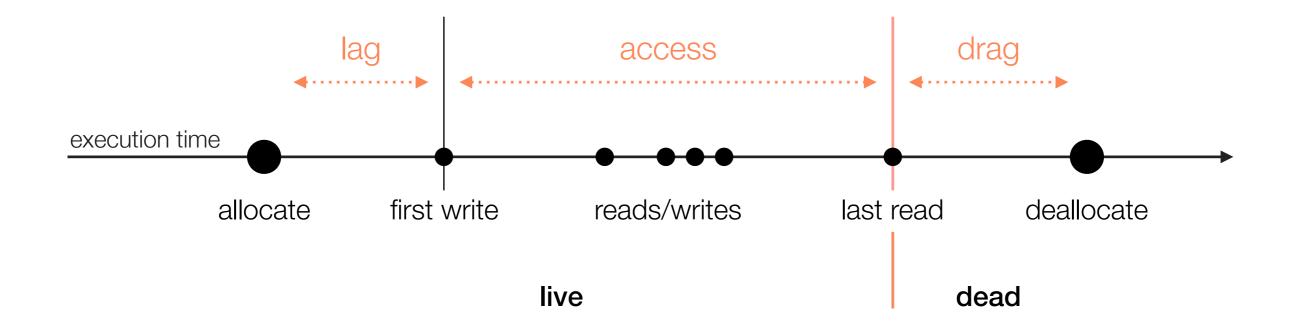
Liveness

- An address can either be live or dead
 - a live address will be read at some point in the future, whereas a dead address will not be read anymore
 - a live address can not be reused



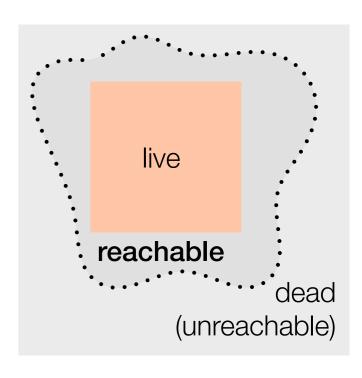
Liveness

- lag time between allocation and first write access
- period of accesses reads and writes
- drag time between last read and deallocation
- ▶ long lag or drag periods → inefficient memory consumption



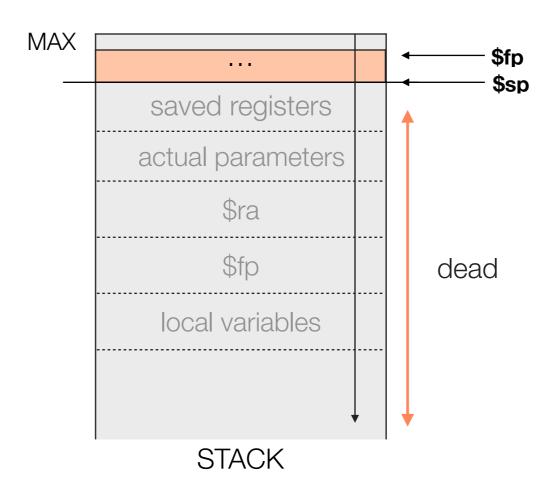
Computing Liveness

- It is undecidable whether an address is live or dead
 - neither the set of dead addresses nor the set of live addresses is computable (lines are blurred)
- We can however compute an approximation of liveness reachability
 - compute superset of live addresses
 - guarantee that everything outside is dead
 - garbage collector use this principle



Liveness Procedures

- It is undecidable whether an address is live or dead but for procedures we can at least make a clear statement about dead memory
 - jumping out of a procedure leaves the addresses below the stack pointer and above the heap dead



Control Flow

Control Flow Statements and Fixup

Control Flow



<u>Control Flow</u> is the **order** in which instructions/statements are executed. It can be described by a **sequence of program counter values**.

- Selfie features 3 control flow statements:
 - while
 - if-then-else
 - return

Expressivity of Control Flow Statements

while and for

same expressivity - one can be rewritten into the other

while/for and if-then-else

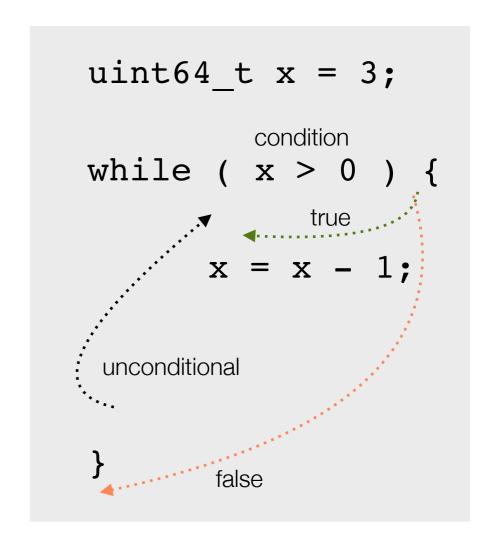
- while/for more expressive than if-then-else
- if-then-else can be rewritten to while/for but not the other way around

rewriting

 same algorithmic complexity - difference in size only a constant factor

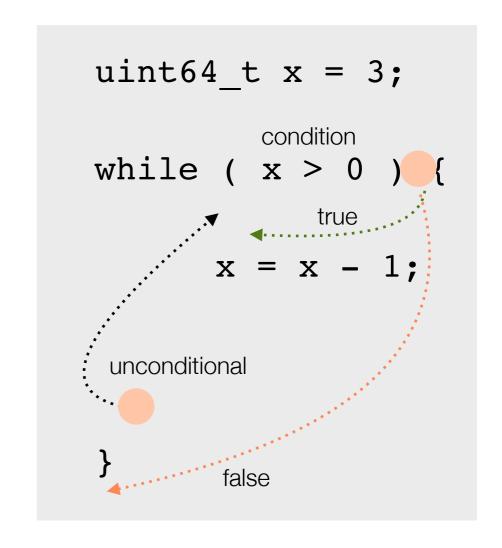
While Loop Semantics

- If the condition evaluates to true the body of the while loop is executed. Otherwise the execution continues after the while loop - <u>conditional</u> branch
- After each execution of the body the expression is evaluated again unconditional branch



While Loop Runtime vs Compile Time

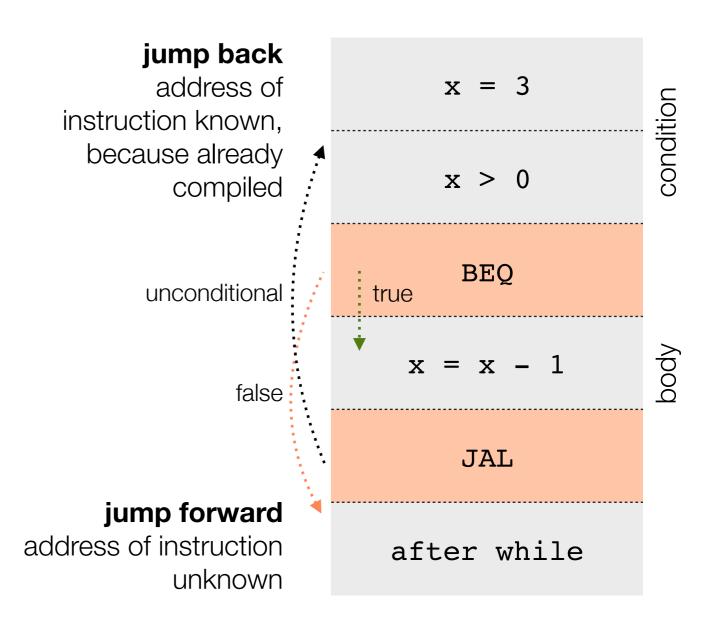
- At compile time this code is parsed once from top to bottom and the instructions are emitted in sequence. Control flow instructions are written into the binary at the right points.
- At runtime those control flow instructions provide the mechanism that allows execution to follow different paths.



While Loop Runtime vs Compile Time

Code is parsed from top to bottom and instructions are emitted in sequence.

Control flow instructions are inserted to enable branching at runtime.



Parsing While-Loop



1. Condition executed before each iteration:

TODO: **remember** address of first instruction (jump back to while)

2. Decide where to continue execution:

TODO: emit **BEQ** instruction to unknown address and remember address of this "wrong" instruction

(jump_forward_to_end)

3. Beginning of body:

TODO: parse body

4. End of body and back to 1:

TODO: emit **JAL** instruction to remembered address (jump_back_to_while)

5. End of while (target address of BEQ known):

TODO: fixup BEQ instruction

Parsing While-Loop



1. Condition executed before each iteration:

TODO: **remember** address of first instruction (jump back to while)

2. Decide where to continue execution:

TODO: emit **BEQ** instruction to unknown address and remember address of this "wrong" instruction

(jump forward to end)

3. Beginning of body:

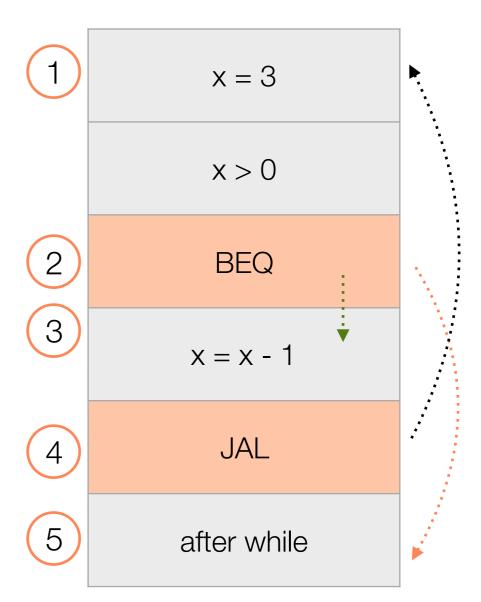
TODO: parse **body**

4. End of body and back to 1:

TODO: emit **JAL** instruction to remembered address (jump back to while)

5. End of while (target address of BEQ known):

TODO: fixup BEQ instruction



Jump and Branch

Jump

- already know absolute address at runtime (strong assumption)
- 20 bit for address

Branch

- addressing relative to program counter
- relocating code becomes simple
- limited branch distance (12 bit for address)

Fixup

A fixup is the process of changing (correcting) the destination **address** of a control flow instruction after it has been emitted.

- Fixups solve a problem that is present in all <u>single-pass</u> compilers for languages featuring procedures or control flow statements.
- Necessary when destination address of instruction is unknown when the instruction has to be emitted.
 - instruction is emitted with wrong address
 - address needs to be fixed up later when the address is known

Fixup

- Fix up a single instruction
 - control flow statements (while, for, if-then-else)
- Fix up multiple instructions
 - control flow statements (if-elseif)
 - return statement
 - procedure calls

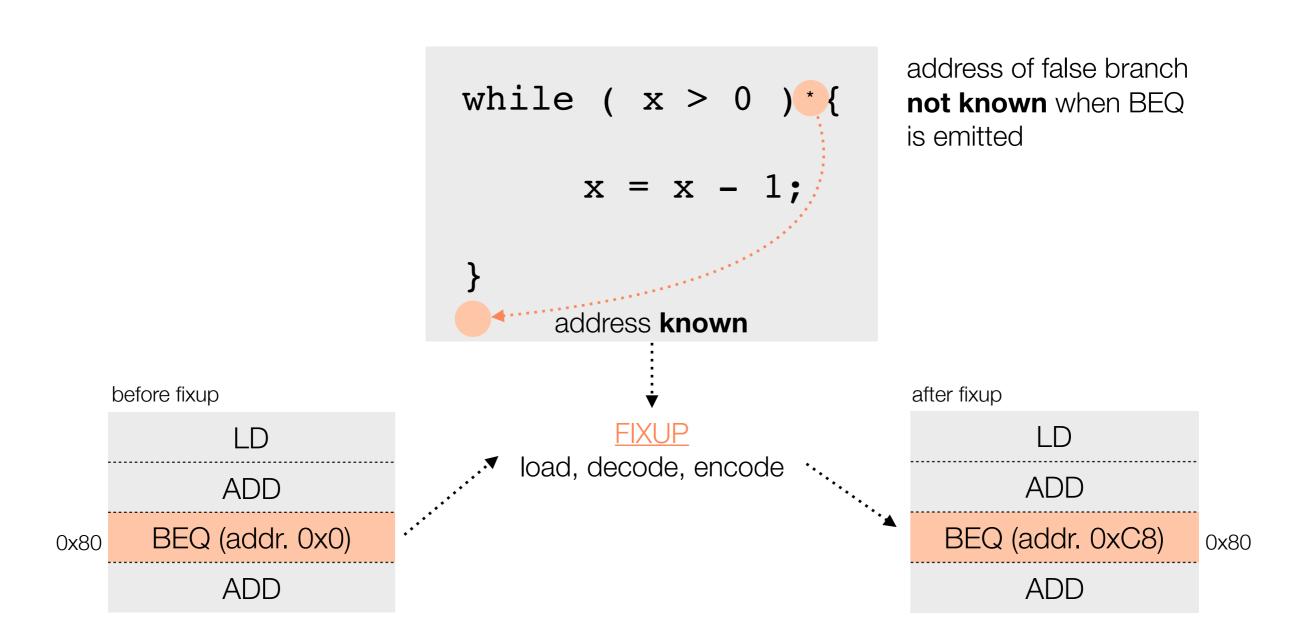
Fixup Single Instruction



- JAL or BEQ instructions emitted when compiling while, for or if-thenelse
- fixed number of instructions to fixup
- How to fix up an instruction
 - emit instruction with wrong address and remember address of this wrong instruction (local variable to enable nesting)
 - 2. fix up instruction as soon as address is known
 - load and decode instruction
 - encode and emit instruction with correct address (at old address)

Fixup Single Instruction





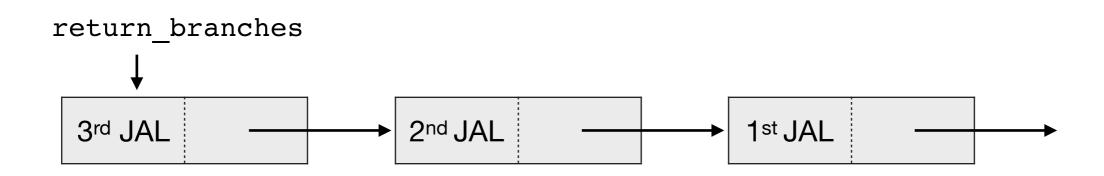
Remember address of wrong instruction:

branch_forward_to_end = 0x80

Fixup Multiple Instructions - Return



- return statement
 - compiling return → JAL to first instruction of epilogue that is not known at this point
 - a procedure might have any number of return statements
- keep a linked list of JAL instructions that will be fixed up later by going through the list and fixing up each instruction (this also deletes the list)



Fixup Linked List of JAL Instructions



Where is this list stored?

- in the binary code, using the instructions to store the list (elegant solution by N. Wirth)
- the bits designated for the destination address of an instruction are used to store the address of the previous instruction
- the head of the list is stored globally (<u>return_branches</u>)
 (no nesting of procedure definitions)

	binary code			
0x80	JAL	0	$\overline{}$	0
				previous
0xC8	JAL	0x80		
				previous
0x120	JAL	0xC8	_	head of list
				<pre>(return_branches)</pre>

Fixup Multiple Instructions - Procedure Calls

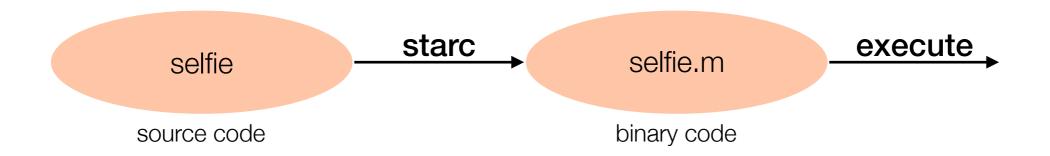


- procedure call
 - compiling call → JAL to first instruction of prologue that might not be known at this point (definition not parsed)
 - a procedure might be called any number of times
- using same principle as for return statements linked list in code
- nesting of procedure calls is possible, therefore storing the head globally is not sufficient
 - use symbol table instead
 - associate the procedures name with the head of the list (procedure address)

Further Concepts & Selfie

Compilation,
Object-Oriented Programming,
Bootstrapping Selfie

Compile Time and Runtime



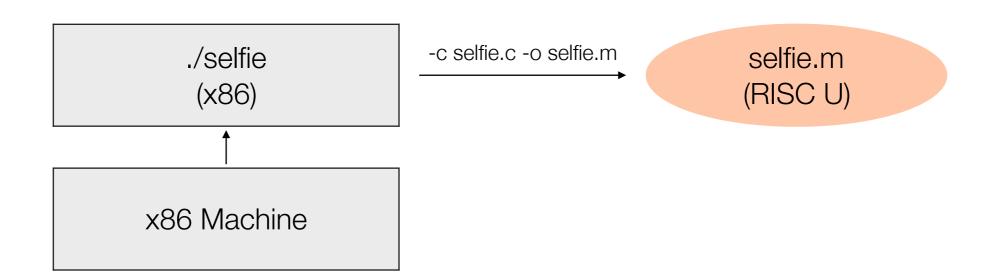
- Compile time program is executed to generate code
 - compilation done once
 - figure out as much as possible before execution
 - infinitely many programs can be created
- Runtime execution of compiled program
 - heap and stack built for/by the program

Compilation

Terminology

Cross Compilation

- A <u>cross compiler</u> outputs binary code that targets a different architecture than its own binary code does.
- for embedded systems, iOS and Android,...

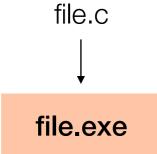


Compiling Source Code

So far we considered our code to be in a single source file.

```
...
uint64_t main() {
   return proc();
}

uint64_t proc() {
   return 1;
}
...
```



Symbolic reference:

reference associated with a name/identifier (ex. proc()).

The address is not known at this point and the symbol table entry for proc2() contains the last jal instruction that needs to be fixed up.

Resolved into **direct address** as soon as address is known.

Direct address: different namespace, the one of machine instructions

Compiling Source Code

- To manage code better it is convenient to have multiple source files contain different parts of the code. However, we still want to use (<u>reference</u>) procedures and variables from those other source files.
- Forward declarations are unavoidable when we reference code from other files.
- But the necessary fixup cannot be performed until the file containing the definition of the referenced procedure/variable is compiled. Therefore a file can still contain <u>symbolic references</u> after it has been compiled.

```
...
uint64_t main() {
  return proc();
}
...
file1.c
symbolic
reference
uint64_t proc() {
  return 1;
}
...
file2.c

file2.c
```

Independent vs Separate Compilation

- (independent) compilation: compiling two separate source files that do not reference each other.
- separate compilation: compiling two separate source files that reference each other.

Linking

The fixup process that resolves symbolic references into direct addresses across object files.

- Compiling source code yields <u>object code</u> that might hold symbolic references.
- Symbolic references are part of unresolved fixup chains that are stored in the symbol table.
- The <u>linker</u> resolves all symbolic references and generates executable binary code.
- Linking can be done at compile time (static) but also at runtime (dynamic loading and linking)

Linking

```
symbolic
                                      reference
          uint64_t main() {
                                                  uint64_t proc() {
              return proc();
                                                      return 1;
Source file
                   file1.c
                                                           file2.c
                                                                separate
                        separate
                        compilation
                                                                compilation
Object file
                   file1.o
                                                           file2.o
                                                                      at compile time
                                        linker
                                                                      at runtime
Executable file
                                       file.exe
```

Object- and Executable Files

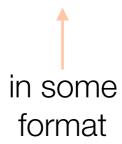
Object files

- might contain symbolic references (associated with name)
- more generic than executable files
- to resolve the references, the information from the symbol table is needed (fixup chains).

object file = binary code + representation of symbol table

Executable File

- no symbolic references
- special case of object file



Dynamic Loading and Linking Java

The file that gets executed does not have to be completely resolved. Under some circumstances the execution of an object file makes more sense.

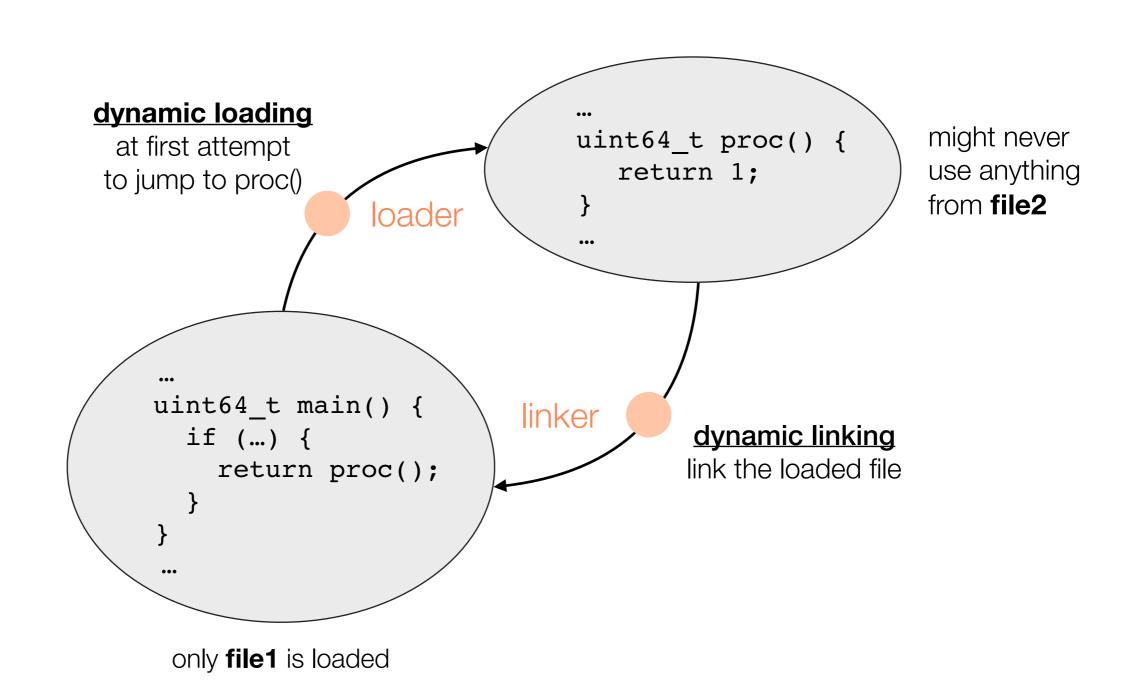
Why?

the referenced code is never used at runtime, so why link it?
 (this is specific to each run)

How?

 Code is loaded and linked at runtime, as soon as the executed program tries to jump to an unresolved address.

Dynamic Loading and Linking Java



Dynamic Loading and Linking Java

- Using this concept we could start from an empty file combined with the information from the symbol table.
- 1. Attempt to load the main() procedure
 - find the code in memory
 - load the code into main memory and execute it
- 2. Attempt to use another unresolved reference
 - find code in memory
 - load code into main memory
 - link code and continue execution

Just-In-Time (JIT) Compilation

A <u>JIT compiler</u> performs **binary translation** of **hotspots** in **bytecode**.

- It is not used with high-level languages.
- It is part of a runtime system such as the <u>JVM</u>.
- Used to increases runtime performance of code.

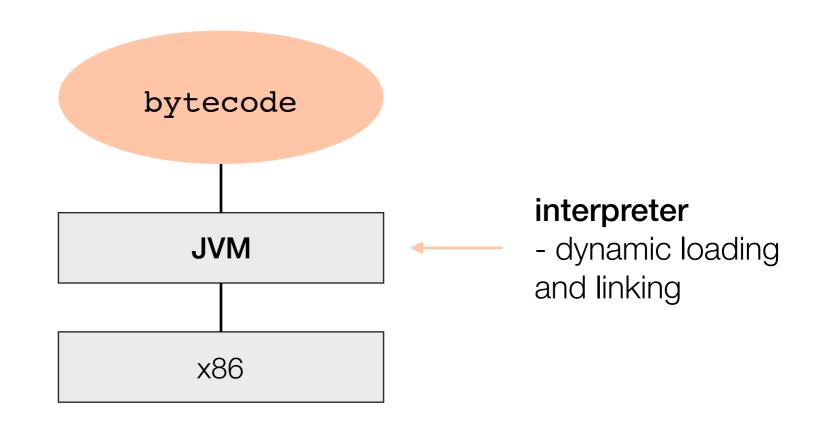
Bytecode



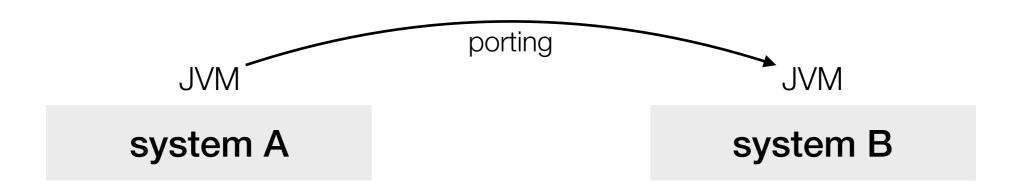
- Bytecode is code for which there exists no processor that can execute it.
 - a (virtual) machine is needed that can load and execute byte code.
- It is independent of the architecture it runs on.

Java Virtual Machine JVM

- The JVM is an emulator of a byte code processor.
 - it executes bytecode through interpretation
- It is a runtime system more complex than an OS.



Bytecode



► C

 change backend of compiler so it produces binary code for a different architecture

Java

- compiler remains unchanged
- port JVM to run on different architectures (portability)

Porting the JVM

 change backend of the compiler that compiles the JVM for system A so it compiles the JVM for system B

The Problem Interpretation is slow, Execution is fast

Ex.: Interpreting a while loop in bytecode hundreds of times is slow.

Make it faster

- instead of interpreting the bytecode feed it a compiler that translates it to binary code that can be executed by the machine it is running on (just-in-time).
- It is <u>hotspots</u> in bytecode that get translated to binary code
- A hotspot is binary code that gets executed often (enough) within a certain amount of time.

Object Orientation

and Polymorphism

Object-Oriented Programming OOP

- <u>object-oriented programming</u> is a programming paradigm
- crucial factors of OOP
 - classes abstract data types
 - structure hierarchy
 - and most important <u>polymorphism</u>

Polymorphism of Procedures

The actual procedure implementation is **selected at runtime according to the data types** of the parameters the procedure is called with.

- The actual type of the parameters is not known at compile time
 - → compiler does not know what procedure to call
- The runtime system needs to decide which procedure to call
 - → determining the data types and choosing the implementation (dynamic binding)
- On the level of procedures we are talking about polymorphism. But in object-oriented languages this concept is implemented using dynamic binding.

Memory safety is a property of program execution that is concerned with memory access violations.

- We distinguish between:
 - spatial safety memory is legally allocated
 - temporal safety no access to undefined memory

- This problem does not exist in Java since you cannot...
 - ... access objects you deleted
 - ... perform illegal access into arrays
- Therefore Java is a memory-safe language
 - You cannot write a program that accesses memory outside its own space (defined and allocated)

What does it need?

- no pointer arithmetic or operators that access memory directly (only internally)
- array range checking at runtime
- using references
- a garbage collector when memory is deallocated

Memory Safety and Garbage Collection

- Deallocating memory
 - using free → memory safety would be lost because dangling pointers might exist
 - a garbage collector ensures memory safety when dealing with memory deallocation

Garbage Collection

and Memory Deallocation

Memory Deallocation

- many but finitely many addresses
- memory deallocation only necessary if addresses need to be reused
 - only dead memory can be reduced
 - not necessary for short running programs
 - tip: only use free in long running programs (server,...)
- manual memory management
 - → programmer explicitly deallocates memory (free)
- <u>automatic memory management</u>
 - → garbage collector reclaims dead memory

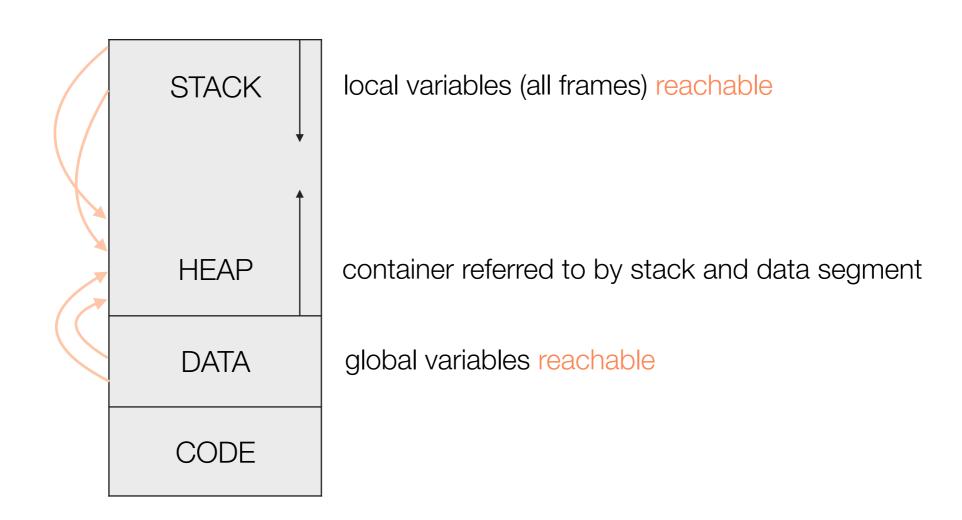
Garbage Collection

A garbage collector automatically **reclaims dead memory** so it can be reused.

- garbage collection is not restricted to object-oriented languages
- a garbage collector tries to find proof that memory is dead
- a garbage collector computes reachability (liveness is undecidable)
- a garbage collector computes the set of <u>unreachable</u> addresses
 - check if an address is still reachable
 - are there any direct or indirect references?

Garbage Collection Check References

- look for references on the stack, heap and in the data segment
- and the <u>transitive closure</u> all possible references (<u>tracing</u>)



Garbage Collection in C

- in C we are able to talk about the actual value of pointer
- conservative garbage collector (tracing)
 - an approach to implement a garbage collector for C and C++
 - see everything on stack and in data segment as pointer (+ transitive closure)

Garbage Collection

Advantages

garbage collection rules out <u>dangling pointers</u>

Disadvantages

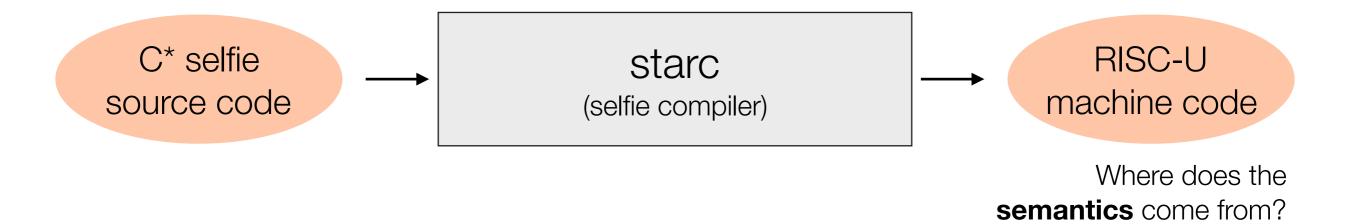
- Garbage collection does not protect against <u>reachable memory leaks</u>
- garbage collections consumes additional resources
 - → might have a negative impact on performance

Reachable Memory Leaks

- memory leaks are caused by incorrect memory management
 - not releasing memory that is not needed or accessible any more
- reachable memory leaks occur when memory is still reachable but not accessed any more
 - hold on to references (i.e. hoard in hash table)

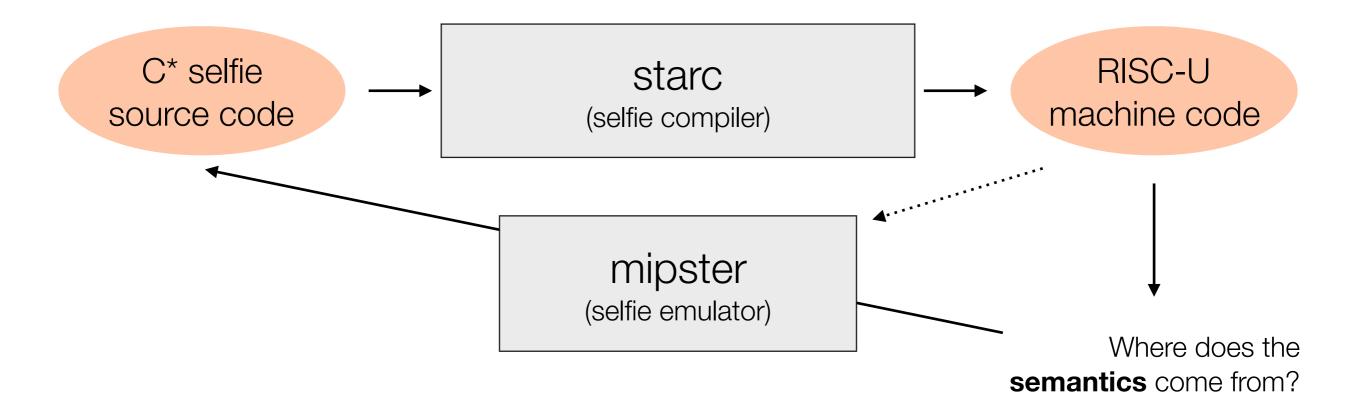
Bootstrapping Selfie

Semantics of Selfie?



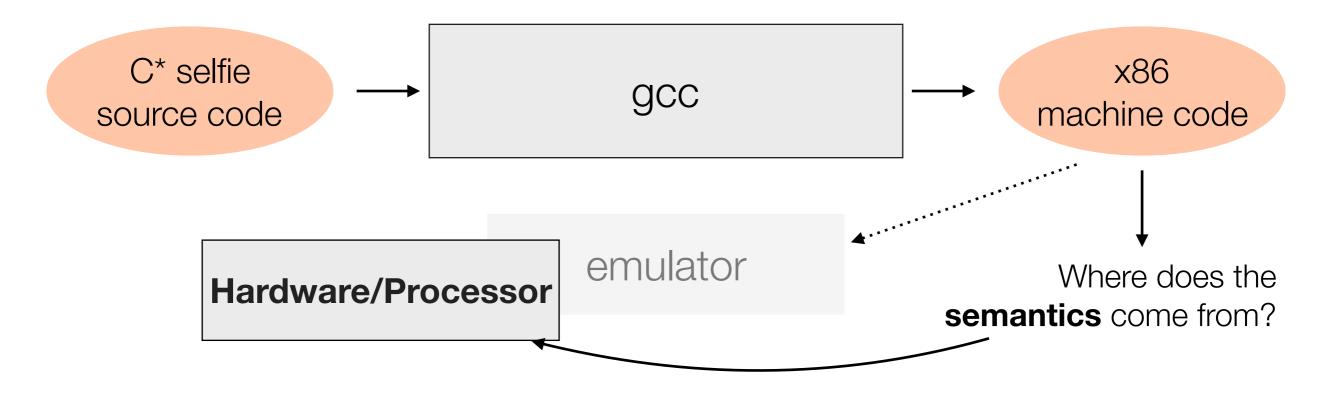
- "The semantics of an instruction is determined by how the processor implements it"
 - Semantics of RISC-U
- Machine code is a sequence of instruction, therefore the processor that executes it determines its semantics.
- Execute RISC-U code when only x86 processor available?
- Fortunately, selfie also features an emulator that can interpret RISC-U instructions.
 - selfie defines the semantics of selfie → cycle
 - we need executable of selfie to compile and interpret selfie.
 Where does this executable come from? → cycle

Semantics of selfie? Cycle



- We define the semantics in selfie:
 - how code is generated and interpreted by selfies RISC-U emulator
- The semantics of selfie.c is in selfie.c cycle

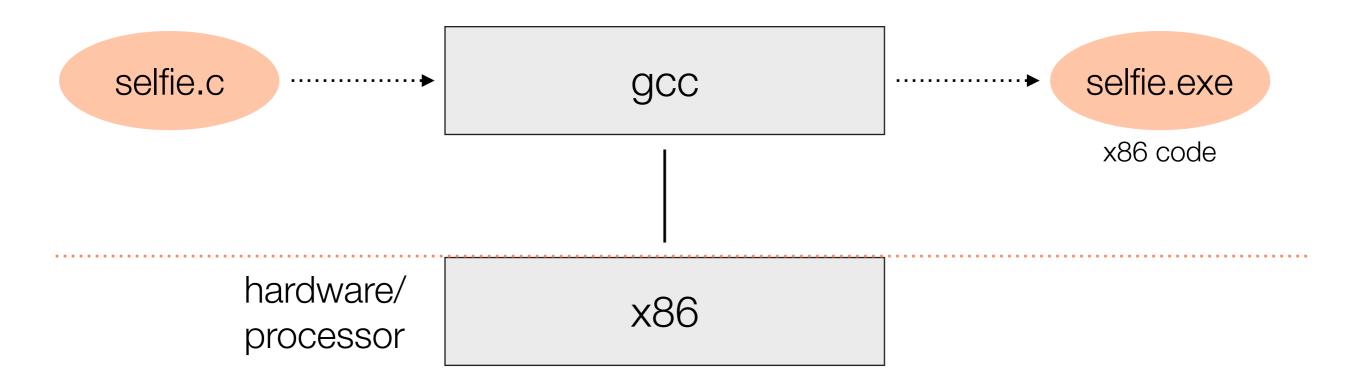
Breaking the Cycle



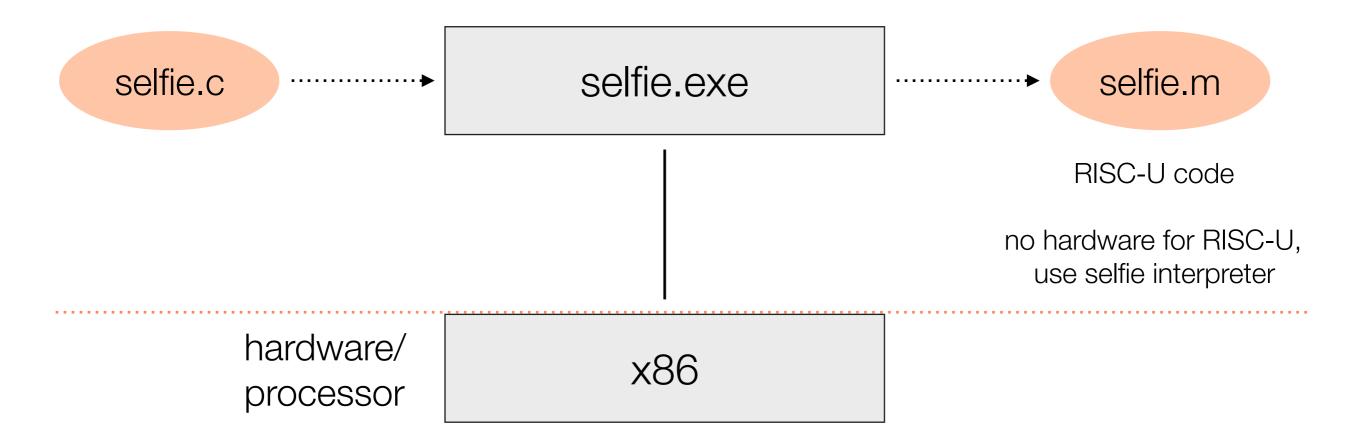
- We break the cycle when we compile the first selfie.c with gcc to get our first executable of selfie.
- x86 code does not need an emulator. It is executed directly on hardware.
- ► Therefore the semantics is ultimately defined by the system/hardware it runs on.
- Processor errors might affect the semantics of code.

We only have the **selfie.c source file**

Execute gcc
 compile selfie.c to get an x86 executable version of selfie

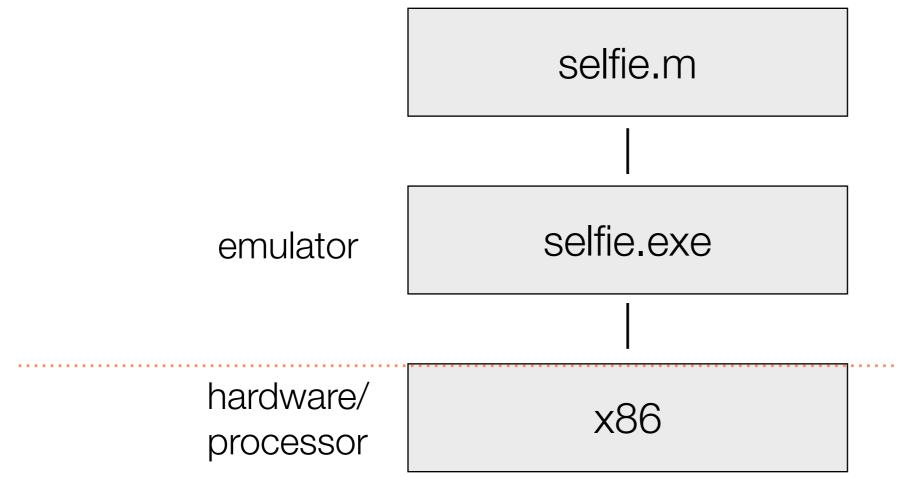


- Execute selfie.exe compile selfie.c to get an RISC-U executable version of selfie
 - → self-compilation

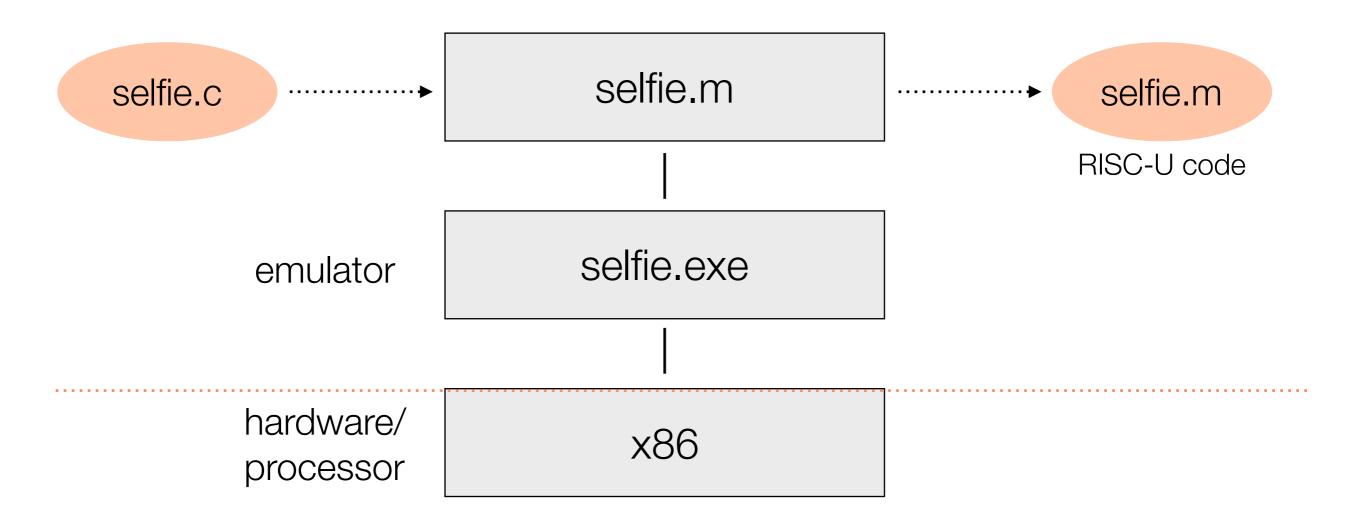


3. Execute **selfie.exe** start the selfie emulator and have it **interpret** selfie.m

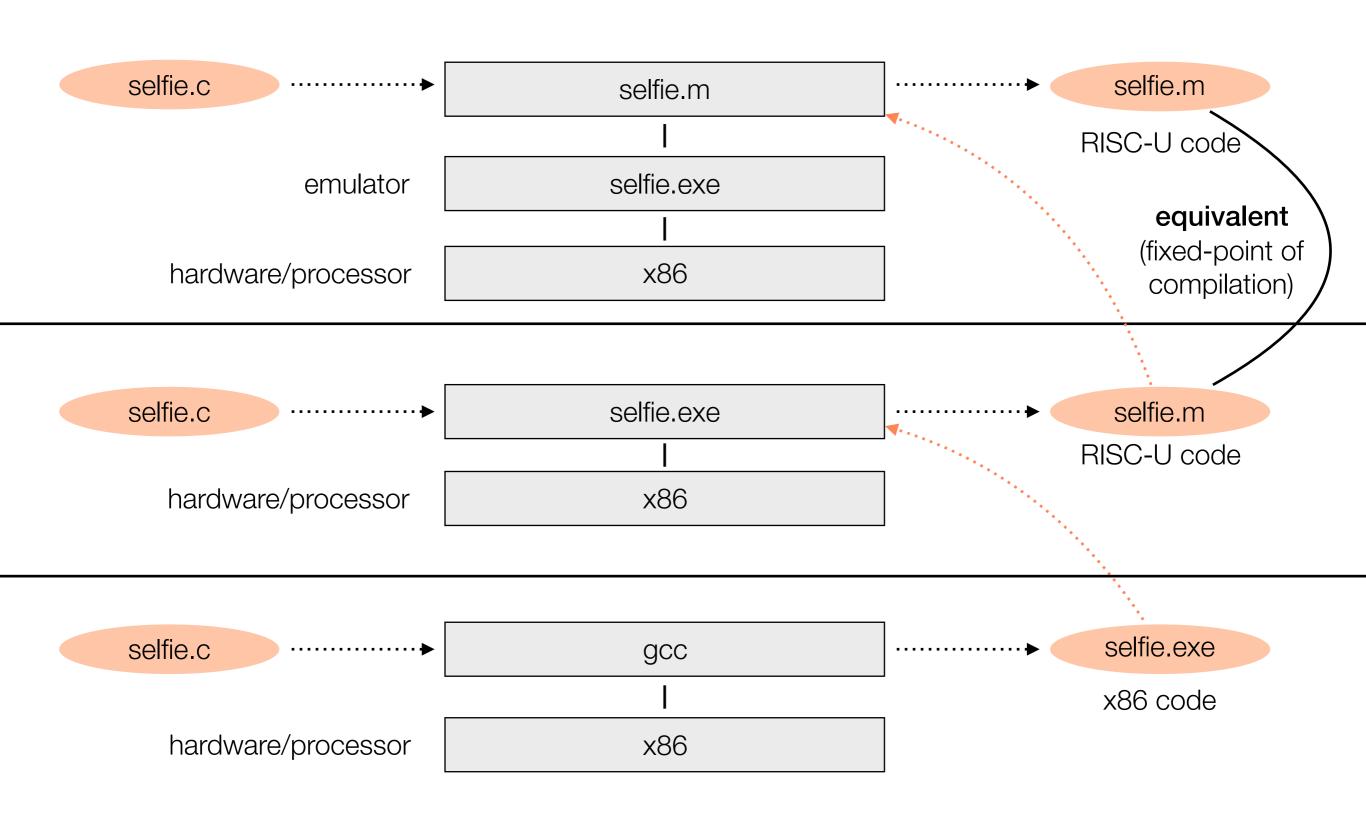
→ self-execution



- 4. **Interpret** selfie.m **compile** selfie.c to get an RISC-U executable version of selfie
 - → again self-compilation



The Whole Picture



Command Line

	execute selfie.exe	compile selfie.c	•	output RISC-U code into selfie.m	
2.	./selfie	-c selfie.c	-o self:	ie.m	
3.	./selfie	-l selfie.m	-m 1		
4.	./selfie	-l selfie.m	-m 1	-c s	elfie.c
		load selfie.m and interpret it on mipster			compile selfie.c

► Compiled code can be executed directly:

./selfie -c selfie.m -m 1 -c selfie.c

Introduction to Operating Systems

Operating Systems

An operating system manages the **execution** of a program on a machine.

- We will extend and refine this definition and build an understanding of operating systems as we extend the minimal operating system support already implemented in selfie.
- However, before we address operating systems, we will learn how a program is executed by selfie.

Previously on...

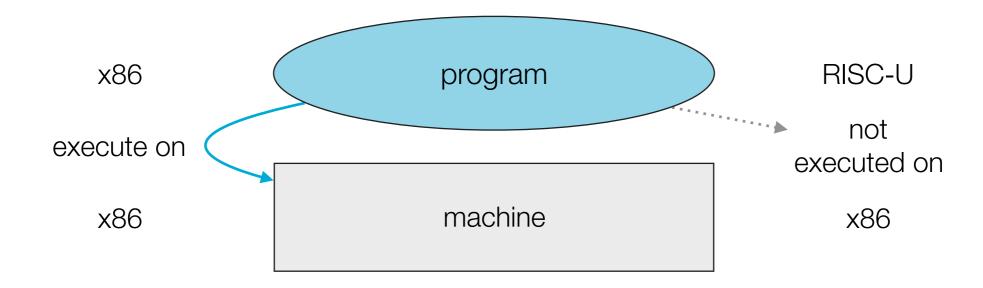
 Introduction to compilers explained how source code written in C* is translated into an executable RISC-U binary.



Now we want to execute this compiled program.

Executing a Program

- The machine provides resources for the execution of a program (CPU, memory, I/O devices).
- The machine can only execute machine code that is written in its machine language, which is defined by the machine's <u>instruction set</u> <u>architecture</u> (x86, ARM, MIPS, RISC-V,...).
- We cannot execute RISC-U code directly on an x86 processor. Therefore, we will use **mipster**, a simple emulator for RISC-U code, as we did when explaining <u>bootstrapping</u>.

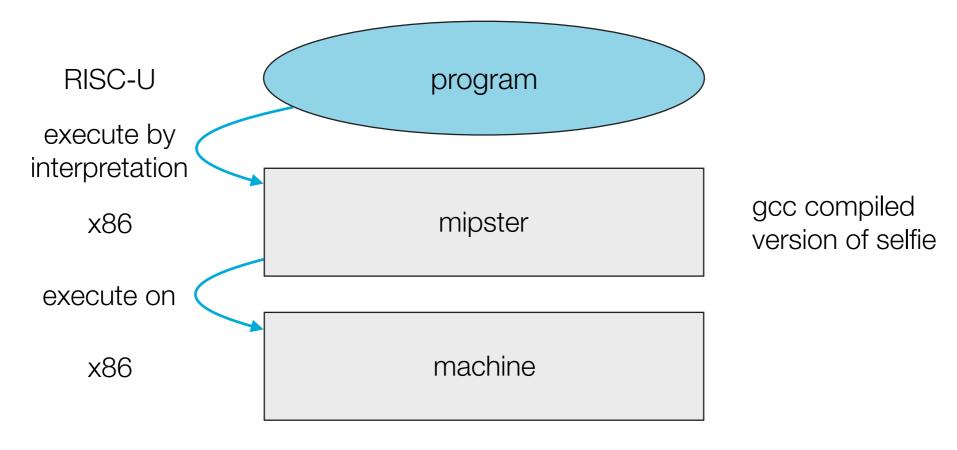


Mipster

Emulation, Context, Process

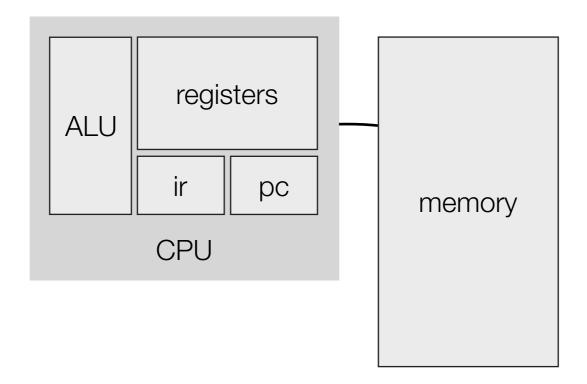
Mipster A RISC-U Emulator

- Mipster is an <u>emulator</u> that
 - emulates a RISC-U processor in software
 - can execute RISC-U code by interpreting it
- To understand how mipster works we take a look at the machine it emulates.



Mipster A RISC-U Emulator

- emulates <u>Von Neumann machine</u>
 - mipster creates a machine instance with 32 registers and 0-64 MB of memory
- How does emulation in selfie work?
 - 1. selfie executed with -m option starts mipster
 - 2. mipster creates a machine instance in software
 - 3. mipster starts emulation by implementing fetch-decode-execute cycle (fetch instruction, decode instruction, execute instruction by interpreting it)



Mipster A RISC-U Emulator

We will use the following example to look at each step that is necessary to start mipster and have it emulate a program.

The following slides explain the design and implementation of mipster and are best studied with the actual code next to them. However, it is not necessary to understand every little detail at this point.

1. Execute Selfie main()

```
./selfie -c program.c -m 1
```

- Executing ./selfie starts the main() procedure where:
 - the selfie system is initialized → <u>init_selfie(...)</u>, <u>init_library()</u>
 - selfie is run \rightarrow selfie()

1. Execute Selfie selfie()

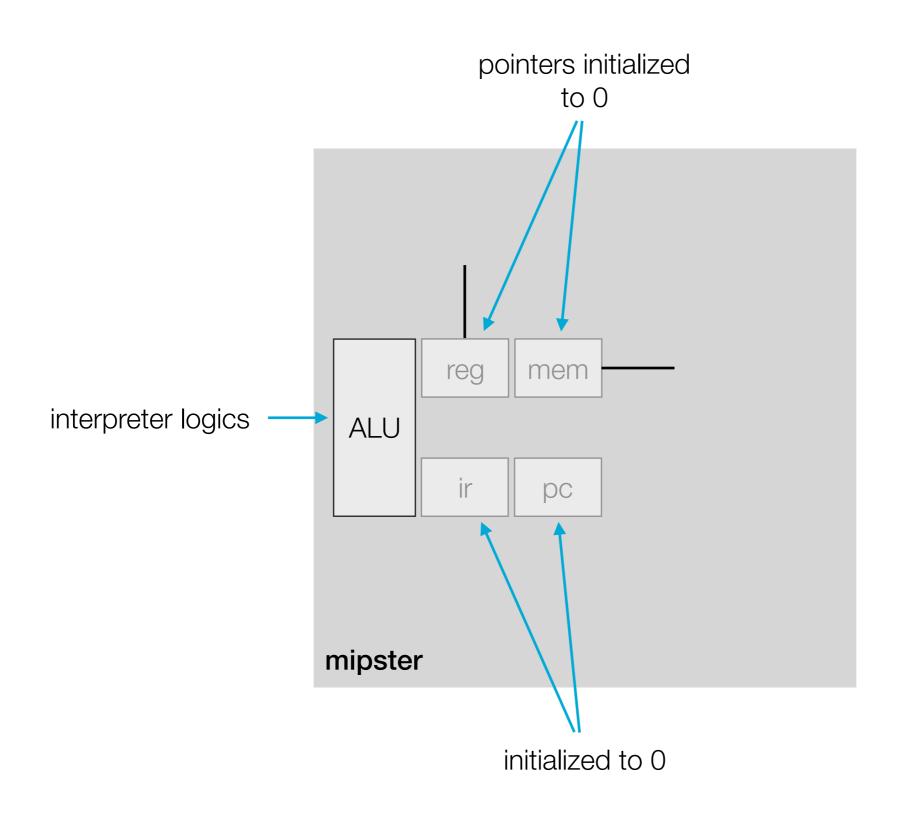
```
./selfie -c program.c -m 1
```

- selfie() This is selfie and all it can do.
- The arguments provided to ./selfie determine what part of selfie is executed.
 - -c option starts the compiler → selfie compile()
 - -s option disassembles binary code → selfie disassemble()
 - -1 option loads a binary → selfie_load()
 - options to execute code on different machines → selfie run(machine)
 (-m option starts mipster)
 - invalid options print help → print usage()

- <u>selfie_run(...)</u> initializes mipster and creates a machine instance
- machine state is represented by a set of global variables
- two main steps
 - setting up mipster
 - creating something called a <u>context</u>

```
./selfie -c program.c -m 1
```

- flags are set depending on the machine that is run
- initialize memory size of machine → global variable
 - init memory(atoi(peek argument()))
- reset/initialize mipster → global variables program counter(pc), instruction register(ir), pointer to memory(pt), registers
 - reset interpreter()
 - reset_microkernel()

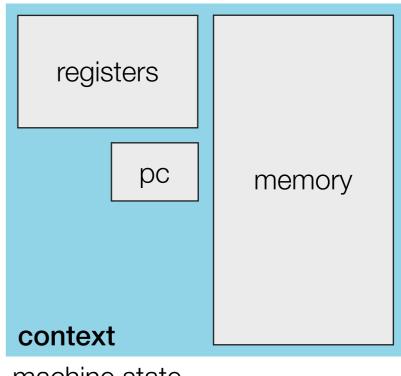


This is the bones of the emulator

Context For Now

For now - A container that holds part of the machine state.

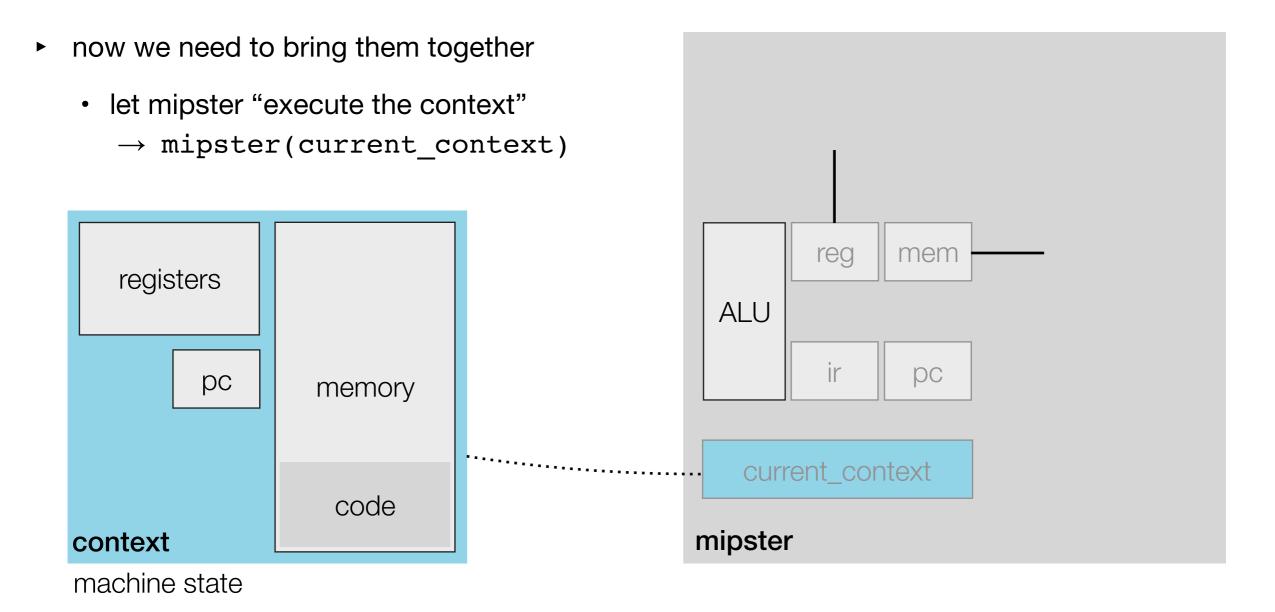
- selfies context is explained in more detail <u>later</u>
- stores
 - program counter
 - registers
 - memory (page table)
- a context is uniquely identified by its address



machine state

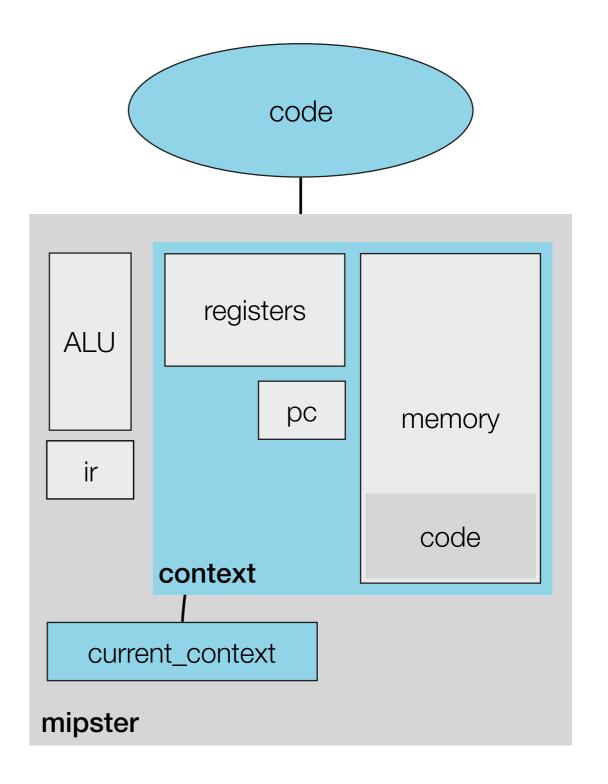
- create a context
 - allocate space for the context and its components (memory, registers) →
 allocate_context()
 - the machine for which the context is created is the parent of that context
 - MY_CONTEXT to the machine
- ▶ upload binary into the current context → up_load_binary()
- ▶ set binary name as first argument that will be passed to context
 → set_argument()
- pass name and remaining arguments to context
 - arguments for the emulated program
 - \rightarrow up load arguments()

- so far we have
 - an new and "empty" machine
 - a context representing part of the machine state



3. Start Emulation

- mipster is provided with the context it is supposed to execute
 - mipster(to_context)
- a context switch is performed
 - load the context into mipster and execute it
 - mipster_switch(to_context)



"Why create a context in the first place and not set up mipster right away?"

"The concept of a context is part of the operating system functionality already provided by selfie. It will make what comes next a lot easier."

3. Start Emulation mipster_switch

- ► the heart of mipster is the procedure mipster switch(to context,...)
- it is composed of 3 parts
 - the actual context switch where the contents of the context are loaded into the machine.
 - → do_switch(...)
 - the execution of the context until the occurrence of an exception
 (system call, timer interrupt,...). The implementation of the von Neuman
 cycle.
 - → run until exception()
 - **saving** the context before returning to mipster() after an exception (storing machine state back into context).
 - → save_context(...)

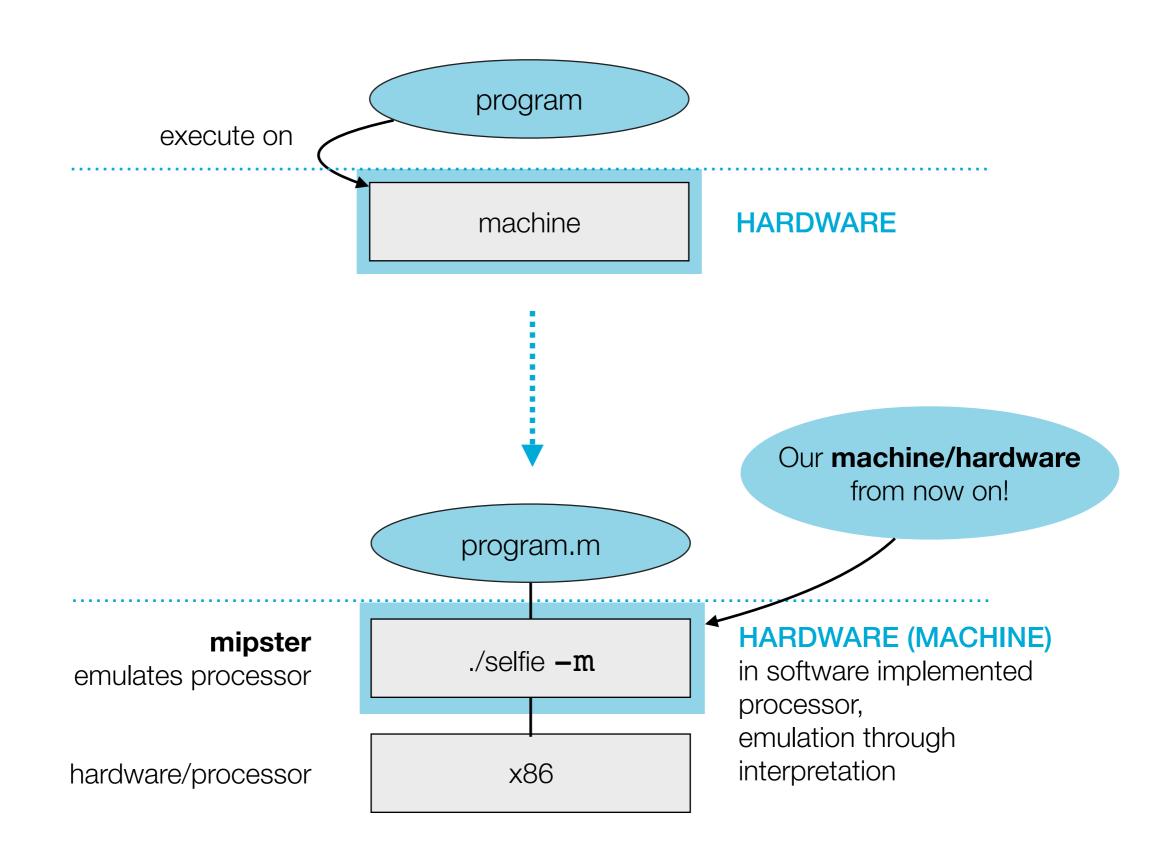
3. Start Emulation mipster(to_context)

```
after TIMESLICE many
                                   timeout = TIMESLICE;
instruction execution will be
interrupted
                                   while (1) {
                                       from context = mipster switch(to context, timeout);
mipster switching to and
executing the context
                                       if (get parent(from context) != MY CONTEXT) {
not important yet (the
created context is always
MY CONTEXT)
                                       else if (handle_exception(from context) == EXIT)
                                           return get exit code(from context);
exception handling - only an.
                                       else {
exception that yields an
                                           to context = from context;
EXIT breaks the loop and
                                           timeout = TIMESLICE;
exits the emulation
renew TIMESLICE and
switch back to context
```

Summary Mipster

- RISC-U code can not be executed on an x86 processor.
- We use mipster, an emulator for a RISC-U processor in software to execute RISC-U code → create a process
- From now on we consider this first mipster (x86 version) to be our machine(hardware).

Summary Mipster



Process

A <u>process</u> is a program **in execution**. More precise, it is an **abstraction** of a running program that is created by the OS.

▶ program → passive

- executable binary
- · sequence of machine instructions somewhere on disk
- "a program is executed"

▶ process → active

- execution of binary code, pc pointing to the next instruction to be executed
- · code in memory and a set of resources available to the process
- processes are isolated, no communication between processes
- "a process executes"

Process Looks familiar?

Resources owned by the process

memory

on disk

- processor state
 - → content of registers, pc, ...
- Context is part of the process
 - can be seen as snapshot of process (in selfie)

CODE CODE

program load into memory process

high STACK **HEAP** DATA CODE

in memory

low

REG

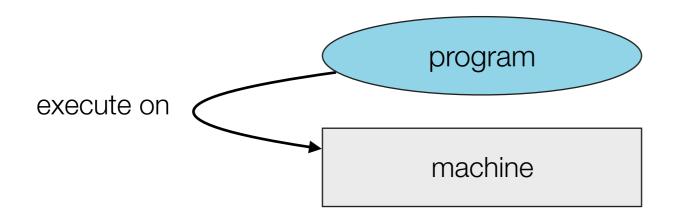
"Now we have enough insight into selfie and mipster to get started with operating systems."

Operating System

Concurrency,
Shared Resources,
Process Model,
Context

Operating System Introduction

- The machine provides resources/hardware for the execution of a program (CPU, memory, I/O devices).
- Execution of a program should be easy, stable, fast,...
- Managing resources is not a big issue if only a single <u>sequential</u> program would be executed at any time (as mipster does).
 - → but we want to execute many, possibly <u>concurrent</u>, programs simultaneously



Concurrency and Operating System

Goal

- execute many programs seemingly at the same time
- Problem is limited resources
 - one physical CPU
 - machine can only execute 1 instruction per core
 - fixed number of cores, but different number of applications
 - one physical memory
 - fixed-sized memory

Our Goal

learn understand how these problems are overcome

Concurrency and Operating System

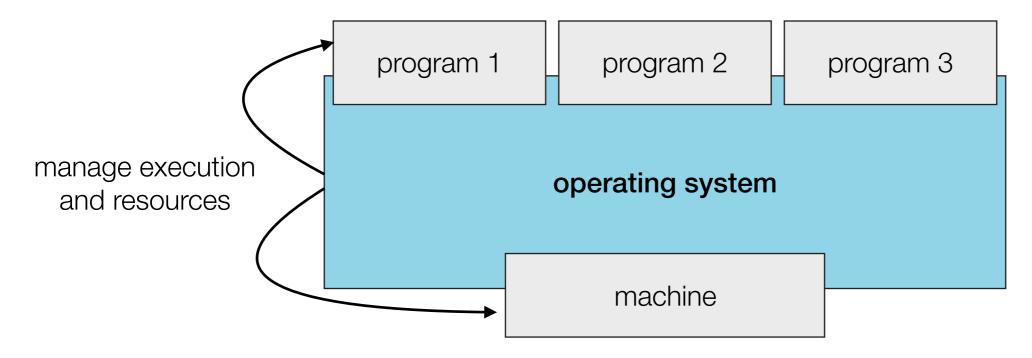
- Where does the need for concurrency come from?
 - machine interacts with real world, where things happen in parallel
 - machine needs to reflect and deal with that parallel mindset of users
 - a concurrent programming model is intuitive

Concurrent or Parallel

- concurrency a property of software
 - illusion of running many processes at a time
 - can be achieved by:
 - time-sharing (single core)
 - execute in parallel (multi-core)
- parallel a property of hardware
 - truly in parallel, simultaneously hardware support essential
 - divide workload to increase performance
 - program itself does not have to be concurrent

Operating System Introduction

- As we know, we can execute many programs on a single machine at the same time, all of which want their fair share in resources → thanks to the OS
- The OS acts as an intermediary between processes and machine,
 - it manages resources.
 - it manages and controls the execution → protection, error management,....
- The OS consists of several <u>components</u> and we are primarily talking about the OS kernel.



Operating System

An operating system manages resources and controls the **execution of many processes** on a machine.

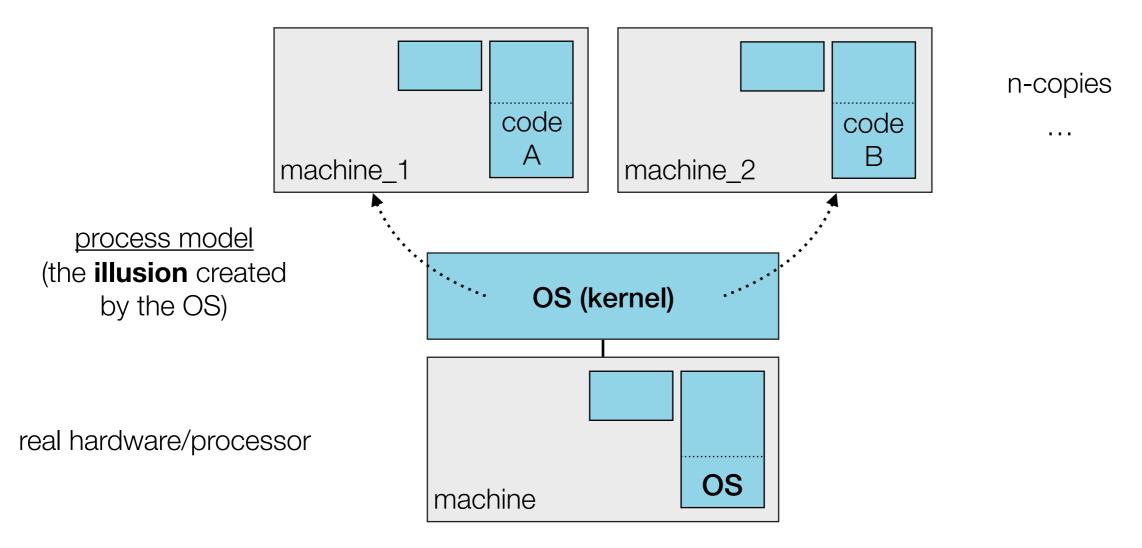
- Selfie (mipster) can not execute more than one program at a time. However it already provides some basic operating system functionality we will built upon.
- Step by step we will extend selfie to enable concurrent execution of programs.

First step:

- Selfie can execute two copies of the same binary concurrently.
 - we will solve the problem of one physical CPU
- To figure out how to implement this feature in selfie, we look at how an OS achieves this goal.

Concurrency in an OS

- The OS kernel creates n instances of the machine it runs on to enable concurrent execution of n processes
 - hardware virtualization (CPU and memory)
- Process is not aware that it runs in such a container.



Process Model

A process model describes the **illusion** that the operating system creates.

- several process models, which differ
 - in how close the illusion created by the OS is to the real machine
 - in the level of <u>temporal</u> and <u>spatial isolation</u> they provide
- system virtualization → an exact copy that is absolutely indistinguishable from the machine
- **2. UNIX process** \rightarrow a subset of the machine
- 3. threads \rightarrow an even smaller subset of the machine

1 CPU

"How could the OS actually execute two processes on a single CPU concurrently?"

"It could behave like mipster, it could interpret code.

The OS could interpret codeA for a while, then stop and interpret codeB for a while, then stop and interpret codeA for a while, then stop and..."

Concurrency in an OS

Interrupt and continue execution:

- It is necessary to save enough of the machines state/process at the point it is interrupted...
- ... so it can be restored when execution is continued some time later.

The actual purpose of a context:

 the minimal set of data saved so execution can be interrupted and continued

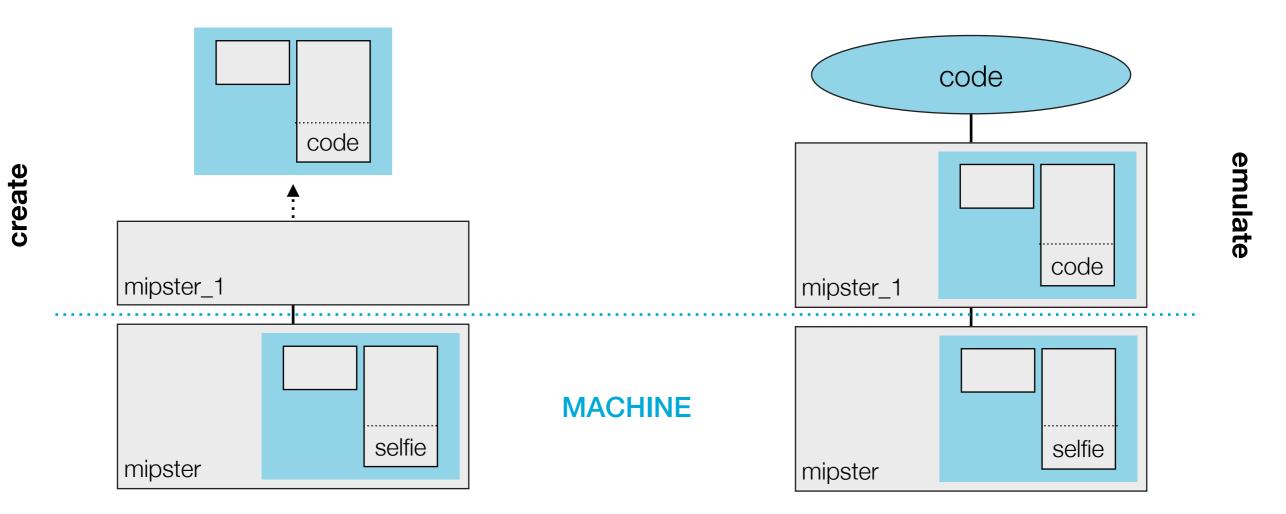
"The OS could interpret codeA for a while, then stop, save codeA's context and interpret codeB for a while, then stop, save codeB's context, restore codeA's context and interpret codeA for a while, then stop, save..."

Concurrency in Selfie

- The machine our OS will run on is the machine instance created by mipster.
 - → this means mipster can create an **instance** of that machine

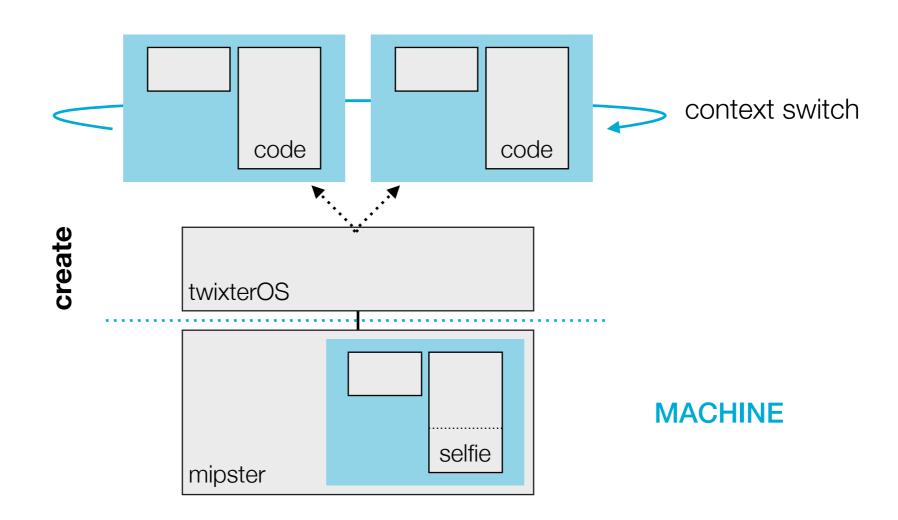
Consider this situation:

- Let our machine (mipster) execute selfie and start mipster, say mipster_1
- mipster_1 creates and emulates one instance of the machine it is run on!
 - we want create and emulate **two instances** of that machine → two contexts



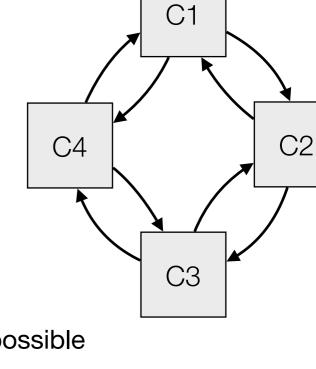
TwixterOS

- First approach towards building an OS will be a system similar to mipster.
 - implement -x option that creates a <u>multi-tasking</u> system (twixterOS) that runs two copies of the same binary → two processes
- Instead of creating and emulating one context, this system creates and emulates 2 contexts concurrently by switching between them after every instruction.

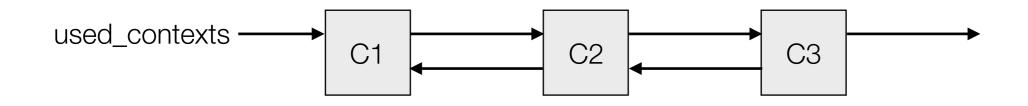


Context

- you can make yourself familiar with the context <u>structure</u> and procedures for managing multiple contexts in selfie
 - <u>creating</u> and <u>allocating</u> a new context
 - searching for a context
 - <u>deleting</u> and <u>freeing</u> a context
 - saving, restoring and caching a context (later)
- selfie maintains two lists to manage contexts
 - used contexts → doubly-linked list of used contexts,
 - free_contexts \rightarrow singly-linked of free contexts
- lists are just one possible way to organize/structure contexts
 - · any structure that can be built from parent, previous and next is possible



circled_contexts



TwixterOS Implementation Tips

- recognize -x option as argument and set up twixterOS (similar to mipster)
- copy or modify selfie_run(...)
 - implement selfie_run_twixterOS(...) that creates two contexts using selfie run(...) as blueprint
 - modify the existing selfie_run(...), such that a second context is created when twixterOS is run
- individual or linked contexts
 - no linking, just passing both contexts as arguments to the twixterOS procedure
 - link the contexts (each beeing the others next) and pass the first context is twixterOS

extra challenges

- make sure that both contexts finish execution with an exit call
- enable loading two or more different binaries and execute them concurrently

TwixterOS Testing

- have twixterOS execute <u>hello-world.c</u> → a program that prints "Hello World!" to the console
- the correct, yet not so nice looking output should be

```
> Hello Wo Hello World! rld!
```

- But why? look for an answer in hello-world.c
 - write(...) 8 bytes at a time
 - switch after one instruction

"Wait...It is rather unlikely that write() corresponds to a single machine instruction. Why is the output not 'HHeelloo...'"

"This is another example of the OS support already provided by mipster and therefore also twixterOS. Let us look for the answer in the code.

Try to find the definition of the write procedure in selfie"

Write in Selfie

- there is no procedure definition for write in the traditional sense
 - only emit_write, implement_write and a declaration of write
- a closer look at emit_write shows that a library table entry for the write procedure is created, followed by its implementation (machine instructions we come back to this soon)
 - part of that implementation is SYSCALL WRITE
- a syscall causes twixterOS to stop interpreting code and return from mipster_switch (as described when mipster was introduced)
- ▶ a syscall is handled by twixterOS → <u>handle exception</u>, handle system call
 - twixterOS <u>implements</u> the syscall, not the process

Write in Selfie

- the process does not write to the console directly
- instead, the process uses a mechanism to have twixterOS execute write on its behalf
- an OS uses this mechanism to solve a problem that we created because we wanted to execute programs concurrently

The Problem

We created two or more processes that are unaware

- that they run in a container
- of each other

and that use machine resources.

- Processes that write to the console or open a file at the same time would cause undefined behavior.
- Therefore, processes can not be allowed to perform certain operations. Instead they
 ask the OS to perform these critical operations for them.
- Solution the OS controls execution and manage resources
 - the OS provides services
 - the process may request those services via system calls
 - process and OS run in different modes

System Calls

Modes and Protection,
API and ABI,
Wrapper Function and Trap,
Bootstrapping

Modes Protection

- Modes with different privileges provide a means of protection and isolation
- CPU has a bit (or more) that is set to its current privilege level (mode).
 Code executed on this CPU is running in this mode.
- An OS is run in kernel mode, it has all the privileges → trusted
 - unrestricted access to hardware and memory
 - allowed to execute any instruction
- A program is run in user mode, it has the least privileges → untrusted
 - not allowed to access hardware and memory directly
 - not allowed to execute certain instructions
 - not allowed to access OS code

Modes Protection

- If a program executes an instruction it is not allowed to (program's privileges level too low), an exception is raised.
- This exception prompts the processor to switch into kernel-mode and execute the OS code that handles the exception.

- Selfie does not implement actual mode switching as described.
 - nothing that indicates the current mode

System Call

System calls are **requests** made by processes **for services** provided by the operating system.

- operating system performs the requested services and returns control to the program
- a switch from user mode into kernel mode takes place

System Call

- ► The OS creates an **abstraction** of the services by providing an **API** that specifies available functions and their parameters and return values.
 - no direct access to actual system call
 - portability and ease of use same API across different systems
- The API is not accessed directly but via a library provided by the OS
 - "not directly": the program does not jump to the function within the OS directly
 - library is the interface to system calls
 - library provides wrapper functions for calls to API
 - wrapper functions conform to ABI

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t count);
```

API and **ABI**

application programming interface - API

- communicate between two pieces of software on source code level
 - relatively hardware-independent
 - exposed parts of software that can be accessed from outside
- abstraction of underlying implementation → provide building blocks to programmer
- API is specification (behavior), library its implementation
- for OS: API describes interface between application and OS (ex. POSIX)

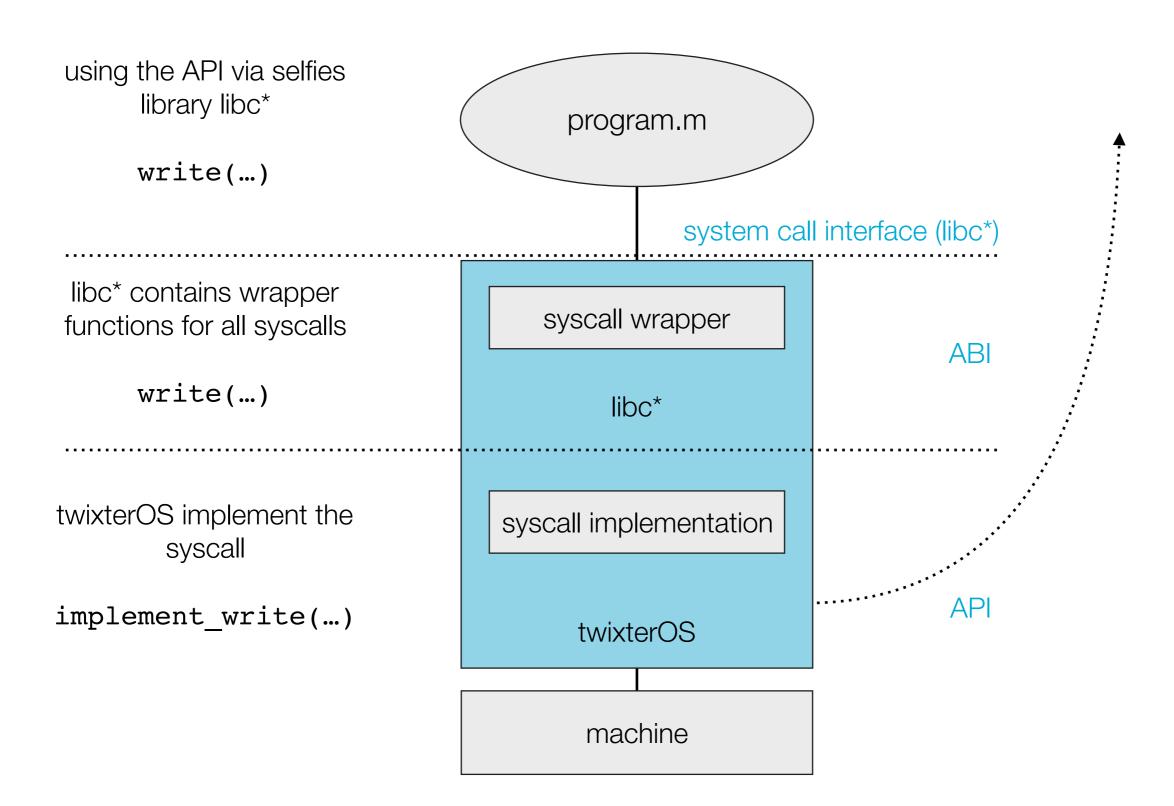
application binary interface - ABI

- communicate between two binary programs
 - hardware-dependent
- conforming to ABI is mostly done by compiler, the OS, a library author,...
- · describes calling convention (passing parameters), syscall interface,...

System Call in Selfie

- Selfie already supports 5 system calls (API)
 - exit, read, write, open and malloc
 - twixterOS has a implementation/definition for each of them
- A program that calls write does **not jump** to the implementation of write in twixterOS.
 - it is not allowed to do so → potential for malicious behavior
- Instead selfie provides an interface as part of the library libc*
 - contains wrapper functions for all system calls (ABI)
 - program jumps to these wrapper functions when executing a system call

System Call in Selfie



Wrapper Function

A <u>wrapper function</u> "wraps" the call to another function and usually performs little additional computation. They are often used to **hide details** and create an extra layer of **abstraction**.

- ▶ wrapper functions are put into the binary at compile time by selfie → emit_write
- prepare actual syscall conform to syscall ABI
 - calling convention
 - copy arguments from stack into registers in which twixterOS expects them to be
 - put unique syscall number into register
- generate a trap, a software generated interrupt to transfer control from process to OS
- interface for syscalls and provide protection to the OS kernel

Trapping Mechanism in Selfie

- ► syscalls and errors throw an exception → throw exception
 - set the exception code
 - set the trap (flag)
- the trap signals the emulator to stop interpreting code and handle the exception
 - transfer control from program to emulator (switch from user to kernel mode)
- after handling the trap, control is returned to the process

"So wrapper functions are put into the binary at compile time by selfie..... What about the gcc compiled version of selfie?"

"That is one missing piece we haven't talked about yet.

The other, as you might have noticed, is that the actual implementation of a syscall in mipster/twixterOS contains that same syscall. (implement_write contains write)".

Bootstrapping Syscalls on Bootlevel 0

- syscall wrapper the same principle applies
 - the compiler (gcc) puts the wrapper code into the selfie binary
- ► How?
 - the compiler sees undefined procedures (the declaration as mentioned before)
 - therefore, the compiler provides the implementation
- Therefore, syscalls on bootlevel 0 are handled by the actual OS running on your machine.

Bootstrapping Syscalls on Bootlevel 0

- Notice how syscalls are passed down?
- Therefore, all syscalls
 reach bootlevel 0 and are
 handled by the actual
 OS running on your
 machine.
- The same way syscalls are passed down, return values are passed up to the program.

```
write(...);
                                                  program.m
syscall
    implement_write(...) {
        write(...);
                                                    mipster
syscall
    implement_write(...) {
        write(...);
                                                    mipster
 syscall
                                                      OS
       OS implements write
                                                 x86 machine
```

Summary OS Introduction

- Being able to run several processes at the same time is a huge gain.
- Unfortunately, there are problems that have to be solved to enable concurrent execution.
- The operating system provides solutions to these problems.

Summary OS Introduction

- We addressed the first problem
 - several processes using the same CPU (and console)
- the OS solves this problem by managing CPU and controlling execution
 - CPU time is shared among processes → time-sharing
 - OS provides services to the process → system call
- concepts and mechanisms used
 - process → program in execution
 - abstraction/illusion → creating machine instances (containers), processes, system call API (hide hardware details), wrapper functions (conform to ABI)
 - different modes → kernel mode for OS and user mode for processes
 - trap → software interrupt to switch between mods

TwixterOS Testing Revisited

- Remember this "problem"?
 - Hello Wo Hello World! rld!
- writing has to be <u>synchronized</u> (coordinated)
- BUT processes are unaware of each other and they have no way of communicating directly with each other
 - the process can not be responsible for checking if it is safe to write
 - OS support is needed
- we can solve this problem with help of the twixterOS by implementing a mechanism called **locking**

A Write Lock Intended Semantics

a variable indicating lock owner

lock call saves caller as lock owner (acquire lock)

unlock call removes caller as lock owner (release lock)

```
uint64 t LOCK;
lock();
    while(*foo != 0) {
        write(1, foo, 8);
        foo = foo + 1;
unlock();
```

iff the lock is held by a process, only the lock owner is allowed to **write** to the console

A Write Lock Implementation Tips

- ▶ uint64 t LOCK is a global variable within the twixterOS
- a process is not allowed to set this variable itself
- the operating has to provide this as a service
 - a lock and unlock system call → API: lock(), unlock(), ABI: syscall number
 - a libc* wrapper function → syscall interface (ABI)
 - handling the syscalls → implementation in twixterOS
- possible pitfalls
 - unlike the other syscalls, lock is not 'passed down'
 - a lock can only be acquired when it is not held (syscall successful)
 - otherwise the process has to wait (syscall failed)
 - the OS sets the program counter back so the process makes the syscall again later
 - only the lock owner can unlock

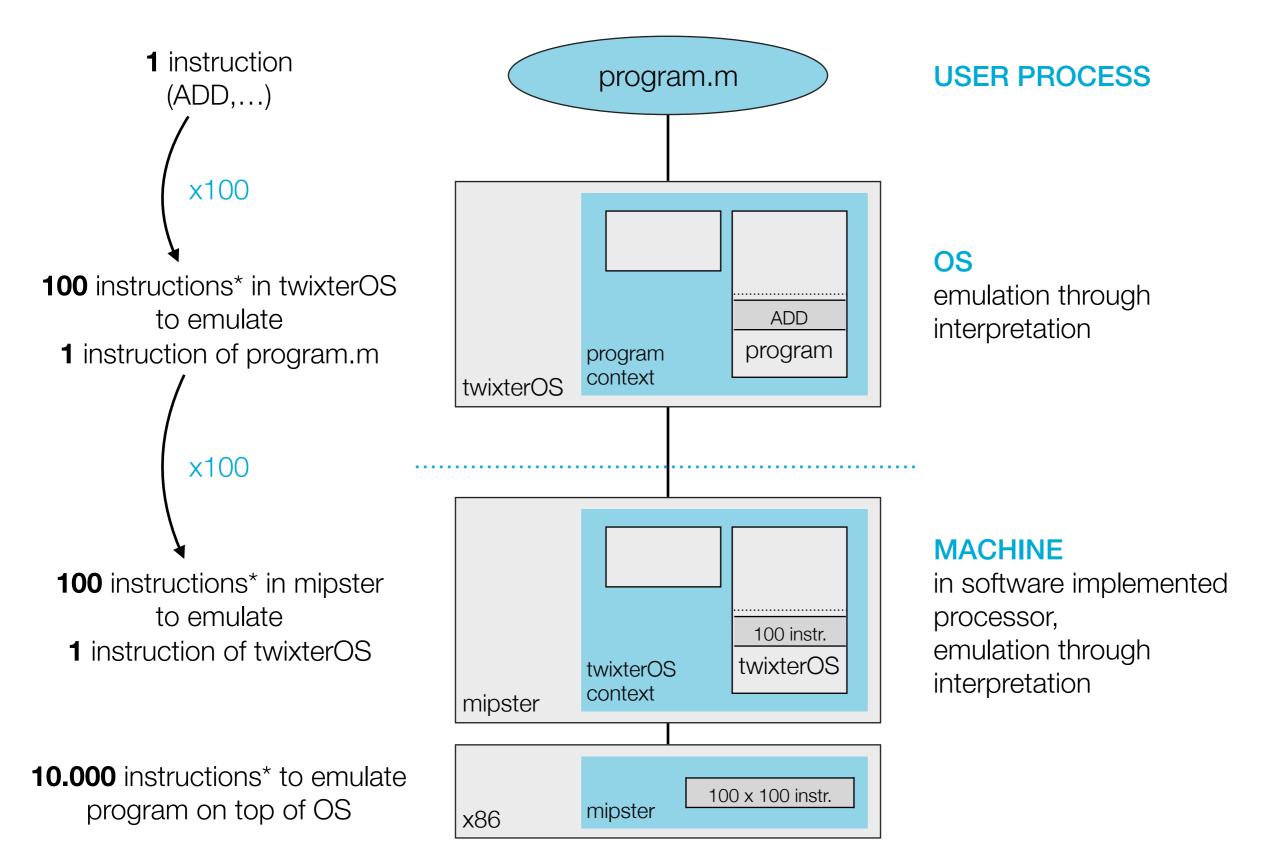
Emulation vs. Virtualization

Operating System

An operating system **solves problems** that arise from concurrent execution of processes. It manages resources and **controls the execution** of many programs on a machine and **abstracts** that machine.

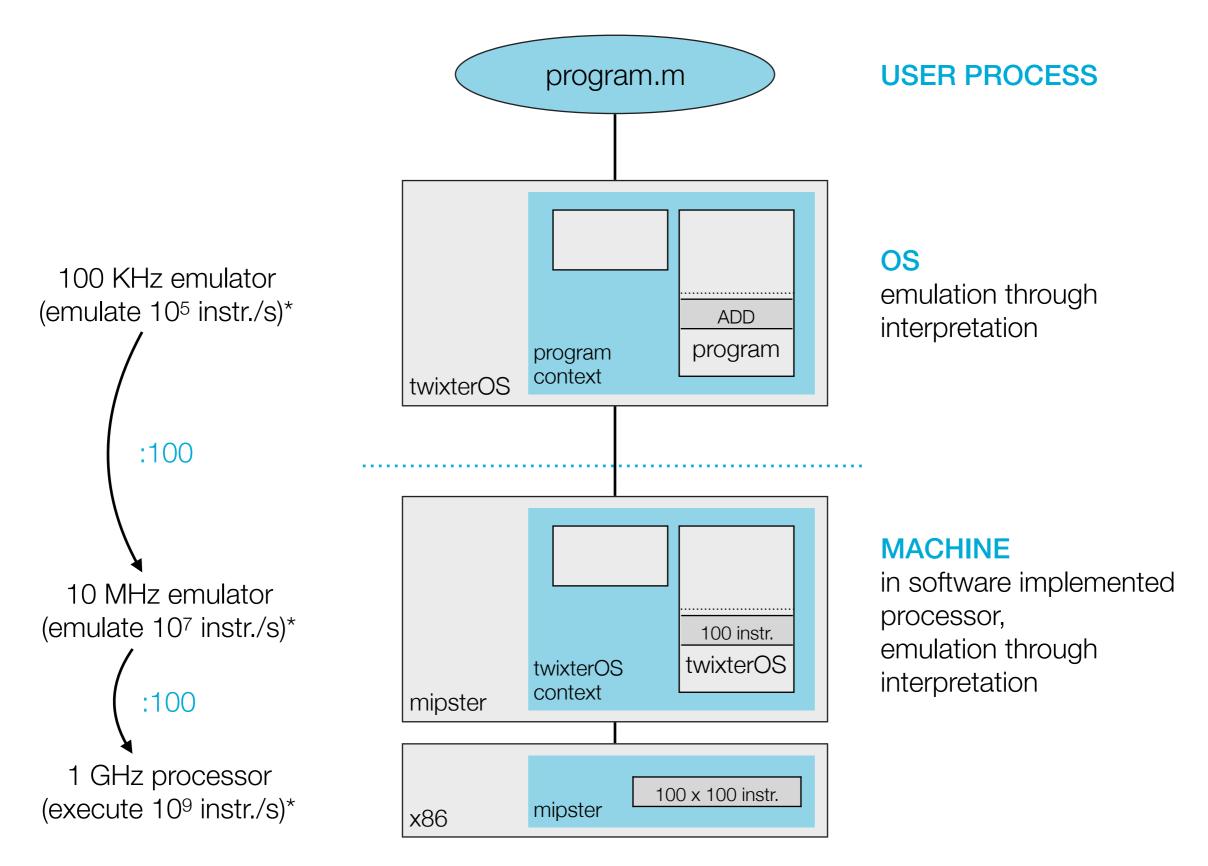
- already a good definition
- BUT twixterOS does not control execution like a 'real' OS kernel
 - controlling execution by interpretation is highly insufficient (exponential in #emulators)
 - more efficient if the program would run directly on the machine

Interpretation is Slow



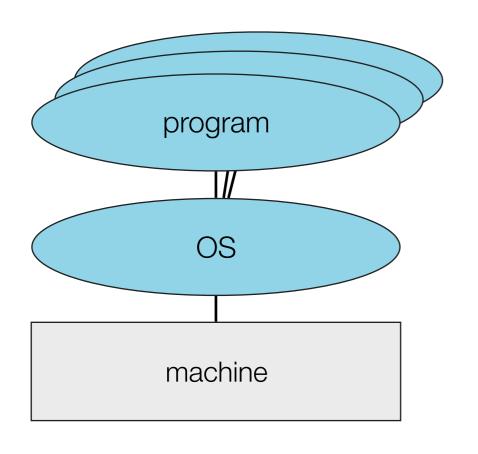
*assumption: 100 instructions to emulate 1

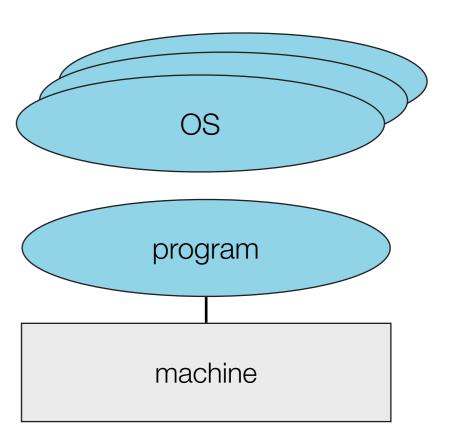
Interpretation is Slow



Interpretation is Slow Solution

- An OS uses a mechanism to allow user processes to execute directly on the machine the OS itself runs on
 - the OS <u>virtualizes</u> the CPU, it 'asks' the machine(CPU) to execute the program on its behalf
- To achieve this, the OS has to give up the machine give up control?





interpretation is slow

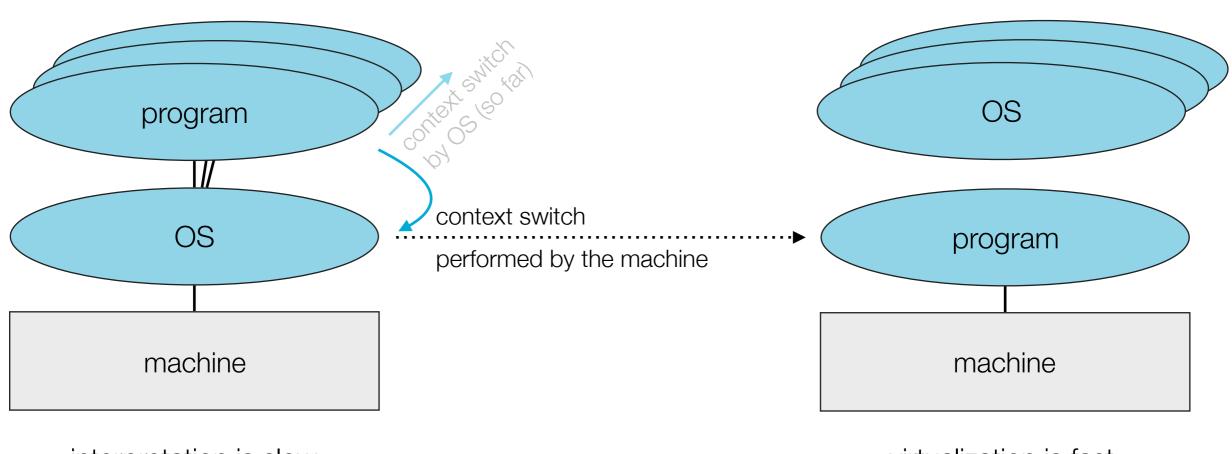
virtualization is fast

CPU Virtualization

- With this concept some questions arise
 - control
 - How does the OS give up the machine (and stay in control)?
 - system calls
 - Are they different now, how does it work?
 - switching between processes
 - Is this different now, how does it work?
- What we discussed so far so still holds. We only left out some details.
- Whenever the OS is needed it regains control of the machine
 - with support of the machine

CPU Virtualization Pass-Over Control

- ► The mechanism used to virtualize the CPU is a context switch
 - initiated by OS
 - for actual switch hardware support necessary



interpretation is slow

virtualization is fast

CPU Virtualization Regain Control

- two methods possible
 - cooperative
 - process gives up CPU voluntarily via syscall (cooperative)
 - syscall → switch to OS
 - preemptive
 - process is forced to give up CPU
 - before giving up the CPU, the OS sets a timer
 - timer causes interrupt → switch to OS

CPU Virtualization Syscall

- ► Emulation → emulator(OS) interprets syscall instruction of program
 - trap into OS
 - trap signals the OS to stop interpreting and handle the syscall
- ▶ Virtualization → machine executes syscall instruction of program
 - trap into machine
 - trap causes the machine to perform a context switch to the OS
 - machine cannot handle syscalls
 - the OS then handles the syscall (and afterwards switches back to the process)

CPU Virtualization Switching

- Whenever the OS regains control it decides which process to execute next (performing context switch)
 - Which next? → we talk about process management soon
 - this switch is performed by the OS (software)

"I understand the general idea but implementing this seems rather complicated and like a lot of work..."

"Yes...getting this to work requires quite a lot of thinking. There are many details that have to be done right.

Fortunately, selfie already implements this mechanism. It is exactly how selfies hypervisor hypster executes a single program.

We can study this mechanism without having to implement it."

Hypster

Hosting, Context Switch

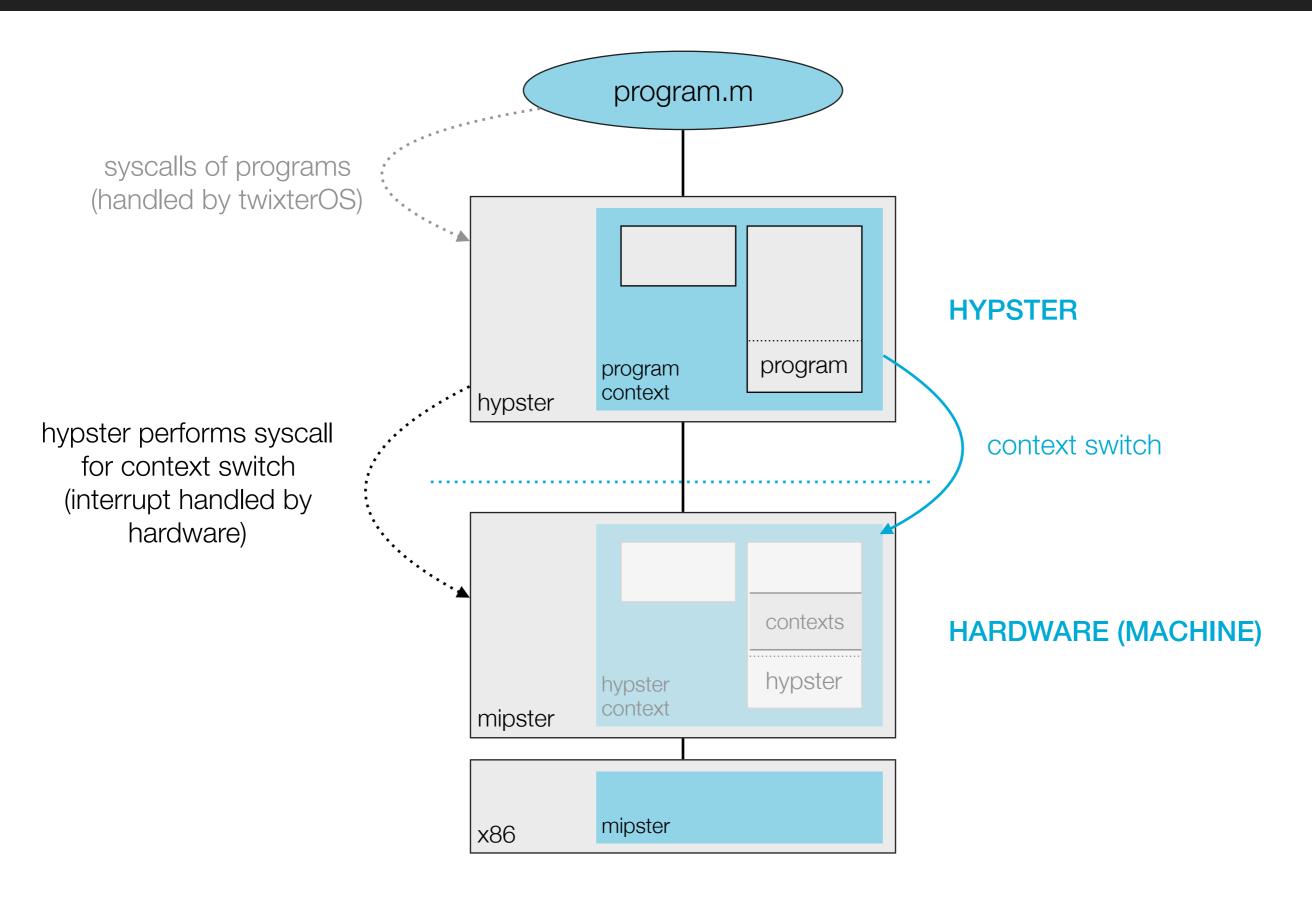
Hypster Hypervisor

- a <u>hypervisor</u> or virtual machine monitor (VMM) creates a virtual machine
- hosts code
 - does not execute code itself
 - uses the interpreter it runs on (has to run on a mipster)
- comparing hypster and mipster
 - hypster_switch instead of mipster_switch
 - no get_parent(from_context) != MY_CONTEXT (soon you will understand why)

Context Switch Hypster vs Mipster

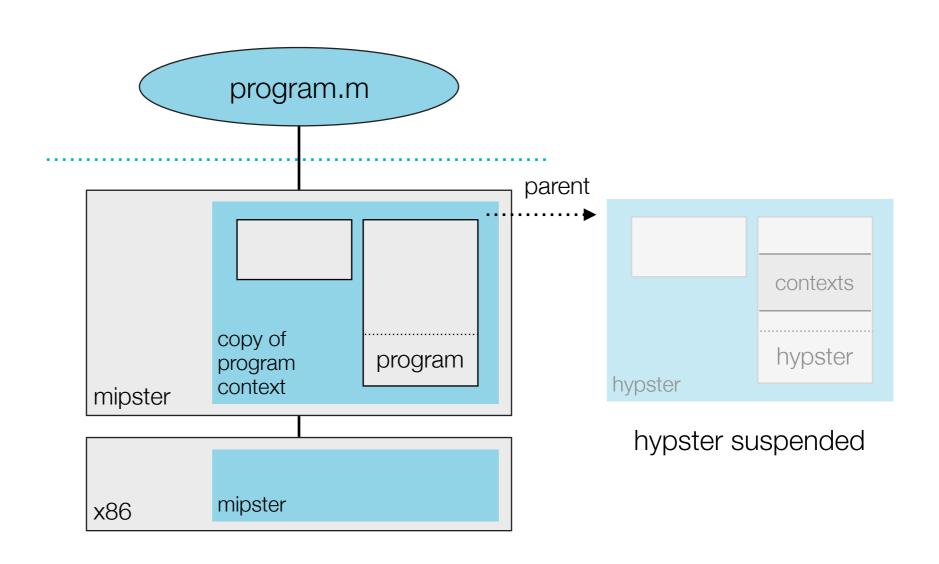
- mipster switch mipster and twixterOS
 - context switch performed by mipster/twixterOS
 - switching implemented in software
- hypster switch hypster
 - syscall by hypster
 - context switch performed by the machine hypster is running on
 - switching implemented in hardware

Context Switch Hypster → Program

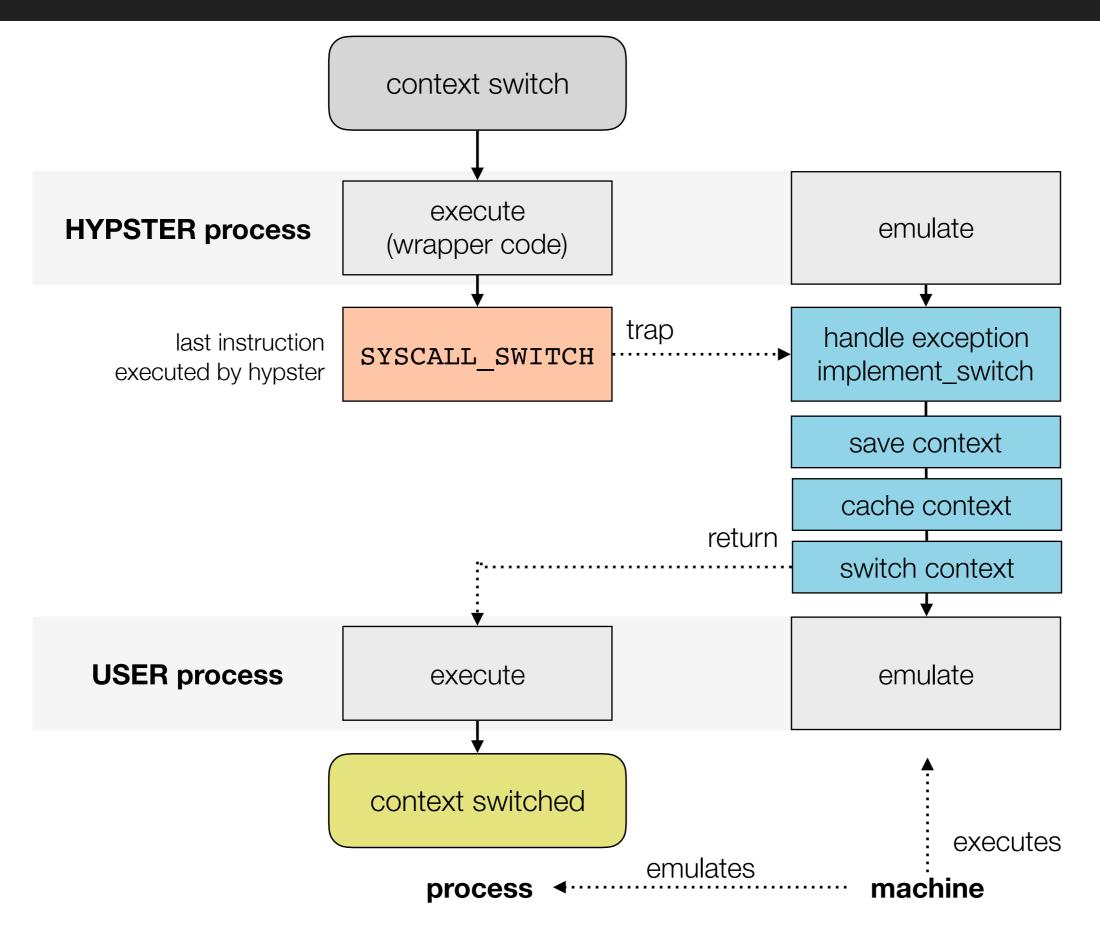


Context Switch Hypster → Program

program running on hardware natively



CPU Virtualization Pass-Over Control in Selfie



CPU Virtualization Pass-Over Control in Selfie

- we have mipster emulating hypster and hypster hosting program
- hypster runs on machine
 - hypster_switch jumps to the code of emit_switch(syscall wrapper)
 - because we search for definition in library table first (compile time)
 - no definition in library table on bootlevel 0 → search global table (fall back to mipster_switch)
- machine implementing/handeling the SYSCALL_SWITCH
 - saves the current context (hypster)
 - switches to the cached context (program)
 - continues with emulation (now program)

Cached Context Abstraction

- cached context is simply a deep copy of a context that the machine creates
 - hypster manages the original context
 - the machine uses the copy for execution
- only cached contexts are used for direct execution on machine
- orignial and copy are synchronized before and after a context switch
 - machine updates the copy before the switch → restore_context
 - hypster modified context (ex. by handling a syscall for the process)
 - machine updates the original after the switch → save context
 - machine executed and thereby modified context

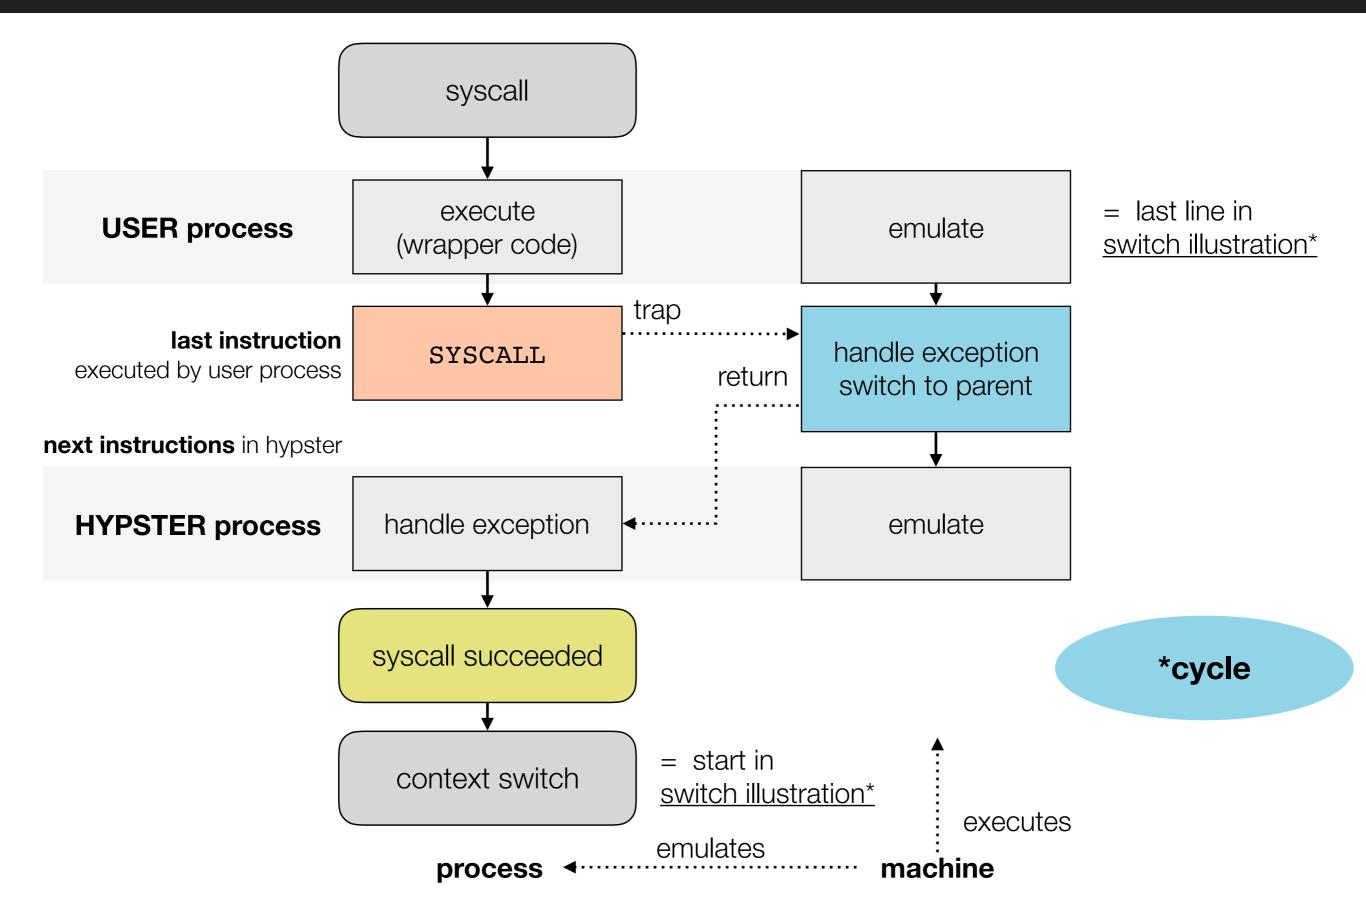
Cached Context Parent

- the parent of a context is the machine the context was created for
 - the original contexts parent is the hypster context (MY_CONTEXT(0) to hypster)
 - the cached contexts parent is also the hypster context
 (cannot be MY_CONTEXT(0) to the machine)
- So how does the machine know who the cached contexts parent is?
 - the parent initiates the switch to execute its context
 - machine knows which context executed syscall switch
 - the context it was executing (current context)

CPU Virtualization Syscall in Selfie

- only a contexts parent can hande syscalls
- ► We add a little detail to the <u>trapping mechanism</u> in selfie:
- 1. program runs on machine and makes syscall → trap
- 2. trap causes the machine to perform context switch to hypster
 - machine cannot handle syscall → machine is **not** the contexts **parent**
 - switch to parent → get_parent(from_context != MY_CONTEXT)
 - syscall/exception information remains in context
- 3. hypster runs on the machine again
 - machine executes next instructions of hypster code
 - the instructions after the switch syscall in wrapper function (return from hypster_switch and handle exception)
- 4. after handeling the exception hypster switches back to process

CPU Virtualization Syscall in Selfie



get_parent(from_context != MY_CONTEXT)

- only a contexts parent can handle that contexts syscalls
- machine can execute a context that is not its own
 - from_context could be a cached context executed on behalf of hypster
 - therefore the machine checks the contexts parent
- hypster does not execute any context
 - hypster only handles the syscalls of its own context
 - from_context can never be a context that is not hypsters context
 - therefore hypster does not check the contexts parent

Summary Emulation vs. Virtualization

- ▶ Problem → interpreting code is slow, hosting code is fast
 - take out the middle man and execute directly on hardware
- the OS virtualizes the CPU
 - OS gives up control → context switch
 - has mechanisms to regain control → cooperative vs. preemptive
- mipster switch and hypster switch
 - switch implemented in software
 - switch with hardware support

Memory Management

Address Space, Memory Virtualization

Operating Systems

An operating system **solves problems** that arise from concurrent execution of processes. It **manages resources** and controls the execution of many programs on a machine and **abstracts** that machine.

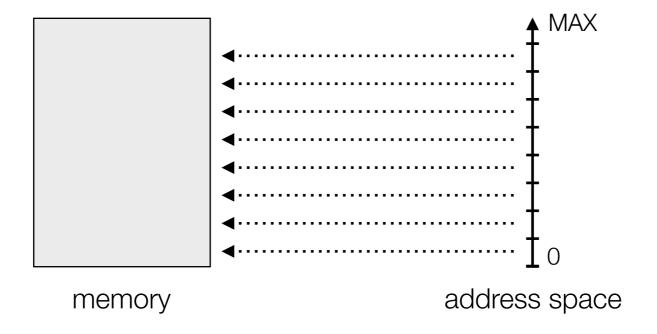
- Selfie can now execute two or more programs at the same time efficiently
 - the problem of having only one physical CPU was solved by virtualizing the CPU using a time-sharing approach
 - the challenge of sharing one physical memory was already solved by the operating system support implemented in selfie
- Nevertheless, we will discuss how an OS manages one physical memory and how see how this is implemented in selfie.

1 Memory

Address Space

An <u>address space</u> is a range of **memory addresses.** It is an **abstraction** of physical memory created by the OS.

- a contiguous block of numbers (<u>addresses</u>) ordered from low to high
- addresses are necessary to locate pieces of data in memory
 - each address uniquely identifies a piece of data



Memory Management Problem?

- The challenge with one physical memory
 - memory is fixed-sized
 - memory has to be shared among processes

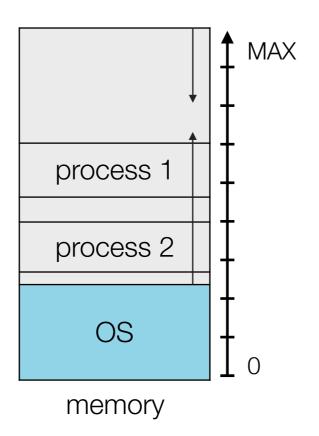
"How could the OS actually execute two processes with one memory concurrently?"

"We said earlier that memory is part of a processes context. This means that the memory gets switched when a context switch occurs.

Does this mean that processes (contexts) are stored somewhere on disk when they are not executing?"

Memory Management Shared Memory

- Possible solution
 - only the running process is in main memory, others reside on disk
 - huge drawback → context switches are slow (memory to disk)
- Better solution
 - keep processes in memory and switch between them



Memory Management Shared Memory

- keeping several problems in memory creates new problems
 - portability
 - → a process should run independent of where it is loaded into memory
 - protection
 - → a process should not access memory outside its 'own' memory
 - transparency
 - ightarrow a process should be unaware of the fact that memory is shared memory
- all the above boil down to one issue addressing

Memory Management Addressing

portability

 addresses are needed at compile time, but are unknown before program is loaded into memory

protection

no access to addresses outside processes 'own' memory

transparency

process believes it can access every address from 0 to MAX

- creating a layer of abstraction solves these problems
 - → virtualizing memory

Memory Virtualization

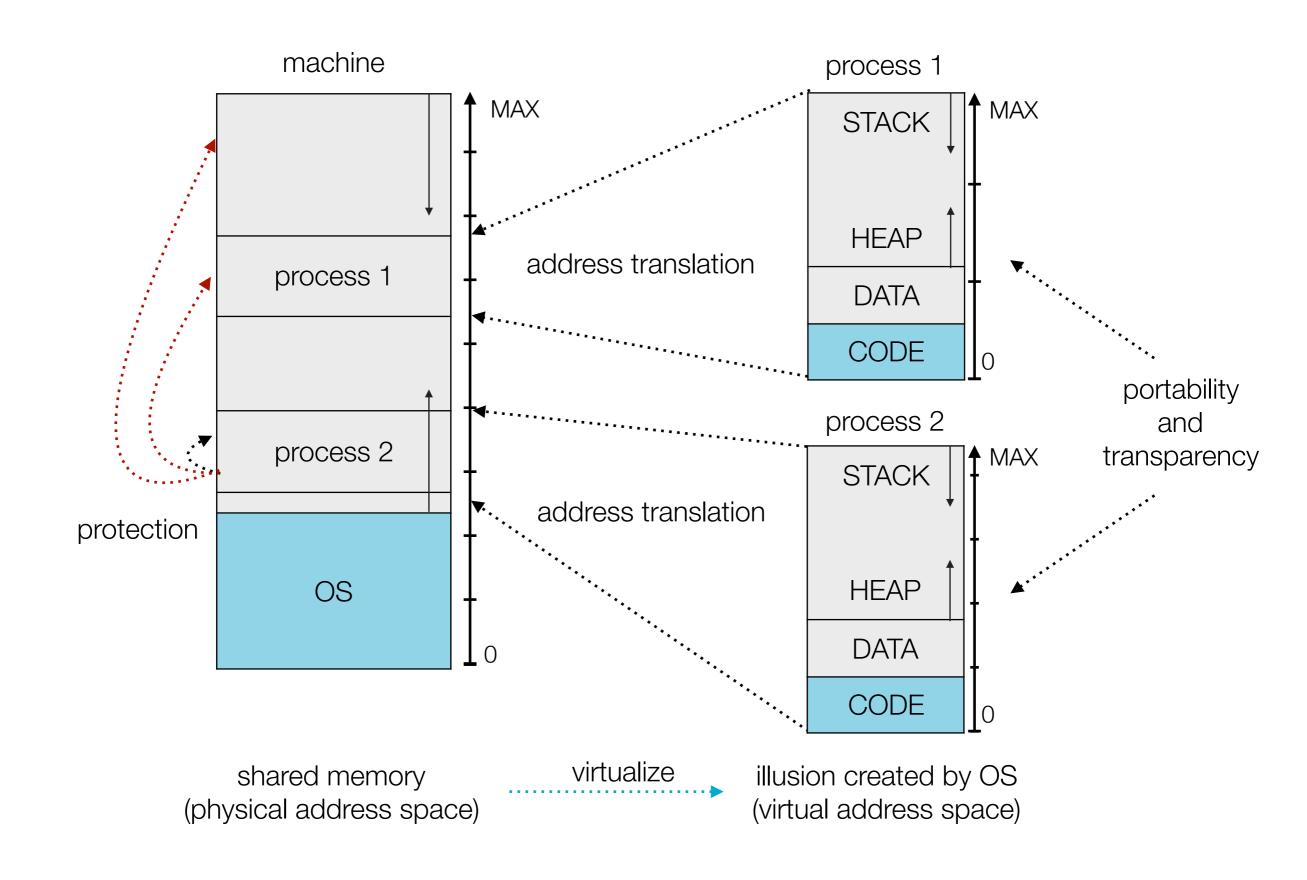
- OS abstracts physical address space using a technique called address translation
- OS provides each process with virtual memory, an illusion of private memory
 - process uses <u>virtual addresses</u>
 (code compiles with virtual addresses)
 - OS assigns physical memory to virtual memory
 - hardware with OS support translates virtual addresses to physical addresses where data is located
- portability and transparency
 - each process is provided the same illusion (same memory layout and virtual address space 0 to MAX)
- protection
 - depends on how OS virtualizes memory

Address Translation

Address Translation is the process of finding the physical address for a given virtual address.

- efficient address translation performed by hardware
 - by part CPU often referred to as memory managing unit(MMU)
 - raises exception/trap when process attempts illegal access
- supported by OS
 - manages memory (assignment, free?)
 - set up hardware so addresses are mapped correctly
 - handles exception raised by hardware (illegal access)
- How addresses are translated depends on memory virtuaiztion technique used by OS

Memory Virtualization



Memory Virtualization Tequniqes

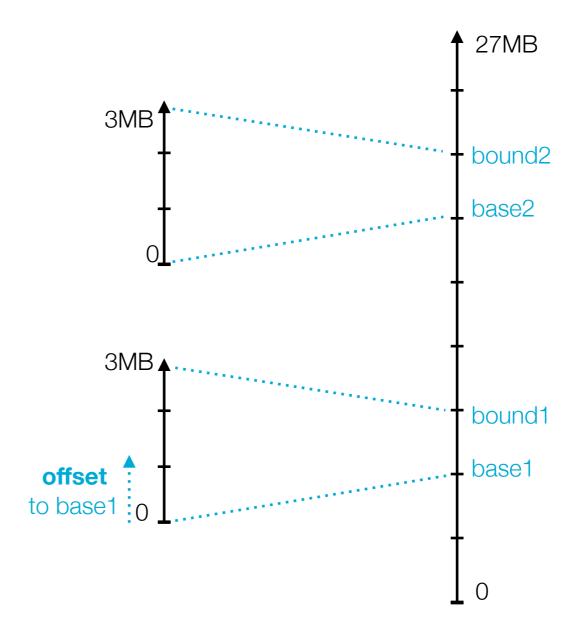
- The OS can use different techniques to virtualize memory
 - base and bound, segmentation
 - paging

- The OS assigns each process a contiguous block(chunk) of physical addresses
 - #assigned addresses = #virtual addresses
- It remembers for each process the lowest and highest address of this block → a base and bound
 - the range of addresses the process can access

- Address translation
 - virtual address is the offset to the base in physical address space

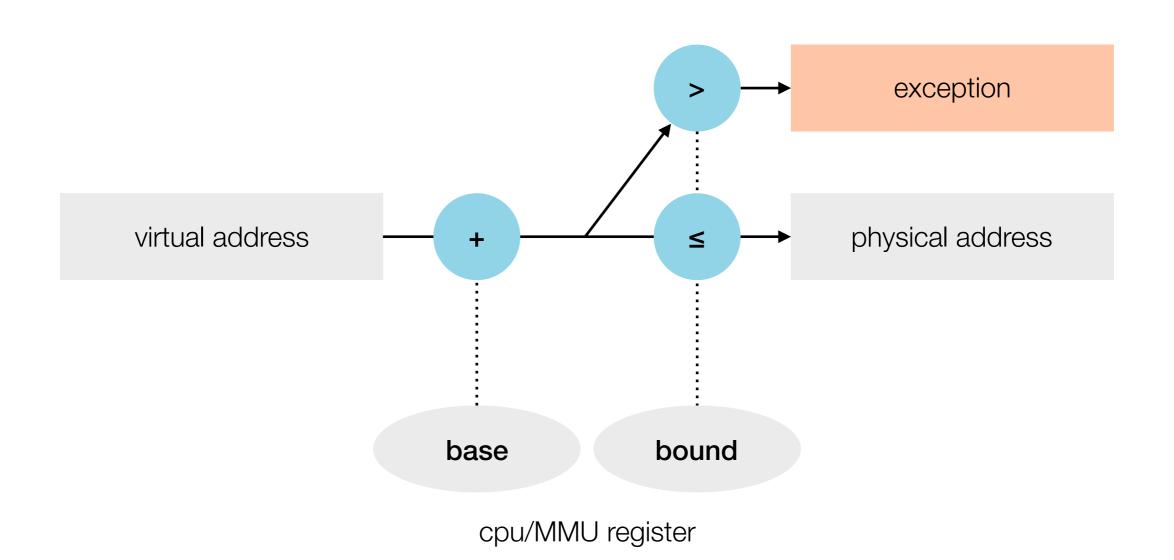
process 1	base1	bound1
process 2	base2	bound2

data structure to store base and bound information





uses vaddr as offset to base



Advantages

address translation is simple and fast

Limitations

- size and number of processes in memory is hardware specific
- potentially a lot of unused memory
 - internal fragmentation gap between heap and stack (improved by segmentation)
 - external fragmentation find contiguous chunk

Memory Virtualization Segmentation

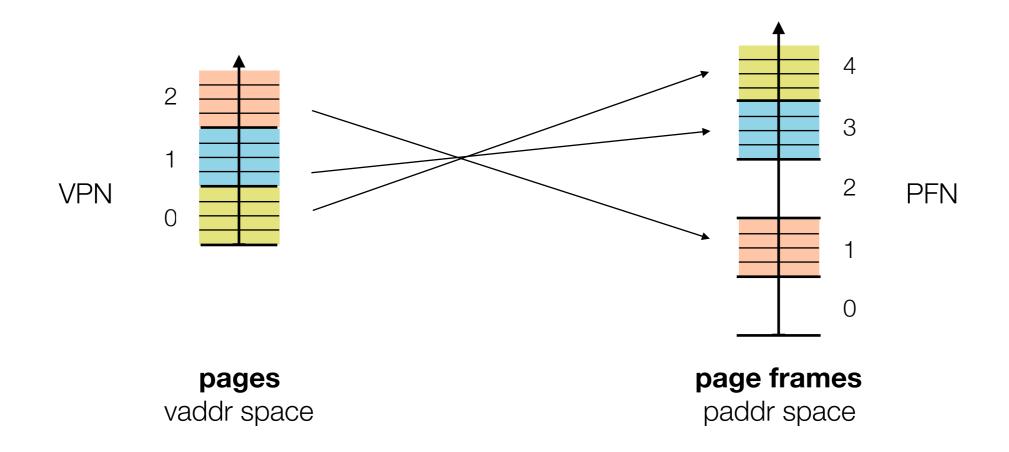
- similar to base and bound
- virtual memory divided into variable-sized logical segments
 - code, data, heap, stack
 - base and bound for each segment
 - variable-sized → external fragmentation
- less internal fragmentation

"Why not chop virtual memory in fixed-sized pieces, wouldn't that help with external fragmentation?"

"Yes, it would. This is exactly the idea behind paging"

Memory Virtualization Paging

- partition address space into fixed-sized chunks
 - virtual address space into pages → identified by virtual page number VPN
 - physical address space into page frames → identified by physical frame number PFN
- ▶ pages and page frame have same size → 4KB worth of addresses
- each virtual page maps to a physical page frame
 - to translate addresses these mappings are simply remembered in page table



Paging Page Table

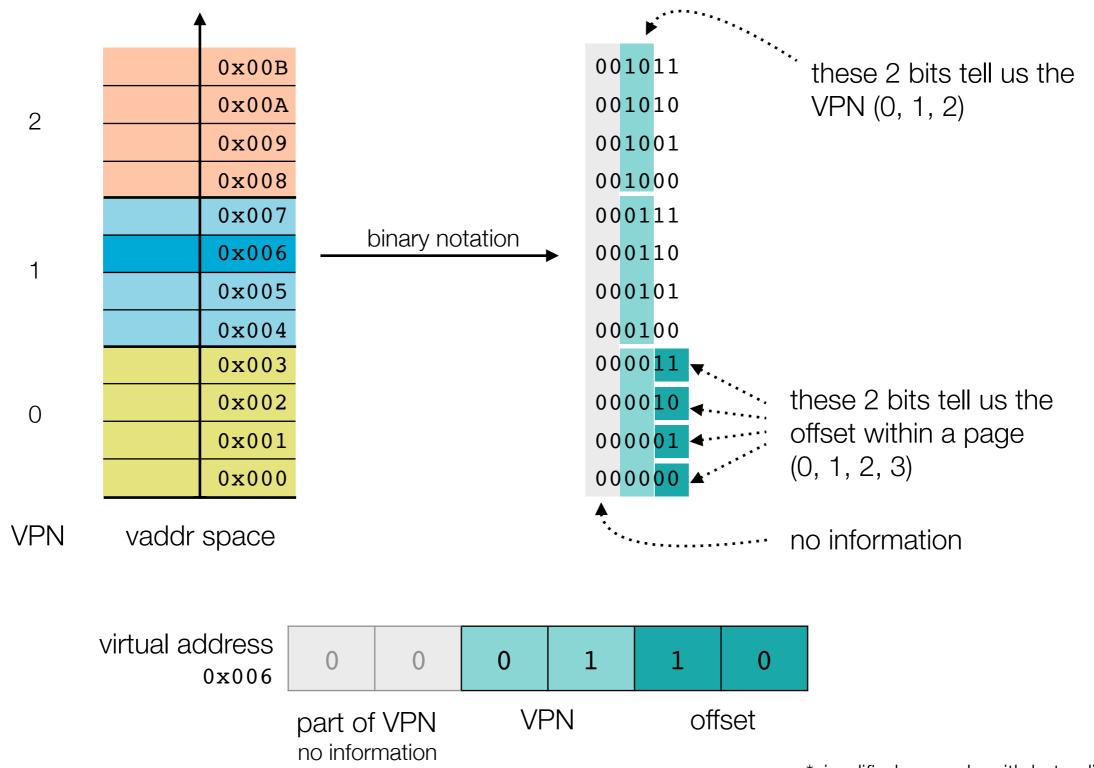
A <u>page table</u> is the data structure used for **address translation**. It stores the mapping from **VPN to PFN** for each process.

- OS maintains one page table per process
- simple implementation
 - array lookup
 - indexed with VPN to find PFN

0	4
1	3
2	1
VPN	PFN

Virtual Address Logical Point of View

there is more to an vaddr than you might think → take a close look at its binary notation*



Virtual Address Do the Math

- to uniquely identify x elements it needs [log₂(x)] bits
 - 4 bytes → 2 bit, 3 pages → 2 bit
- therefore address size (#bits) puts an upper bound on memory size
 - 32-bit cannot address more than 2³² byte of memory
- Using 32-bit virtual address and 4KB pages and page frames (selfie)
 - 12 bit are needed for the offset
 - 20 bit remain for VPN (max. of 2²⁰ pages possible)

Paging Page Table

- size of one page table entry
 - necessary/minimum = #bits needed to identify PFN (depends on size of physical memory)
 - allocate more and use remaining bits to store additional information (dirty bit, present bit,...)
- size of page table (selfie) 4GB of virtual memory
 - #virtual pages * sizeof(page table entry)
 - 8MB (2²⁰ * 8 byte)
 - huge therefore stored in memory
 - requires additional access for every translation

Paging Address Translation

- performed by hardware (MMU) with support of OS
- only VPN gets <u>translated</u>, offset is the same for page and page frame

- 1. **get VPN** from virtual address
 - get_page_of_virtual_address
- 2. **lookup PFN** in page table using VPN as index
 - get frame for page
- 3. **build** physical address
 - <u>calculate</u> using PFN and offset of virtual address

Memory Virtualization Paging

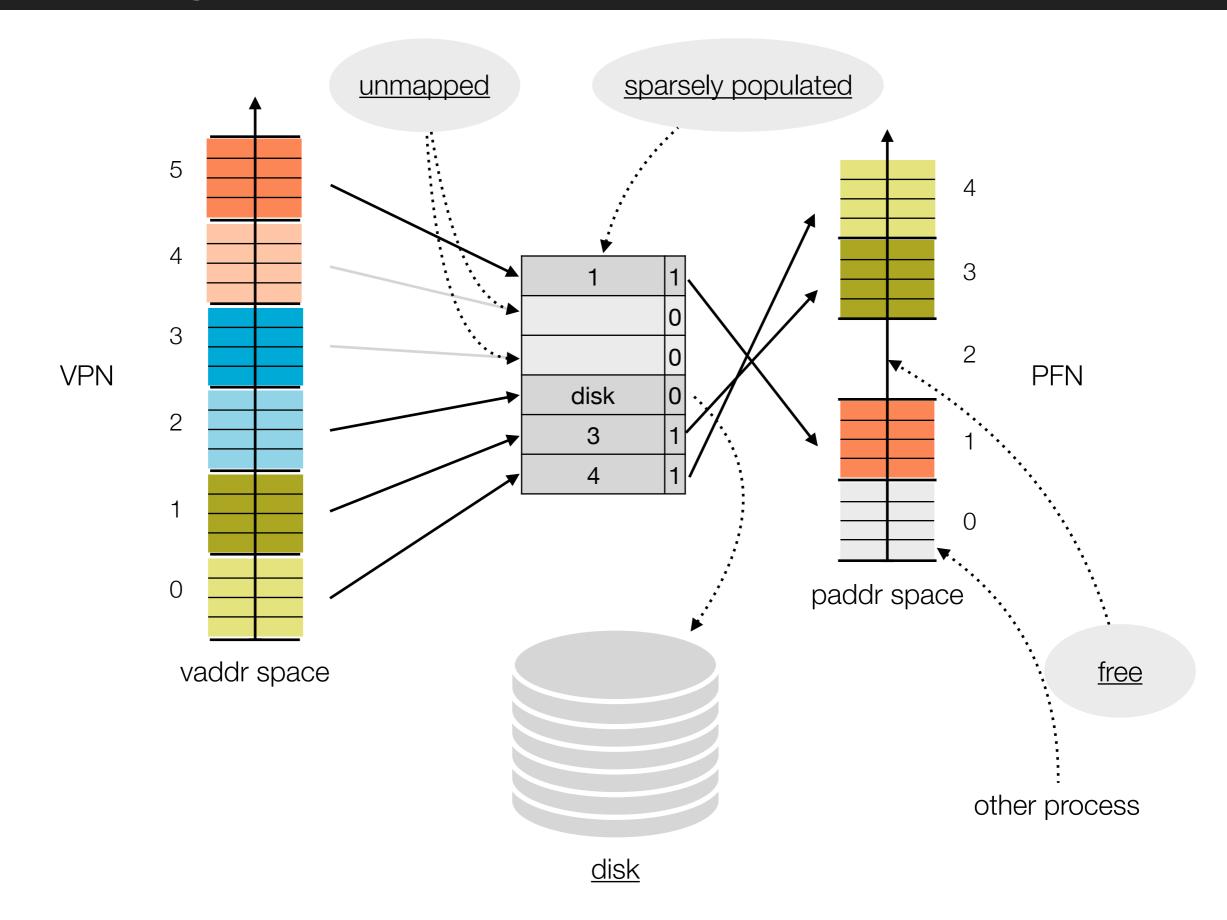
Limitations

- additional memory to store page table
- additional memory access for lookup

Advantages

- managing free memory is easy, external fragmentation is not a problem
- paging is a powerful virtualization technique
 - completely abstracts physical memory (layout and size)
 - virtual memory is not limited by physical memory size (temporarily move pages to disk when not needed)

Memory Virtualization Paging



"A large virtual address space and a huge page table...isn't that a huge waste of memory and disk space?"

"Yes, it most likely is. (unless the process would really need all that memory)

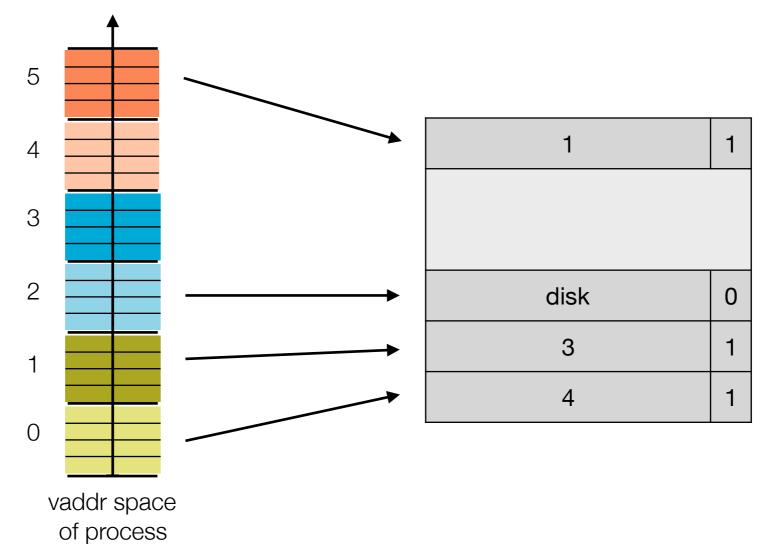
We can get around this easily by not mapping the entire virtual address space. Demand paging is one possibly way of doing so.

Paging on Demand

- not the entire virtual address space is mapped when a process is created
- only pages the process accesses are mapped
 - at process creation different options
 - maybe code segment, data segment, first page of stack (arguments)
 - or nothing is mapped
 - memory the process allocates is not mapped until accessed (lazy loading)
 - might be never → no consumption of physical memory
 - this also applies to the page table → sparsely populated
- when a process attempts to access an unmapped page an exception is raised
 - OS intervenes → allocates a new/free page frame and stores mapping in page table
 - problem → which frames can be used?

Paging on Demand

- page table is allocated in address space of OS
- page table is sparsely populated
 - gap between stack and heap
 - because of how virtual memory is laid out
- therefore page table itself is not entirely mapped



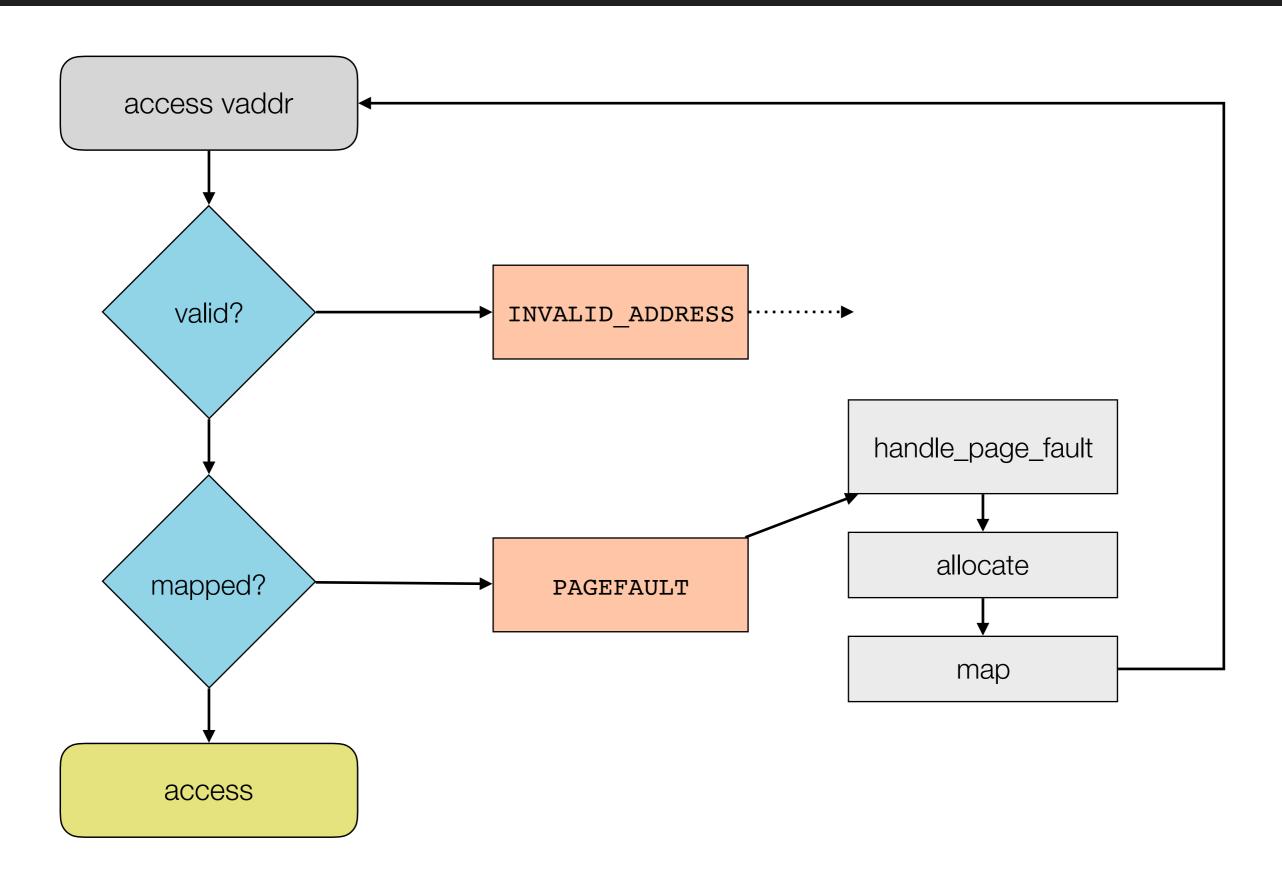
Paging on Demand in Selfie

- process accesses an virtual address (ex. do 1d)
 - check if virtual address is valid
 - within vaddr space, word-addressed
 - is valid virtual address
 - check if page for virtual address is mapped
 - is virtal address mapped, is page mapped
 - if both checks **pass**, translate address and load from physical memory
 - tlb, load virtual memory, load physical memory
 - otherwise an exception is thrown
 - EXCEPTION PAGEFAULT, EXCEPTION INVALID ADDRESS

Paging Page Fault

- page faults are handled by the OS
 - page frame is allocated iff available → palloc
 - virtual page is mapped to this page frame → map_page
 - OS returns control to process
 - pc not increased
 - process executes same instruction again succesfully
 - handle page fault

Paging On-Demand in Selfie



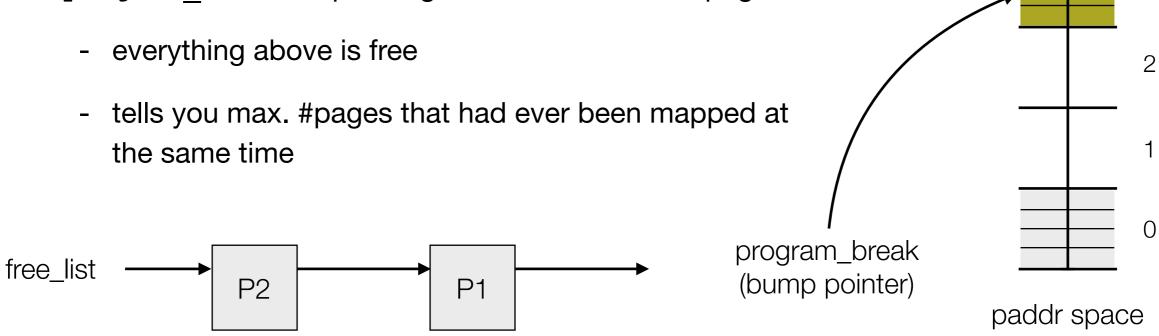
Paging Page Frame Allocation

Problem → Which frames can be used?

- the set of available frames can be partitioned int used and free frames
 - used_page_frame_memory, free_page_frame_memory
- page frame allocator
 - keeps track of free pages
 - used memory is known by the application that uses the allocator (information in page table)
 - using a free list
- in selfie there are only 2 pointers to manage free memory

Paging Page Frame Allocation in Selfie

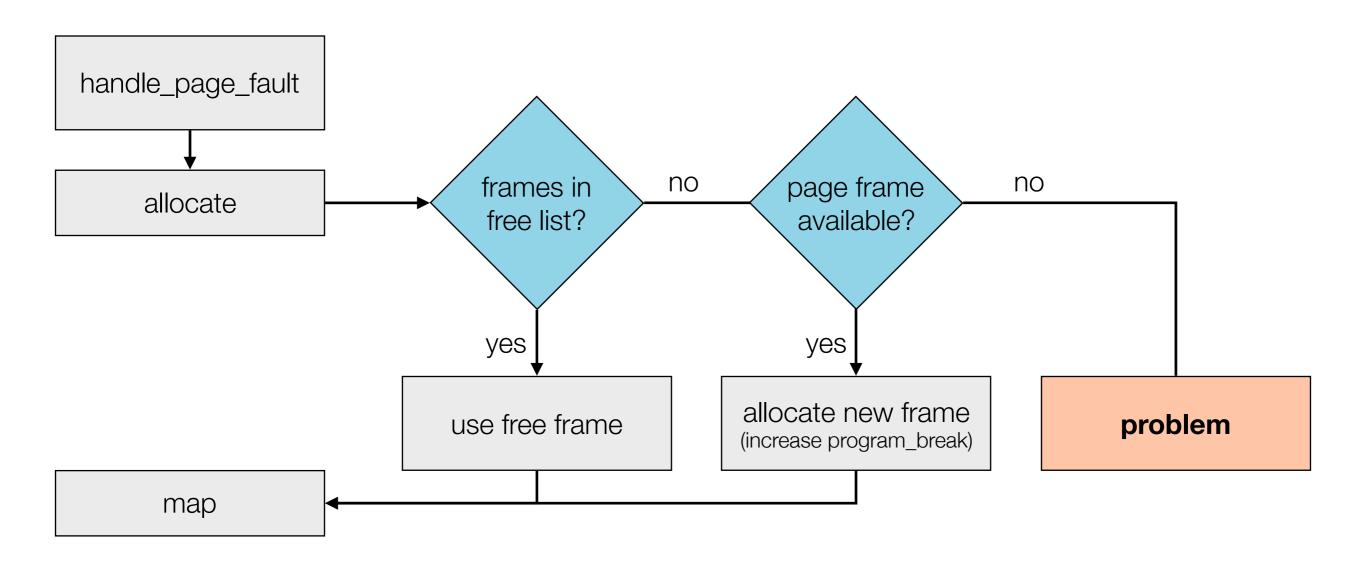
- selfie uses a simple memory allocation concept
 - bump pointer allocation
- selfie maintains 2 pointers to keep track of free memory
 - initialized to 0
 - free_list → pointing to a simply linked list
 - containing pages that became free
 - program_break → pointing to the last allocated page



4

3

Paging Page Frame Allocation in Selfie



Paging Swapping

Problem → All frames are taken but we need a frame!

- ideal → find a frame that will never be used/accessed again and free it
 - unfortunately we cannot make such a statement
- we can however take the data from a frame and store it somewhere else
 - frame can be freed and data can be brought back if needed
- we split the problem:
 - What frame do we free?
 - How and where store data?

Paging Swapping

Page Replacement Problem → What frame do we free?

- best frame would be the one that will be accessed furthest in the future
- best we can do is find an approximation to that frame
- different eviction strategies, depending on assumption we make
 - most recently used (MRU)
 - free the frame that was touched most recently
 - assumption → it probably will not be touched in a long time
 - least recently used (<u>LRU</u>)
 - free the frame that has not been touched the longest
 - assumption → bet on **locality**

Paging Bet on Locality

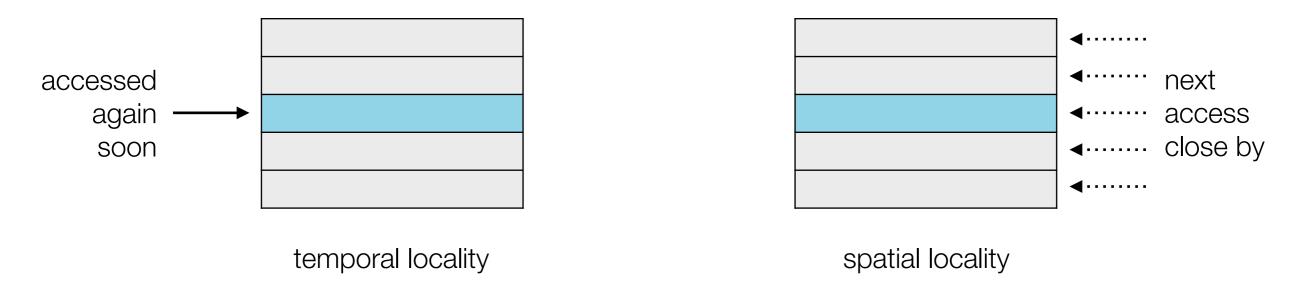
choosing LRU we bet on a property of code - <u>locality</u>

temporal locality

 when a memory location (address) is accessed, it is likely that the same memory location is accessed again in the near future

spacial locality

- when a memory location is accessed, it is likely that the next memory access happens in the 'neighbourhood'
- most code shows spatial locality → not much jumping around



Paging Swapping

Swapping → How and where store data so we can get it back?

'page out'

- move data (frame content) from main memory to swap space on disk
- mark page table entry as 'not available' and remember address of data in swap space

'page in'

- access a swapped out page → page fault
- move data from swap space back into main memory → page frame allocation (probably causing a 'page out')

Paging Page Fault

A <u>Page Fault</u> is an **exception** that is raised when a process tries to access a page that is **not mapped**.

- page faults decrease performance do 'administrative work'
 - fetch from disk or allocate, map,...
- when virtual memory is overused the number of page faults increases dramatically (page-in, page-out)
- thrashing
 - more time spent on 'administrive work' than on process progress/execution
 - performance of machine degrades or collapses

Summary OS Introduction Continued

- We addressed the second problem
 - several processes using the same fixed size memory
- the OS solves this problem by managing memory
 - memory is shared among processes
 - OS virtualizes memory → illusion of private memory
- concepts and mechanisms used
 - abstraction/illusion → physical address space, virtual address space (illusion of private memory), address translation
 - virtualization techniques → base & bound, segmentation, paging
 - paging → on demand, page faults, page frame allocation, swapping (page replacement strategies),